

COMPUTER VISION AND
PHOTOGRAMMETRY

FEATURE DETECTION

CONTENTS TODAY

- ▶ Feature detection
- ▶ Harris Corners
- ▶ Scale-space, scale invariant.
- ▶ DoG
- ▶ Lab4

LET'S RECAP . . .

FEATURES

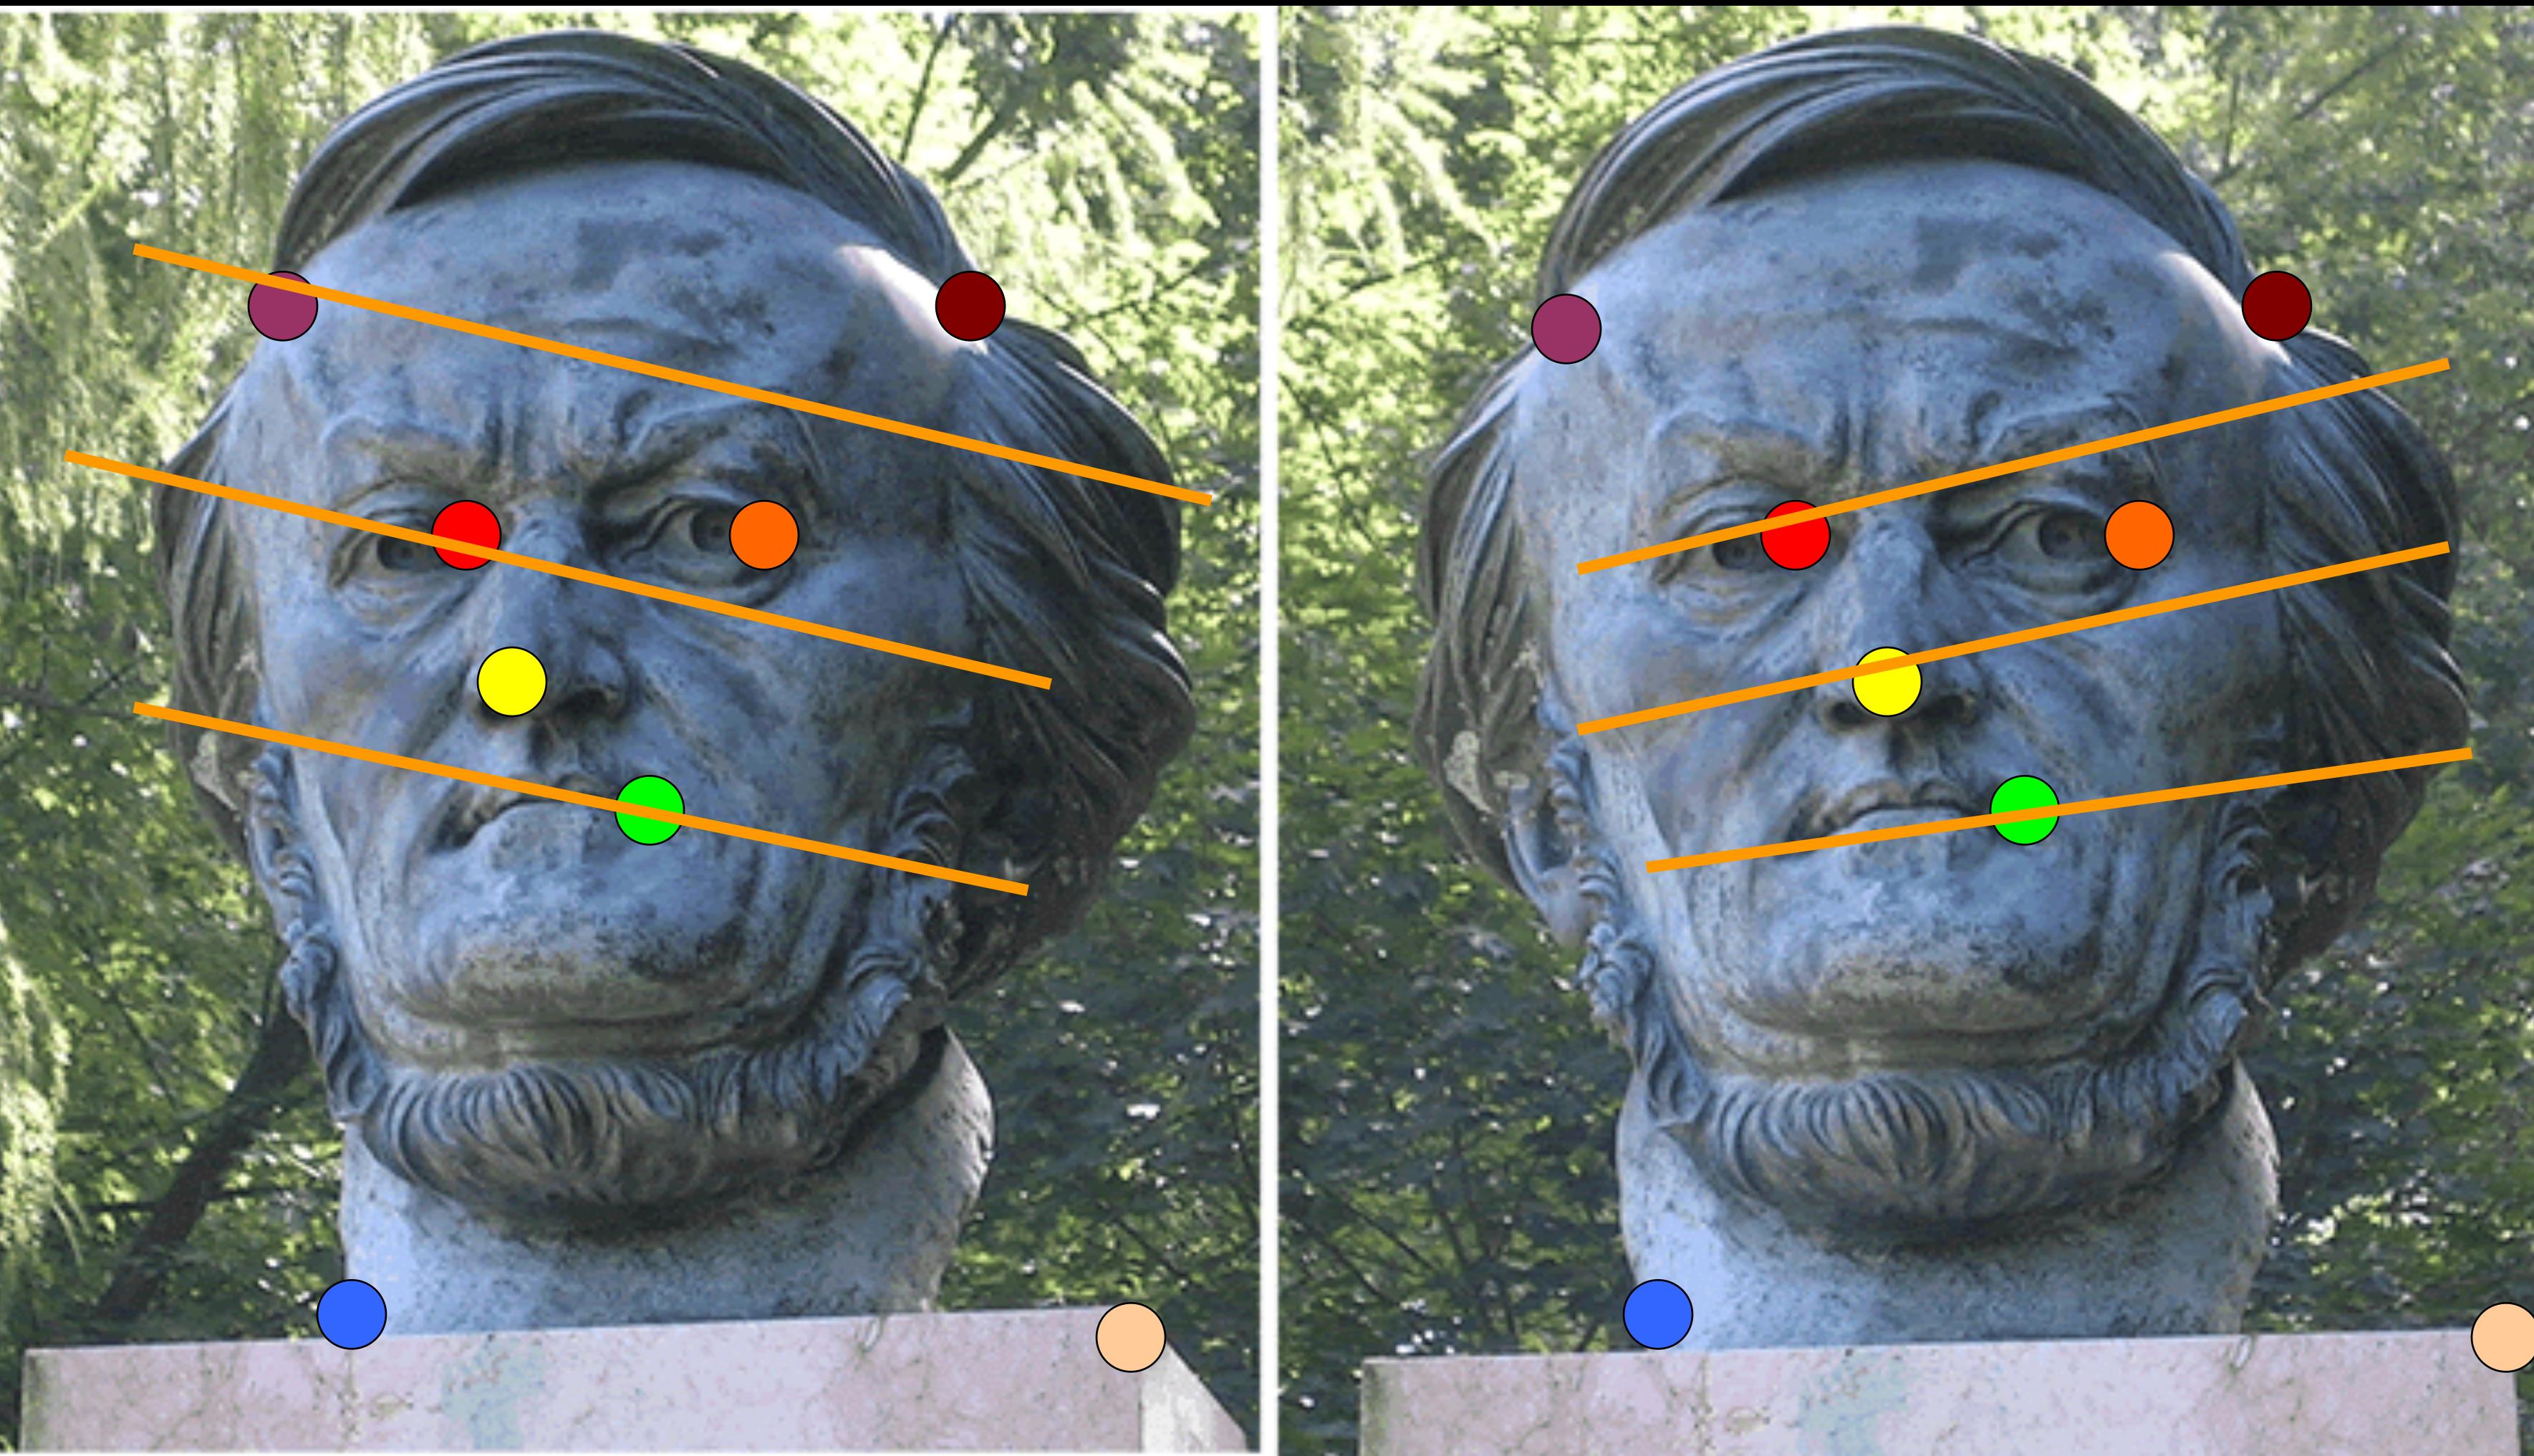
CORRESPONDENCE ACROSS VIEWS

- ▶ Correspondence: matching points, patches, edges, or regions across images



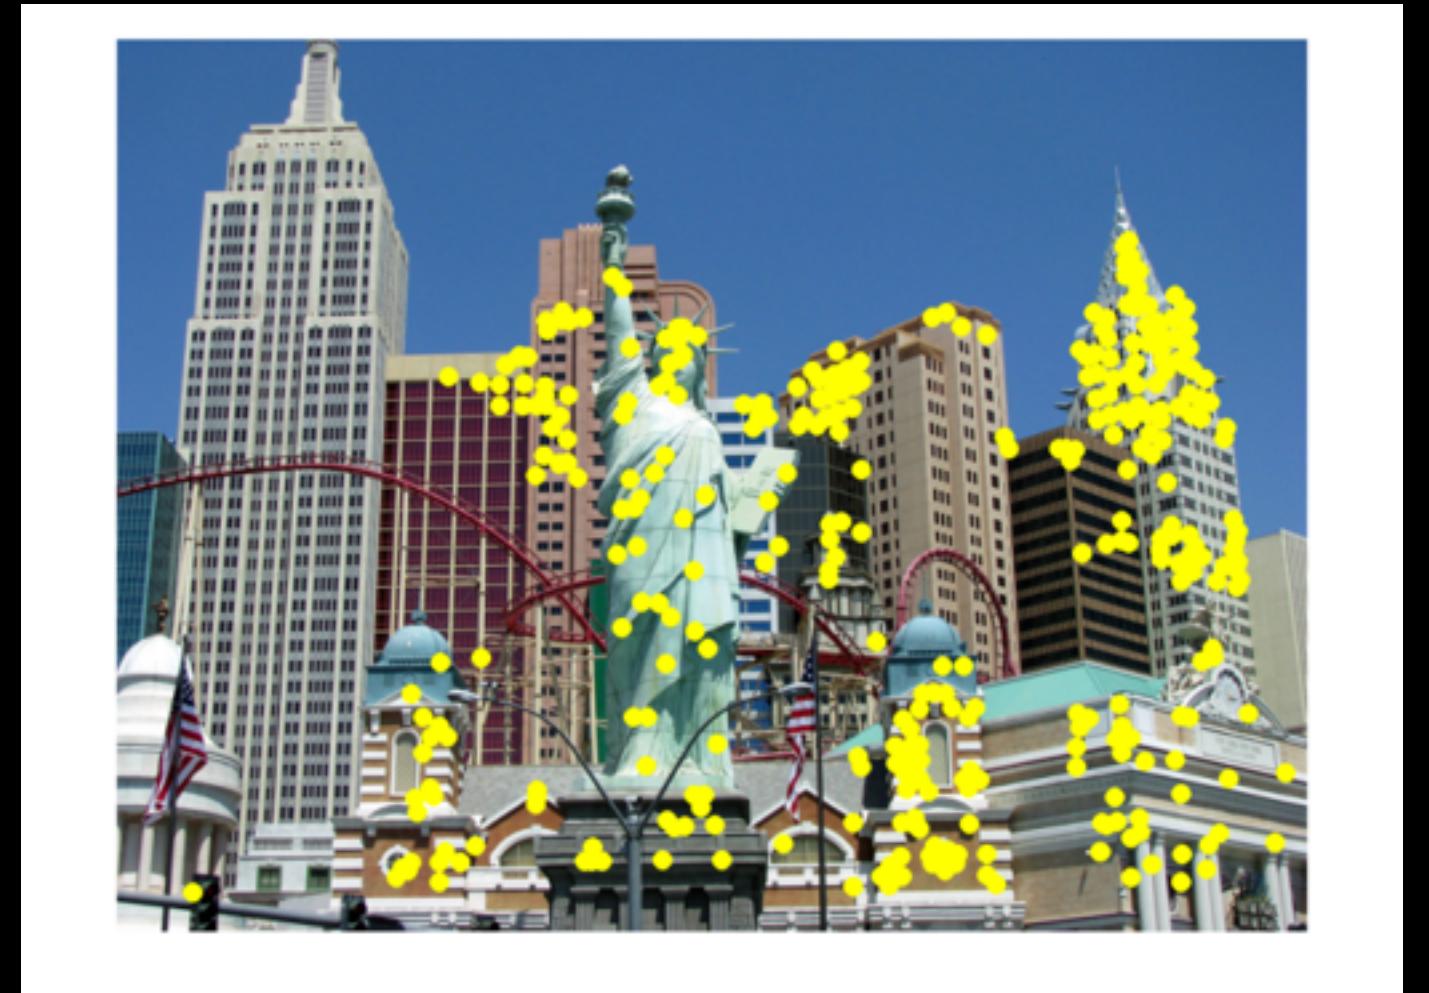
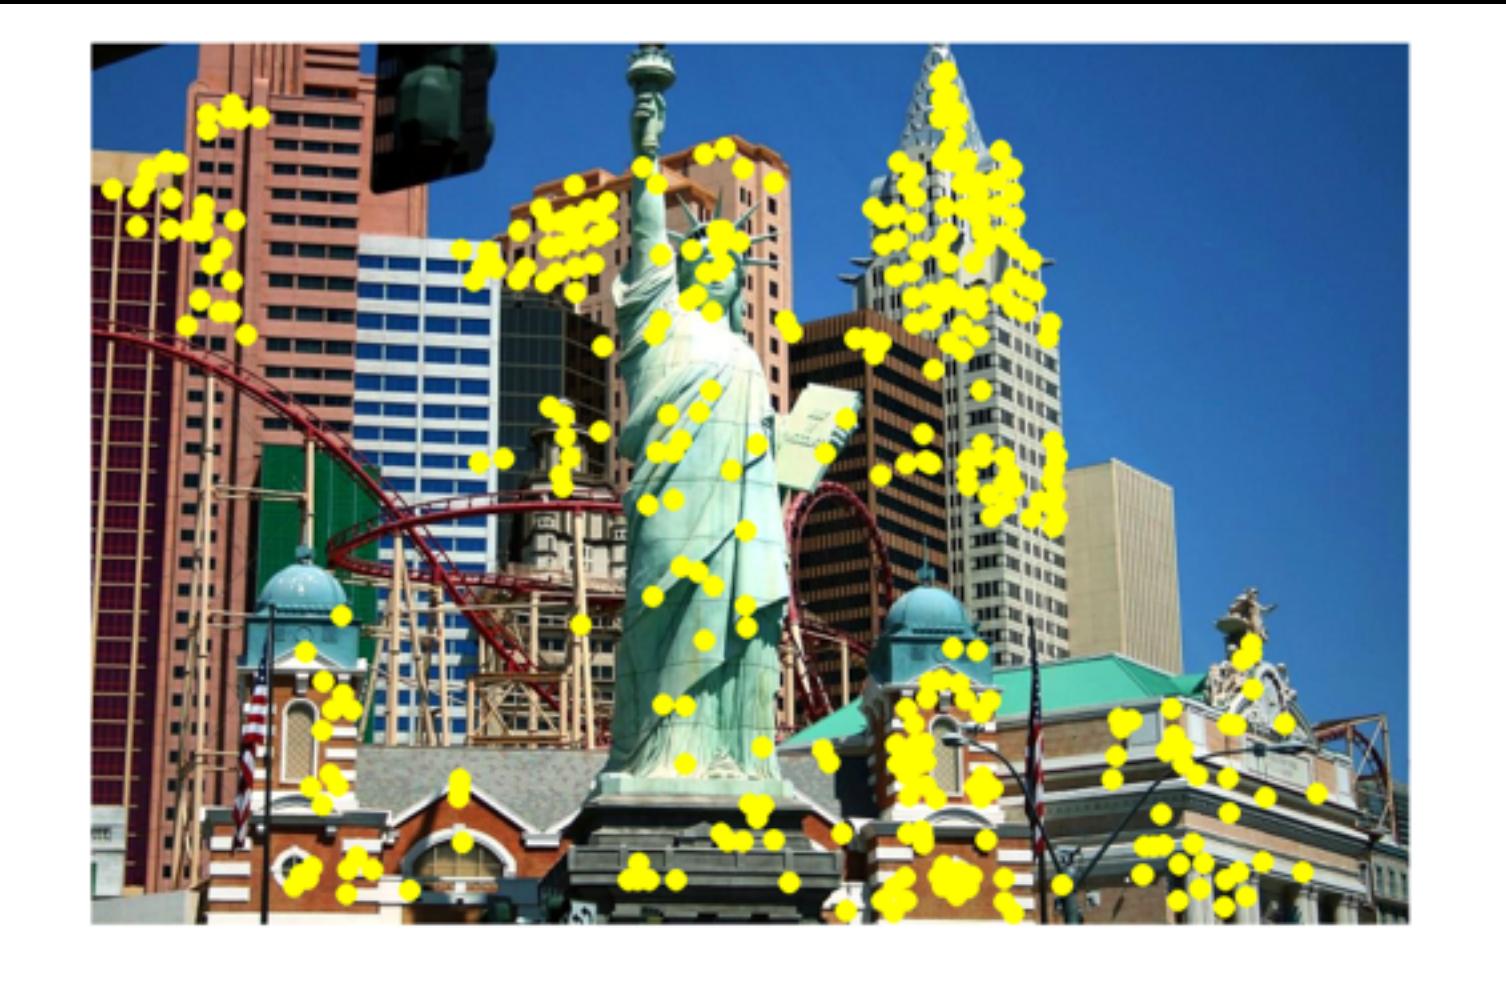
TEXT

EXAMPLE: ESTIMATING “FUNDAMENTAL MATRIX” THAT CORRESPONDS TWO VIEWS



APPLICATIONS

- ▶ Feature points are used for:
 - ▶ Image alignment
 - ▶ 3D reconstruction
- ▶ Motion tracking
- ▶ Robot navigation
- ▶ Indexing and database retrieval
- ▶ Object recognition



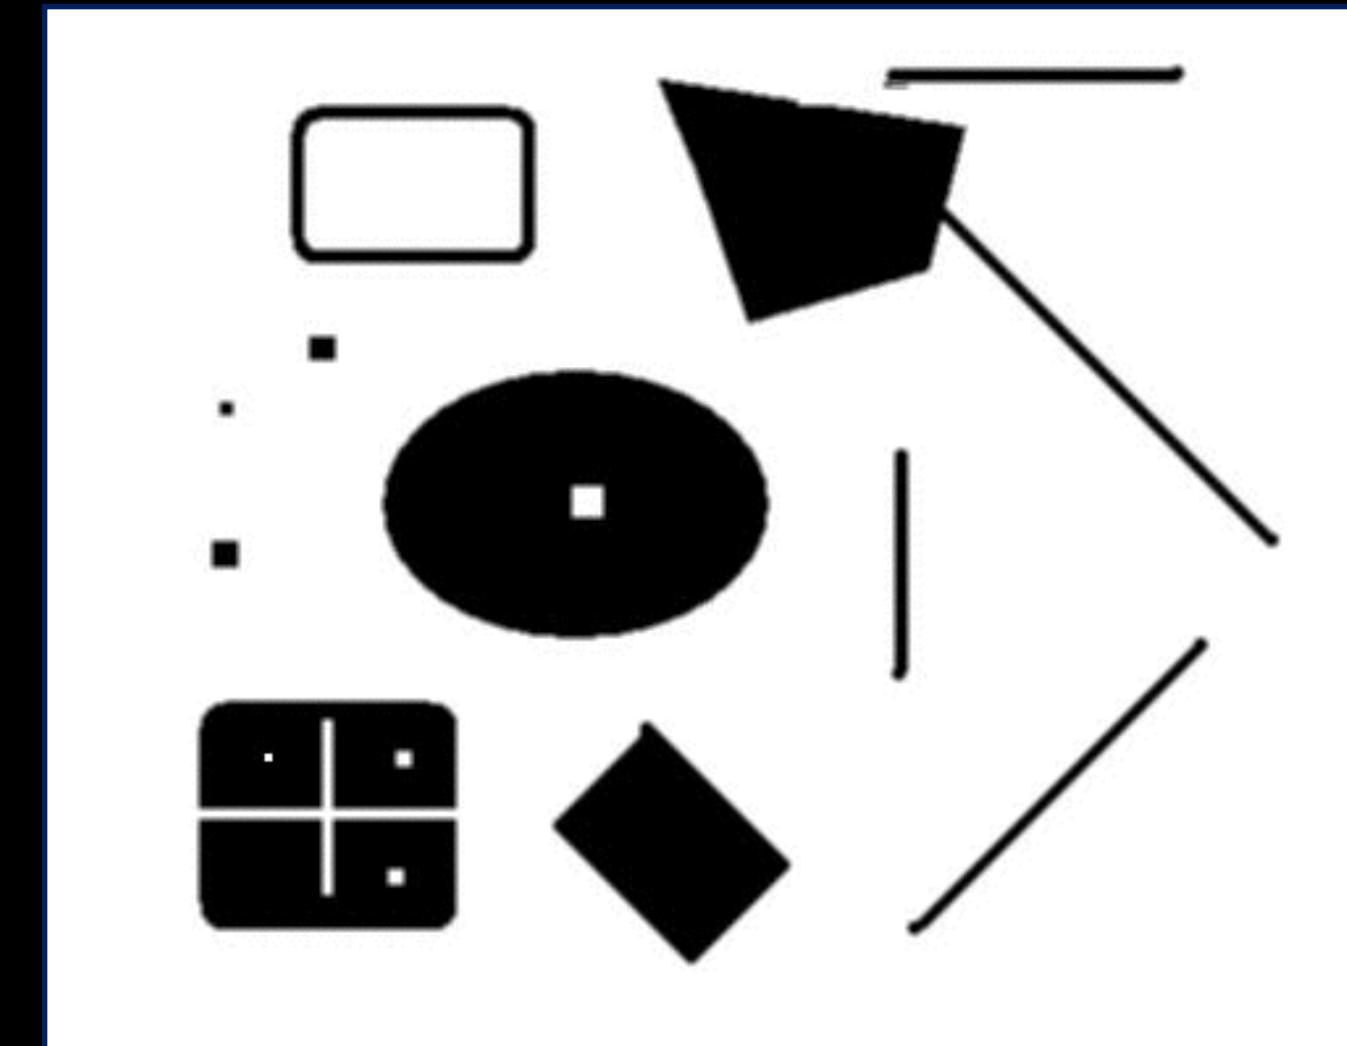
THIS CLASS: INTEREST POINTS AND LOCAL FEATURES

- ▶ Note: “interest points” = “keypoints”, also sometimes called “features”
- ▶ Edges and lines are also features
- ▶ Different applications need different types of features
- ▶ Features with different properties

THIS CLASS: INTEREST POINTS

- ▶ Suppose you have to click on some point, go away and come back after I deform the image, and click on the same points again.
- ▶ Which points would you choose?

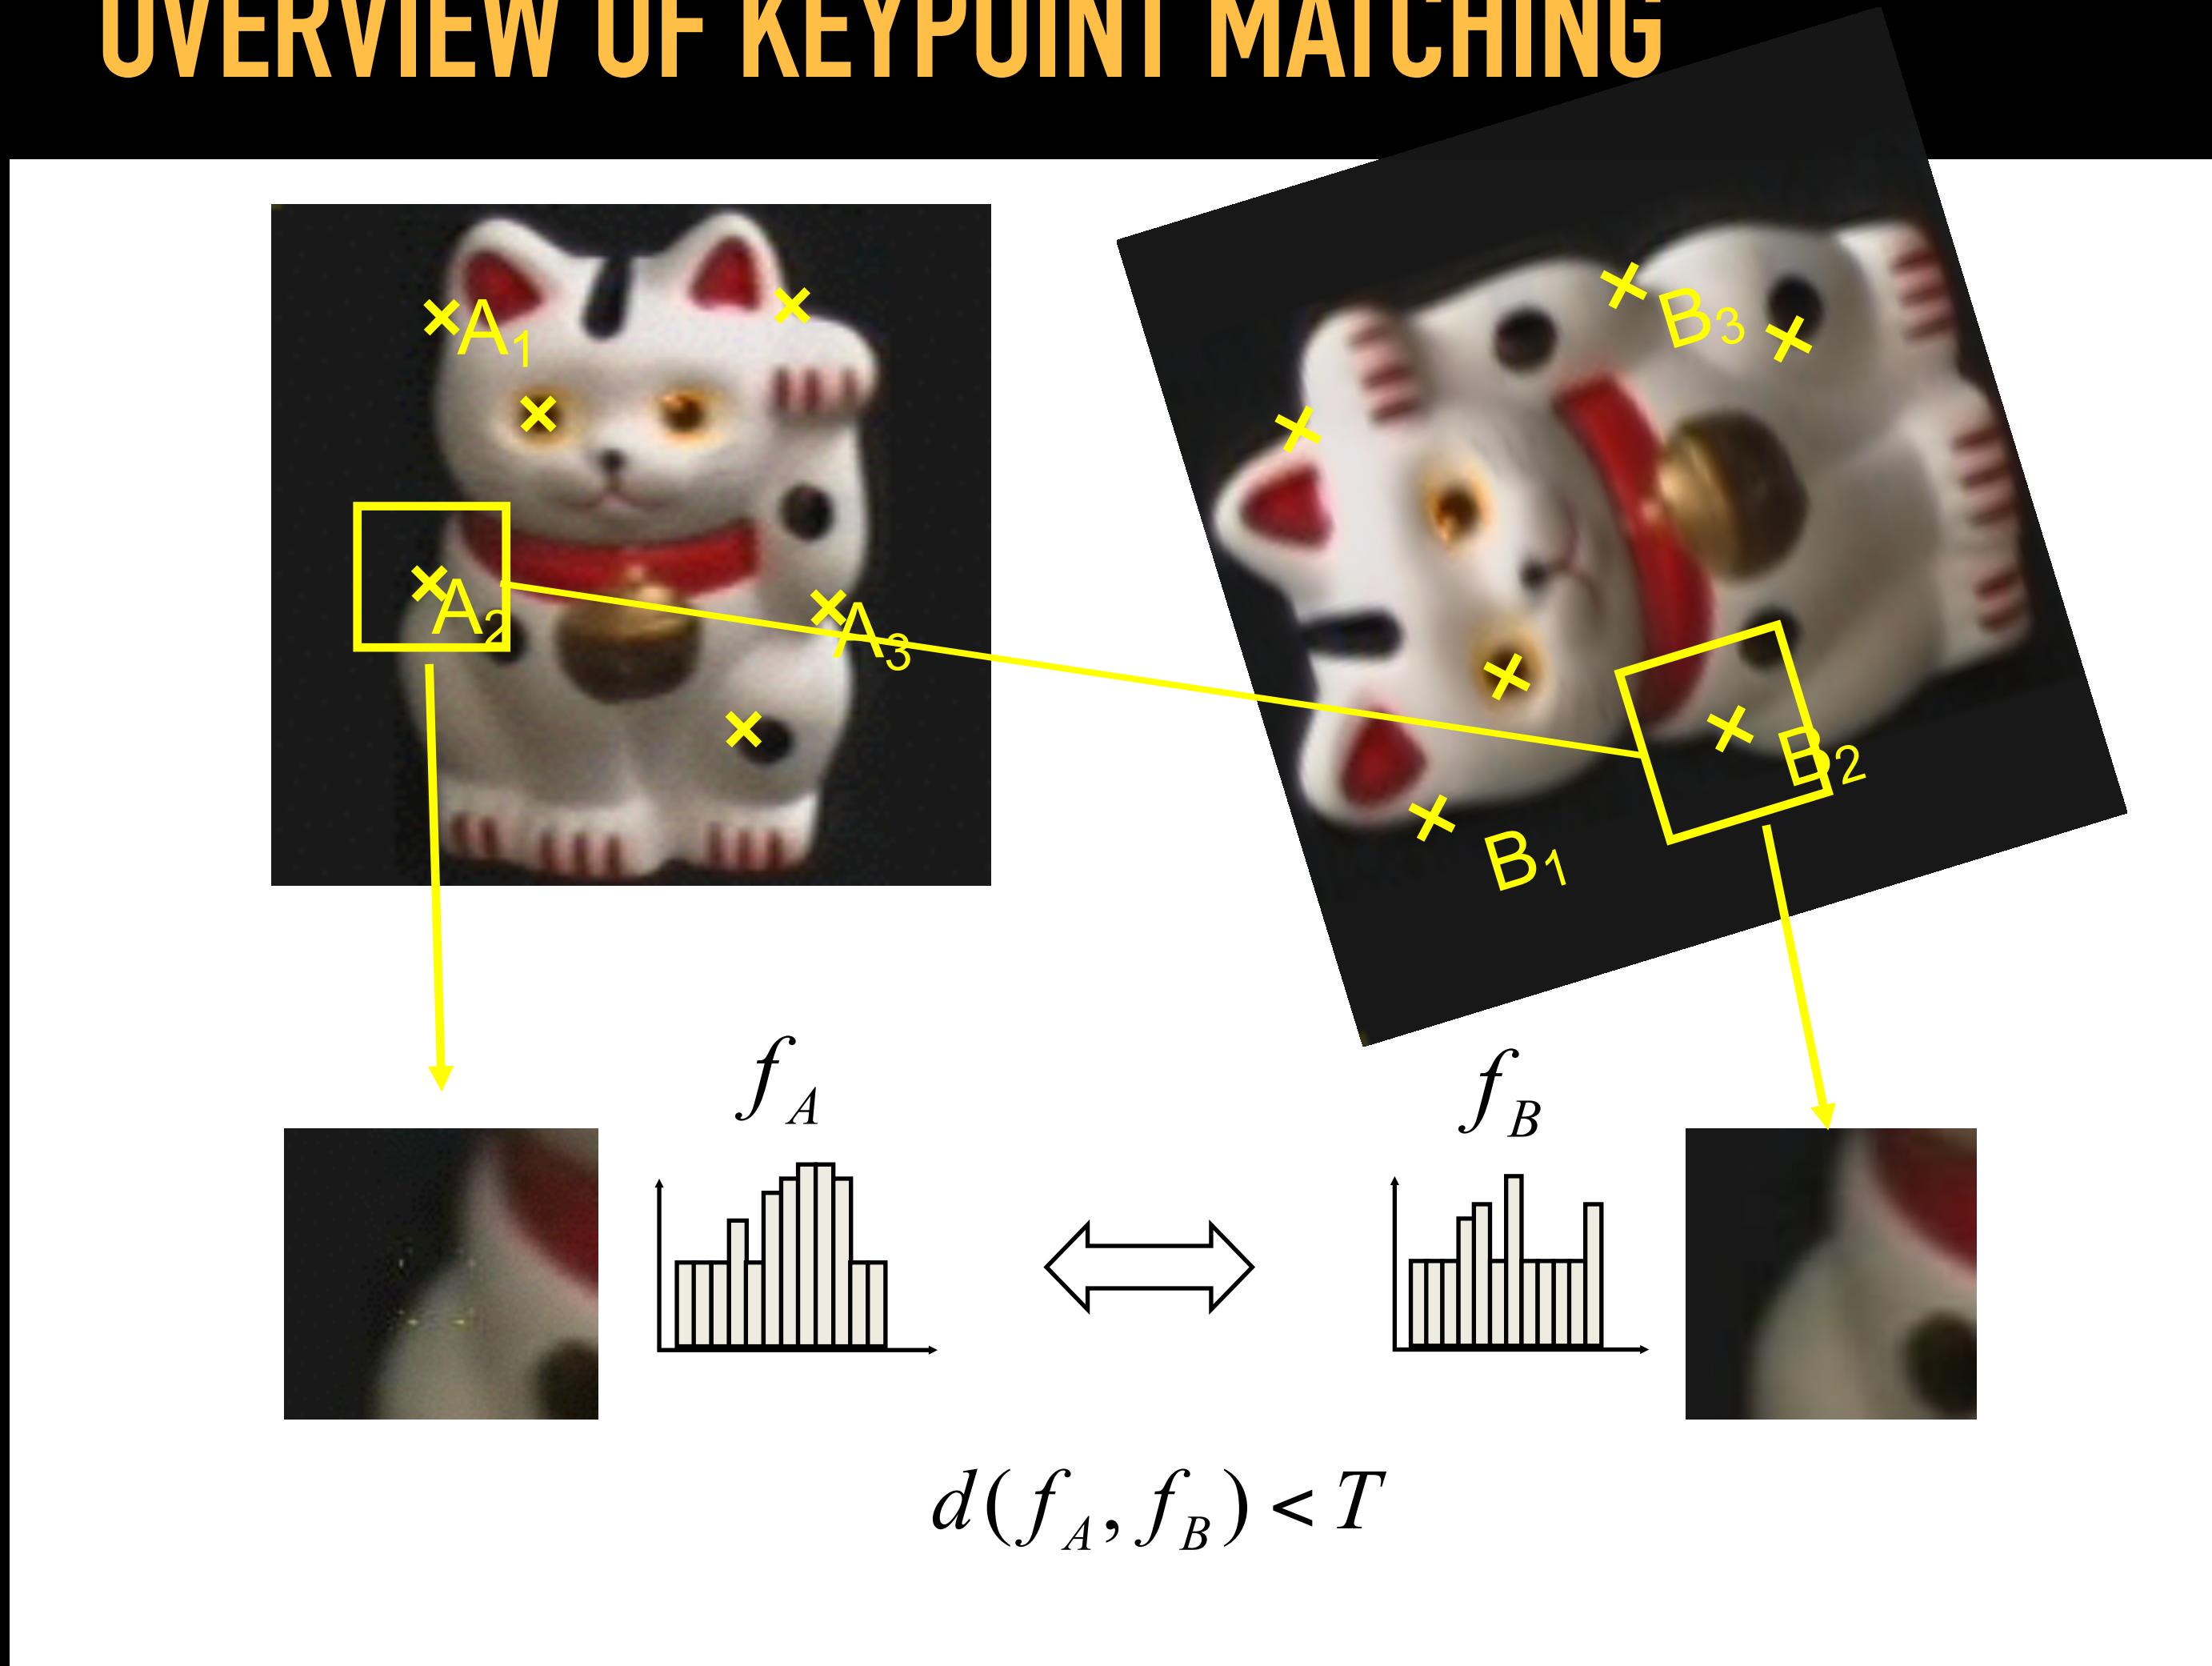
ORIGINAL



DEFORMED



OVERVIEW OF KEYPOINT MATCHING



TEXT

GOALS FOR KEYPOINTS



Detect points that are *repeatable* and *distinctive*

INVARIANT LOCAL FEATURES

- ▶ Image content is transformed into local feature coordinates that are invariant to translation, rotation, scale, and other imaging parameters



WHY EXTRACT FEATURES?

- ▶ Motivation: panorama stitching
- ▶ We have two images - how do we combine them?



CHARACTERISTICS OF GOOD FEATURES

- ▶ **Repeatability**

- ▶ The same feature can be found in several images despite geometric and photometric transformations

- ▶ **Saliency**

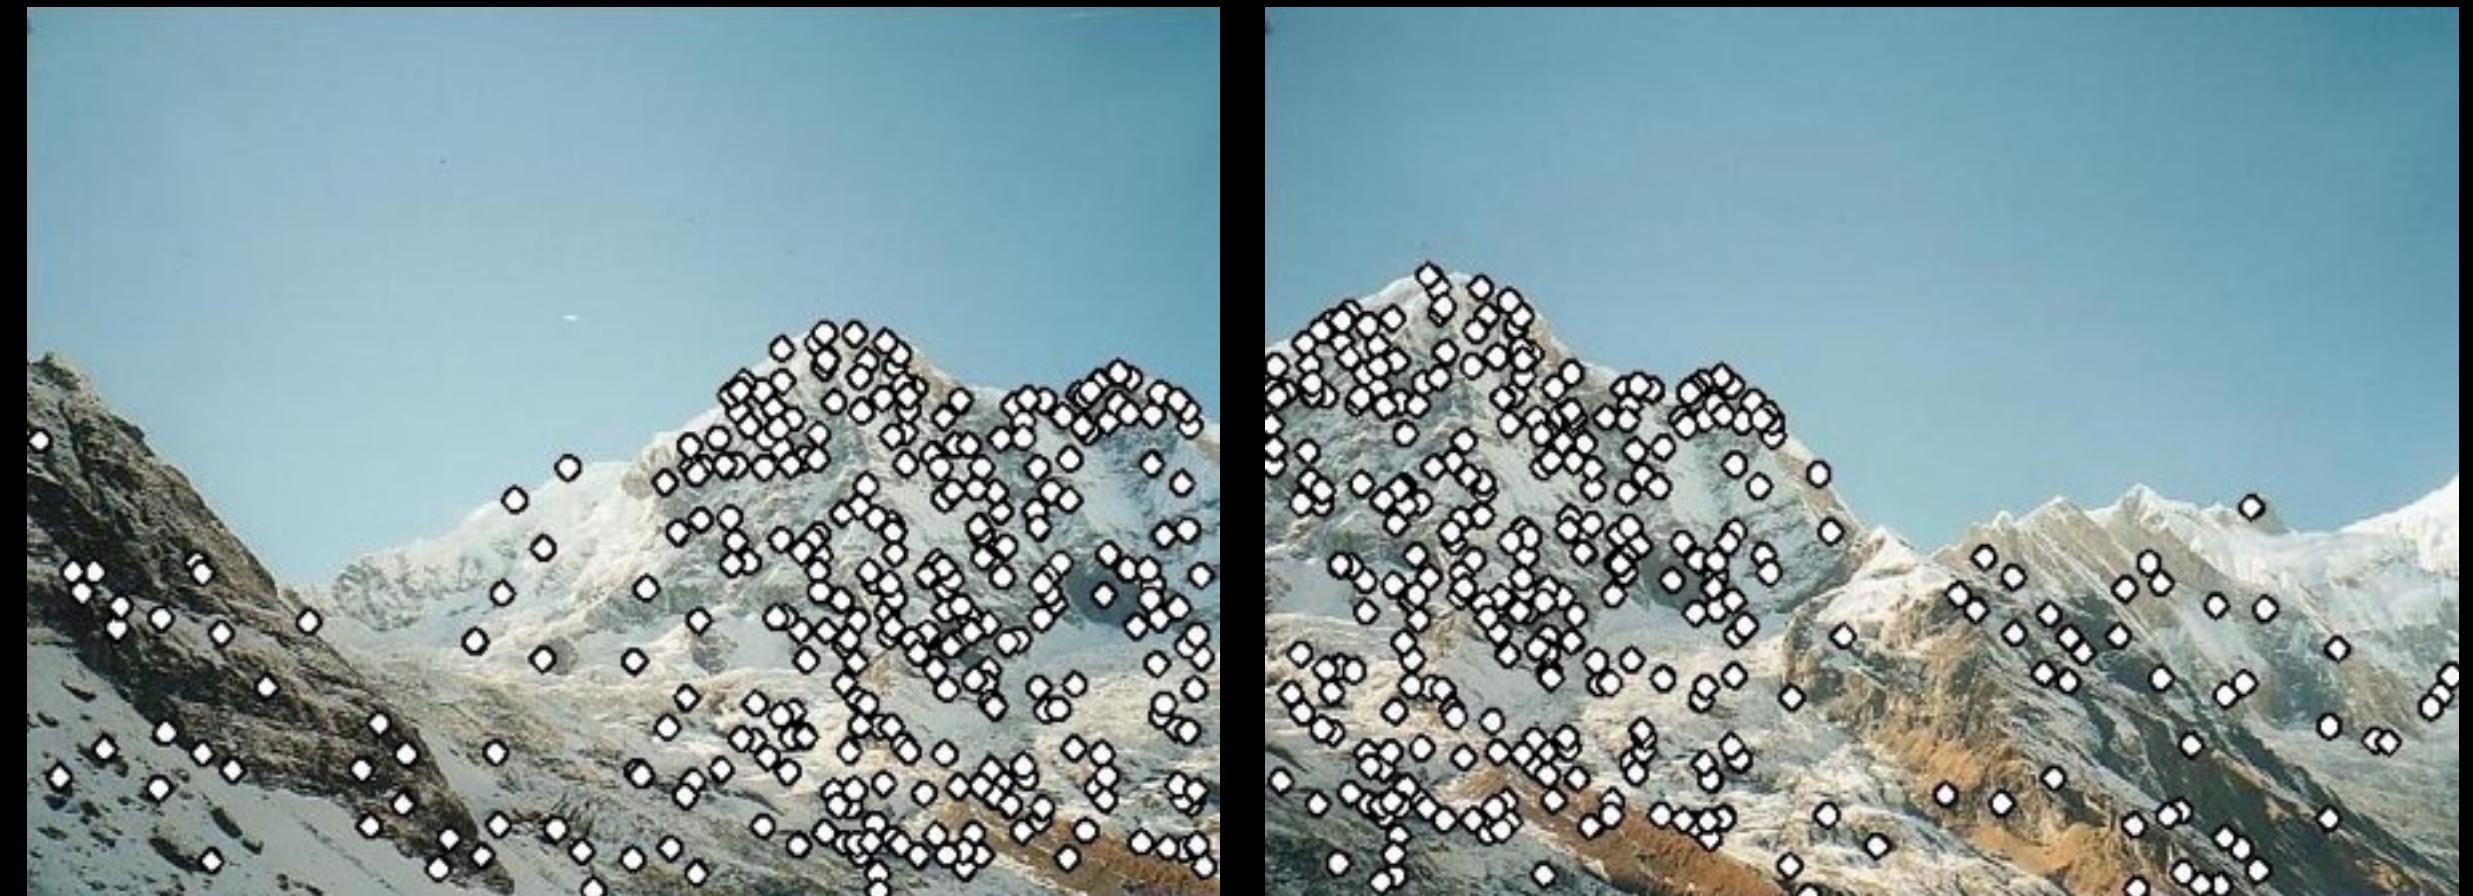
- ▶ Each feature is distinctive

- ▶ **Compactness and efficiency**

- ▶ Many fewer features than image pixels

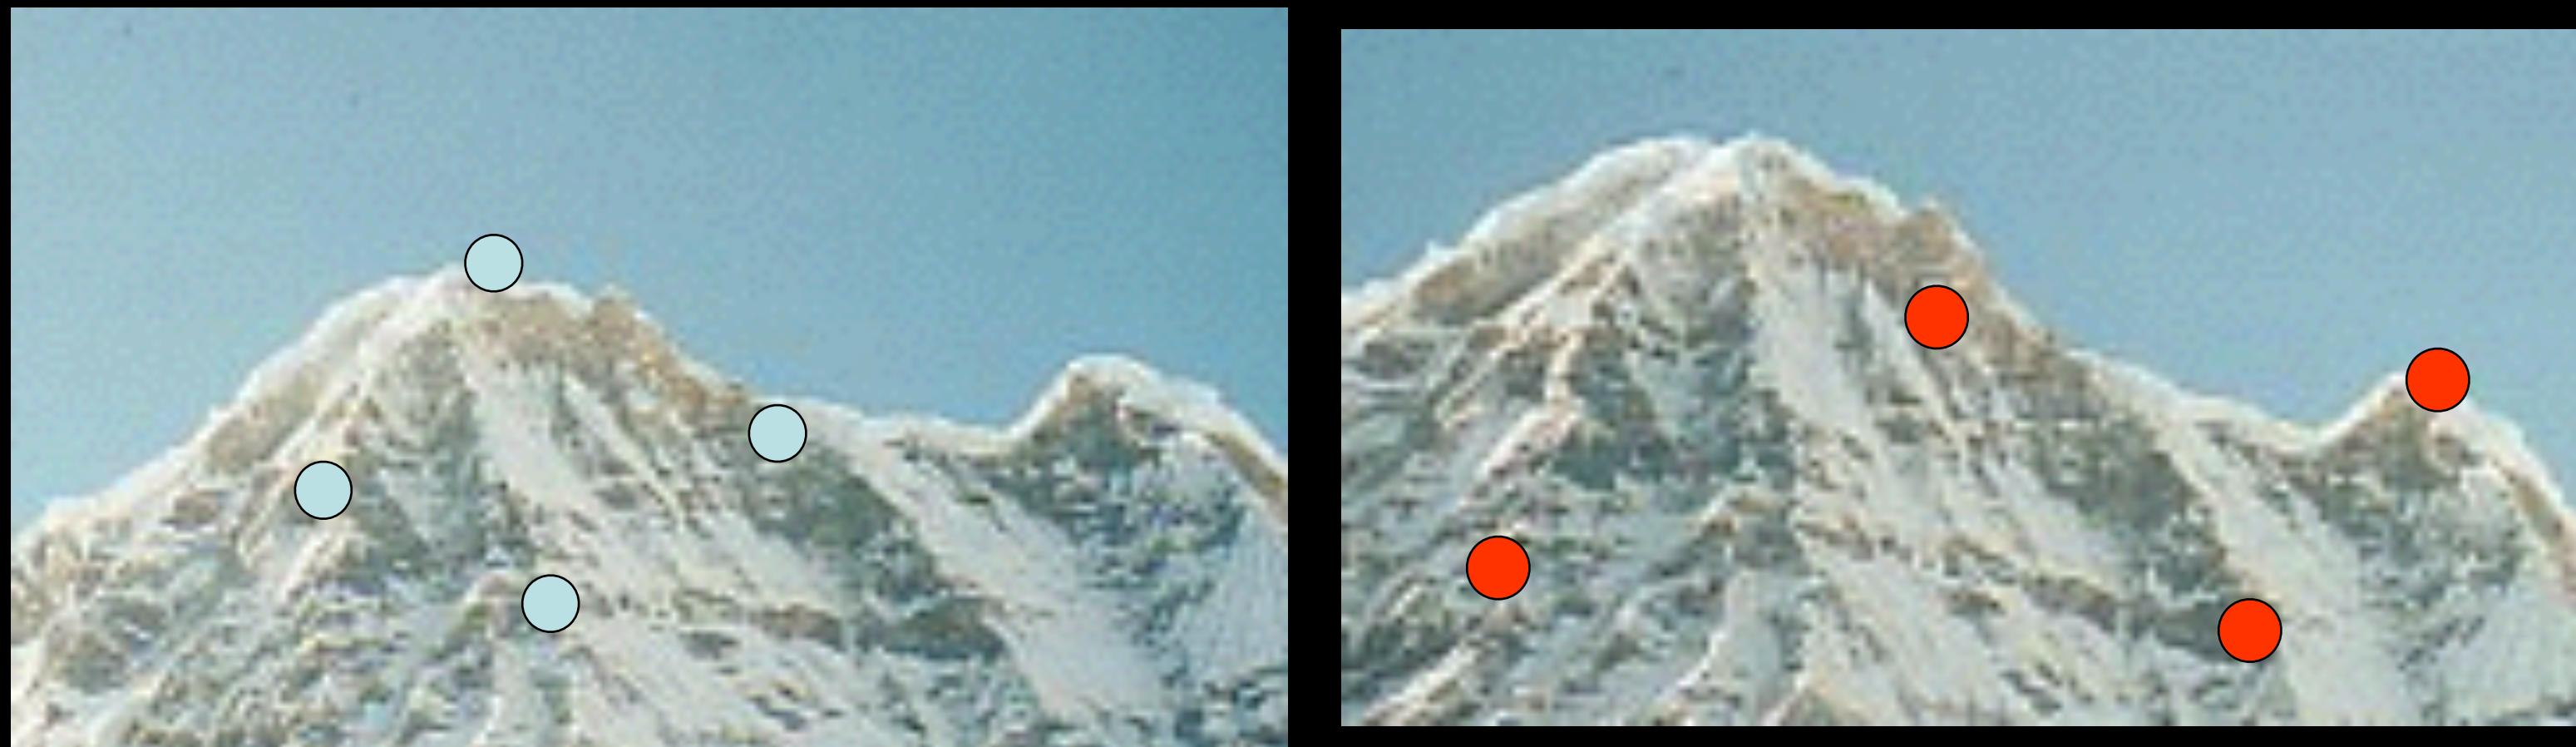
- ▶ **Locality**

- ▶ A feature occupies a relatively small area of the image; robust to clutter and occlusion



GOAL: INTEREST OPERATOR REPEATABILITY

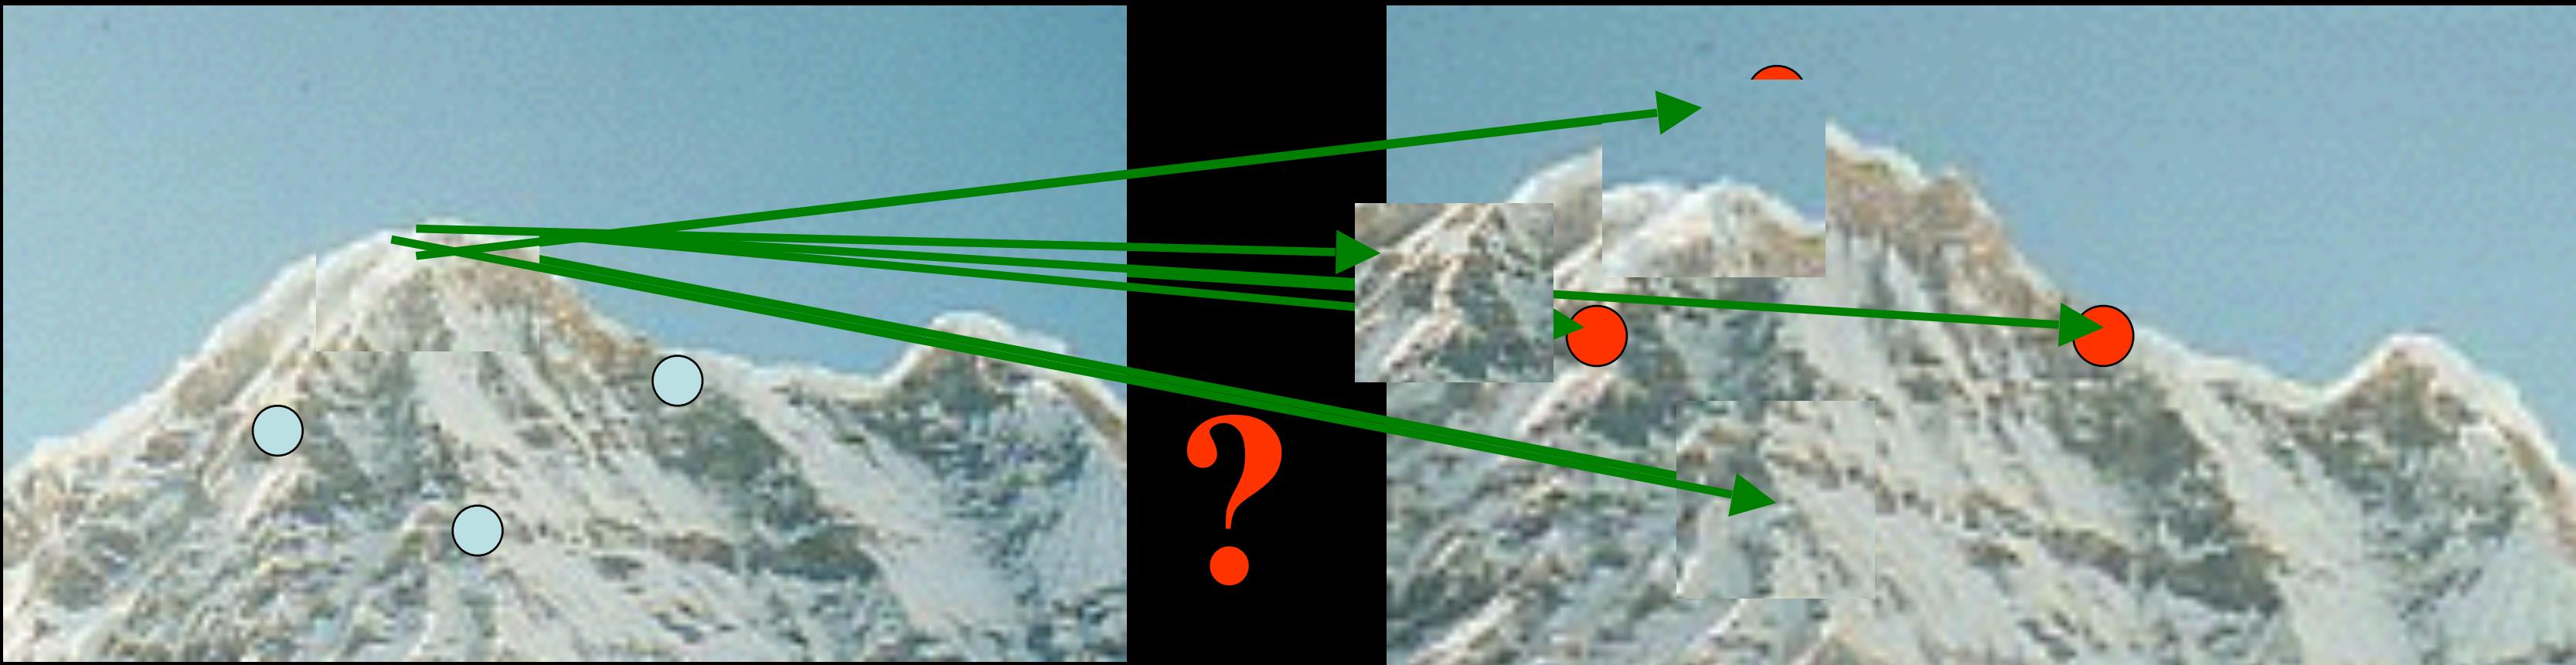
We want to detect (at least some of) the same points in both images.



Yet we have to be able to run the detection procedure independently per image.

GOAL: DESCRIPTOR DISTINCTIVENESS

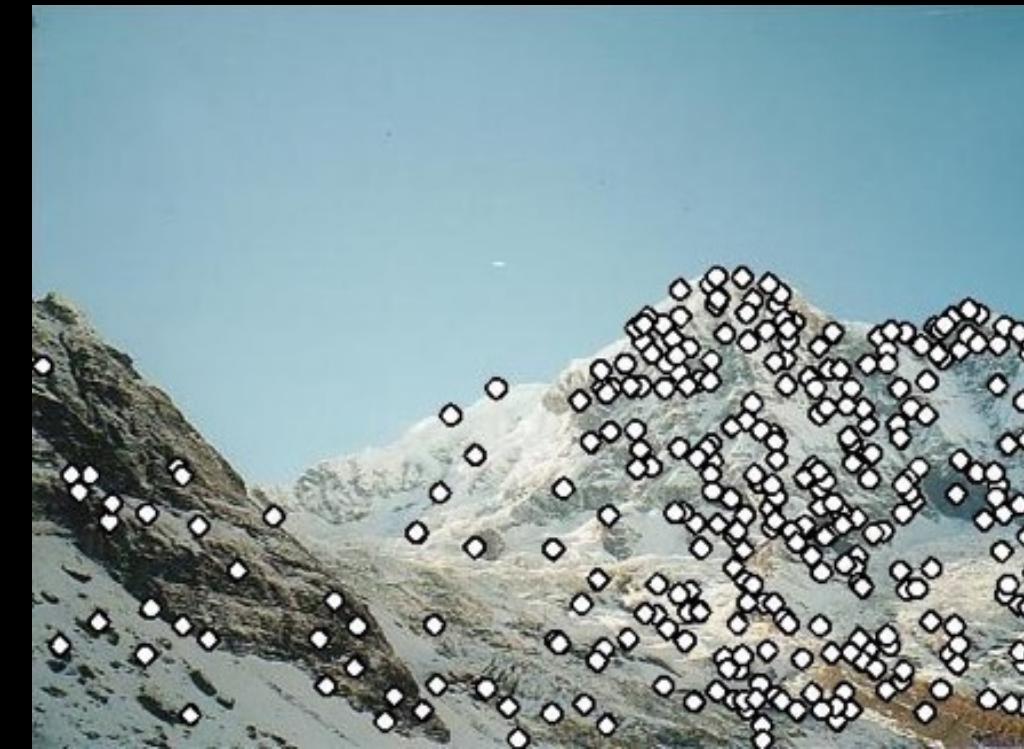
We want to be able to reliably determine which point goes with which.



Must provide some invariance to geometric and photometric differences between the two views.

LOCAL FEATURES: MAIN COMPONENTS

- ▶ Detection: Identify the interest points
- ▶ Description: Extract vector feature descriptor surrounding each interest point.
- ▶ Matching: Determine correspondence between descriptors in two views

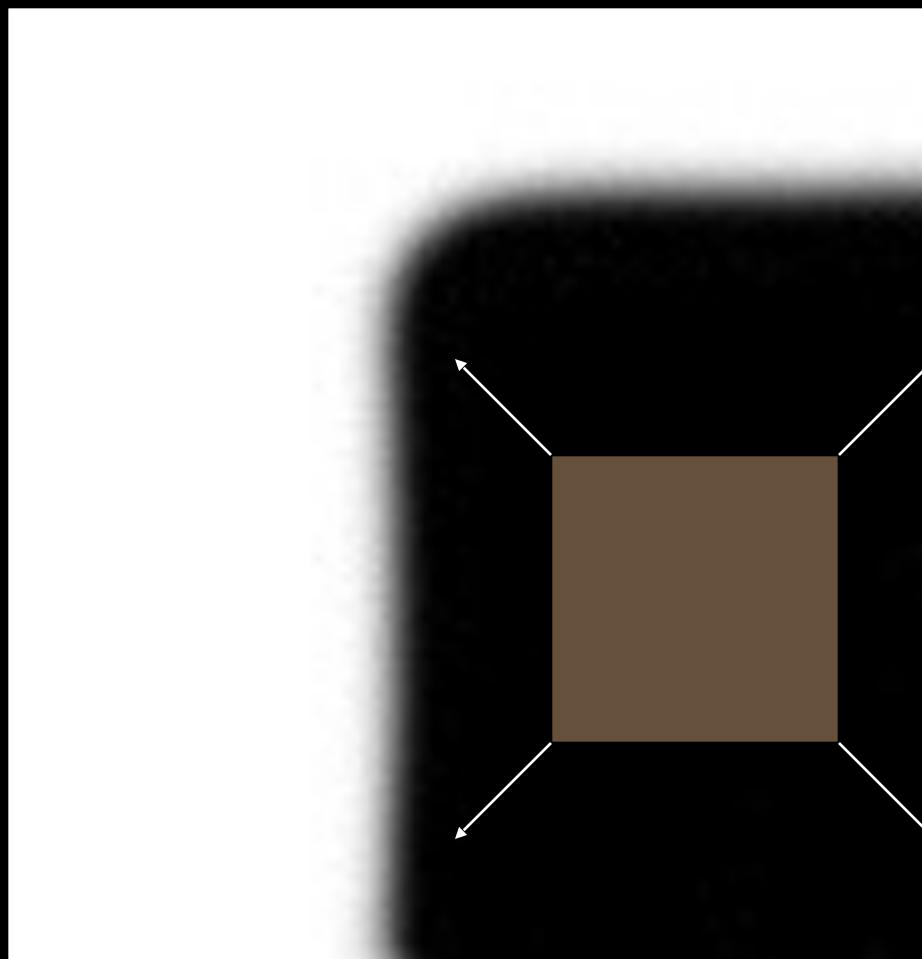


MANY EXISTING DETECTORS AVAILABLE

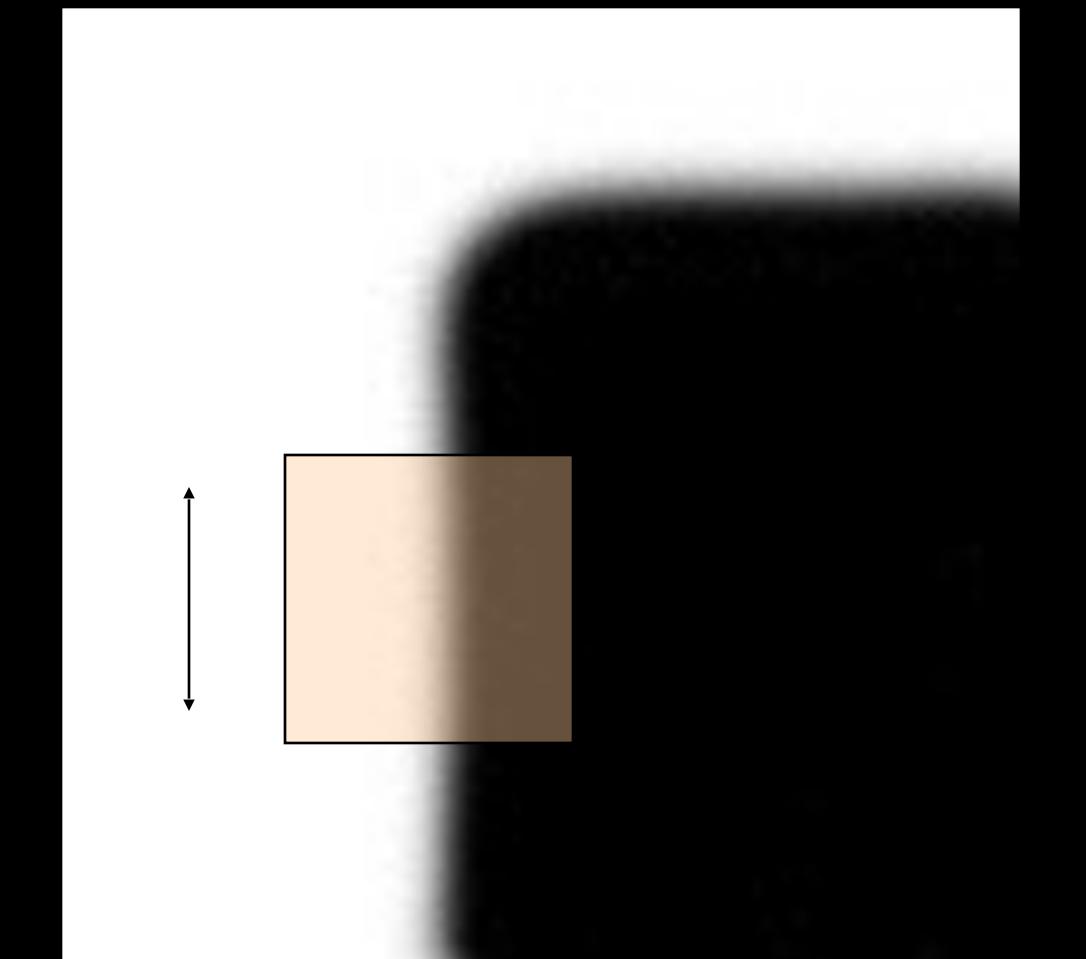
- ▶ Hessian & Harris [Beaudet '78], [Harris '88]
- ▶ Laplacian, DoG [Lindeberg '98], [Lowe 1999]
- ▶ Harris-/Hessian-Laplace [Mikolajczyk & Schmid '01]
- ▶ Harris-/Hessian-Affine [Mikolajczyk & Schmid '04]
- ▶ EBR and IBR [Tuytelaars & Van Gool '04]
- ▶ MSER [Matas '02]
- ▶ Salient Regions [Kadir & Brady '01]
- ▶ Others...

CORNER DETECTION: BASIC IDEA

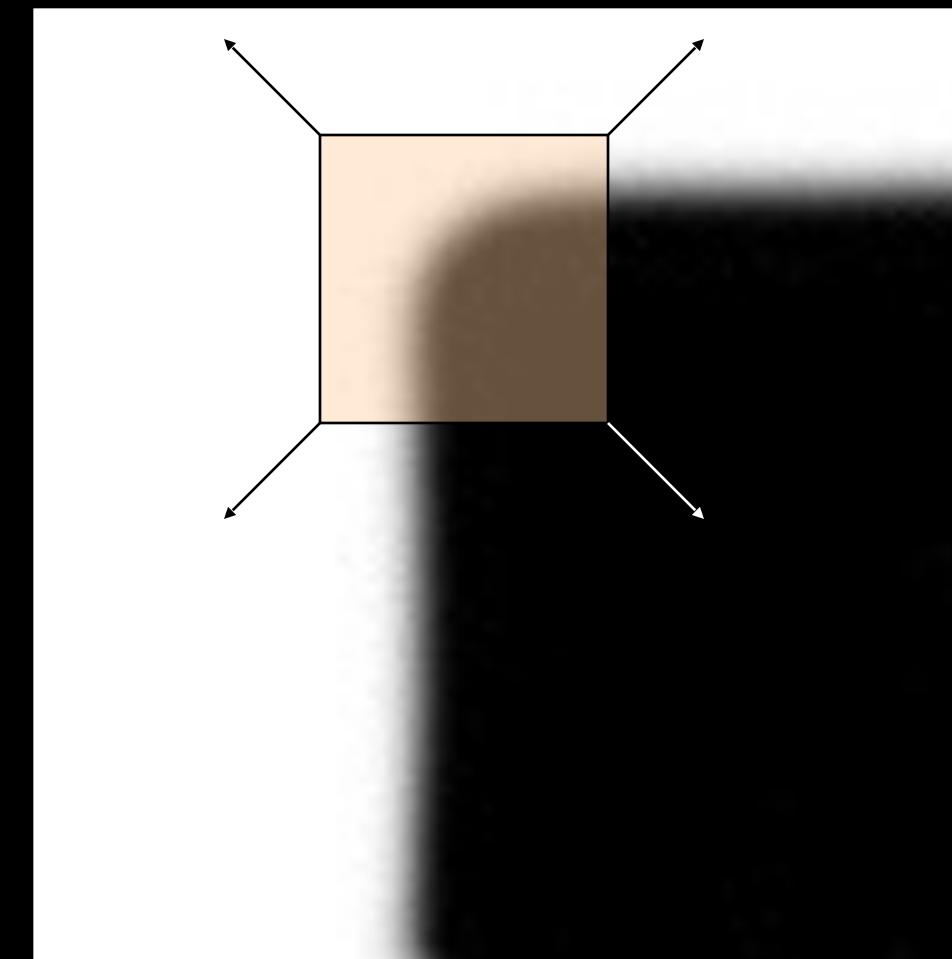
- ▶ We should easily recognize the point by looking through a small window
- ▶ Shifting a window in *any direction* should give *a large change* in intensity



“FLAT” REGION:
NO CHANGE IN ALL
DIRECTIONS



“EDGE”:
NO CHANGE ALONG THE EDGE
DIRECTION



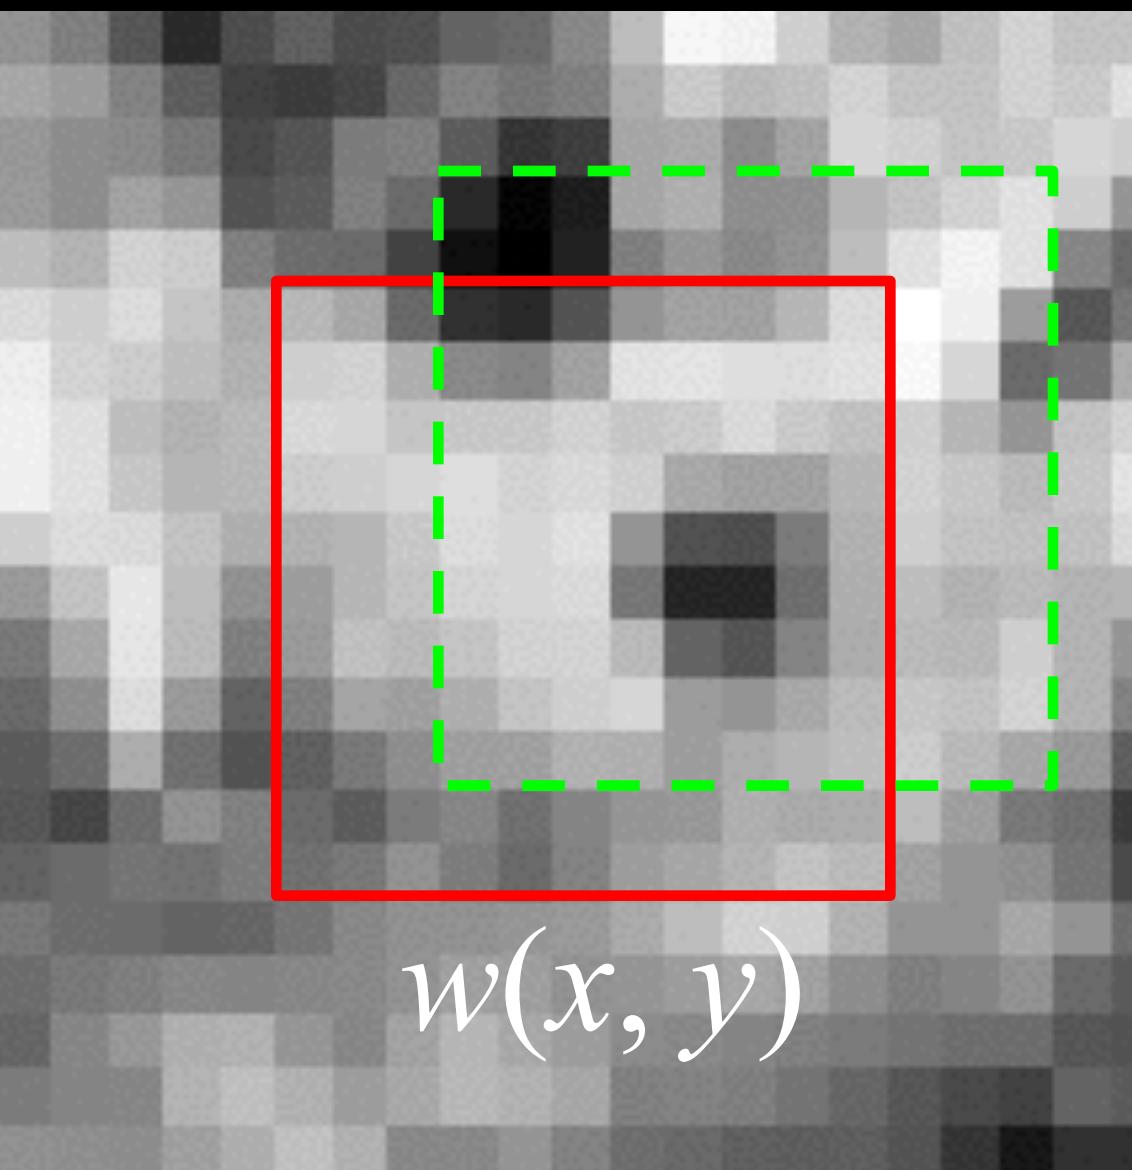
“CORNER”:
SIGNIFICANT CHANGE IN ALL
DIRECTIONS

CORNER DETECTION: MATHEMATICS

- ▶ Change in appearance of window $w(x,y)$ for the shift $[u,v]$:

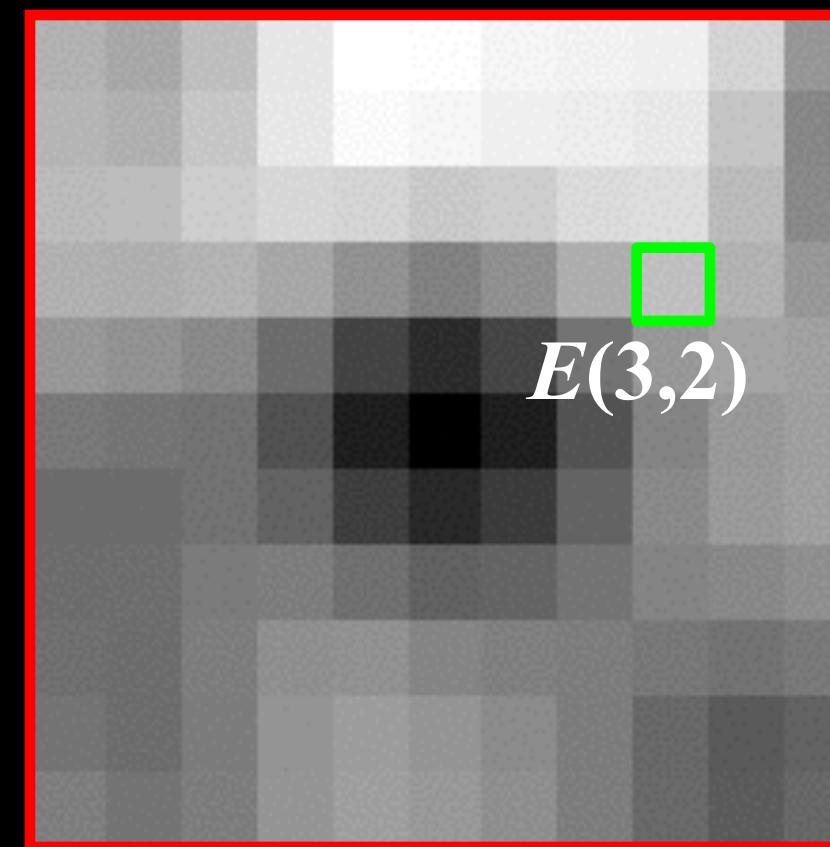
$$E(u, v) = \sum_{x,y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

$I(x, y)$



$w(x, y)$

$E(u, v)$



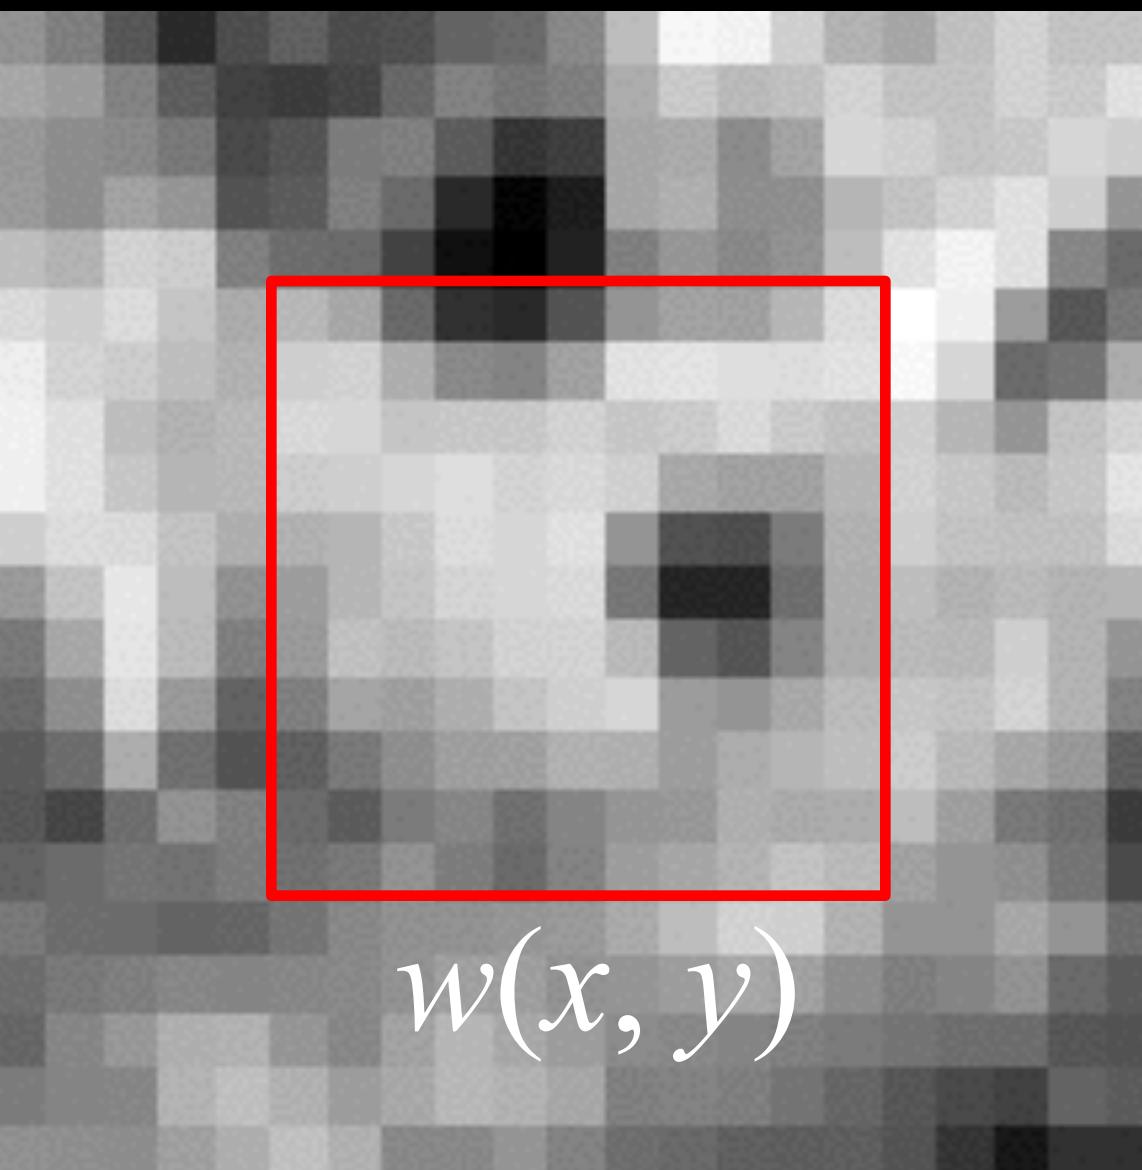
$E(3,2)$

CORNER DETECTION: MATHEMATICS

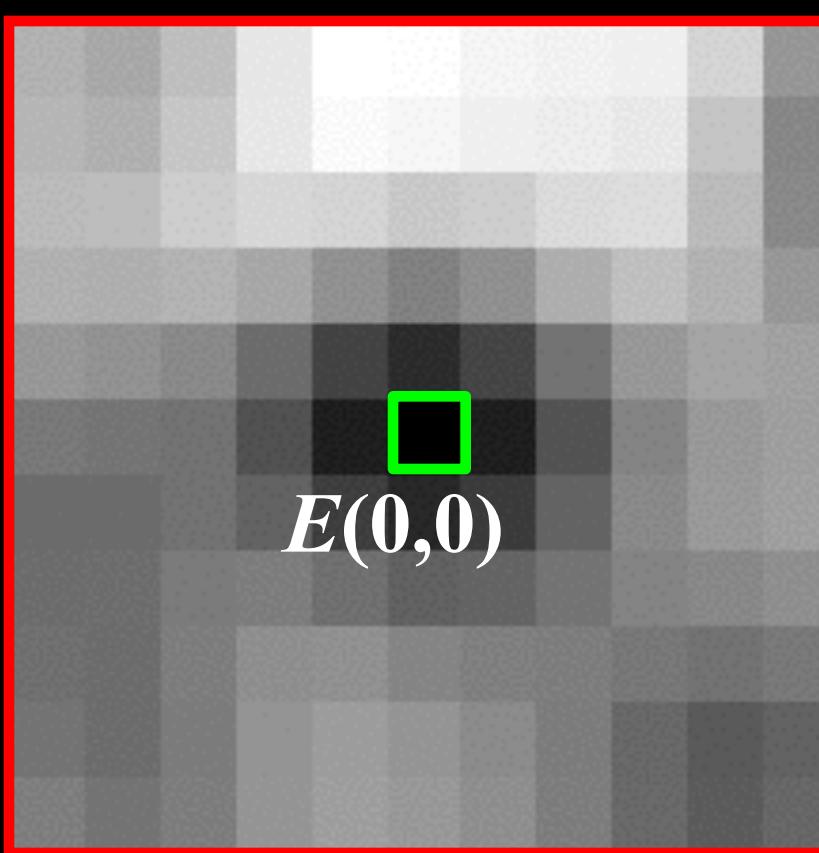
- ▶ Change in appearance of window $w(x,y)$ for the shift $[u,v]$:

$$E(u, v) = \sum_{x,y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

$I(x, y)$



$E(u, v)$



CORNER DETECTION: MATHEMATICS

- ▶ Change in appearance of window $w(x,y)$ for the shift $[u,v]$:

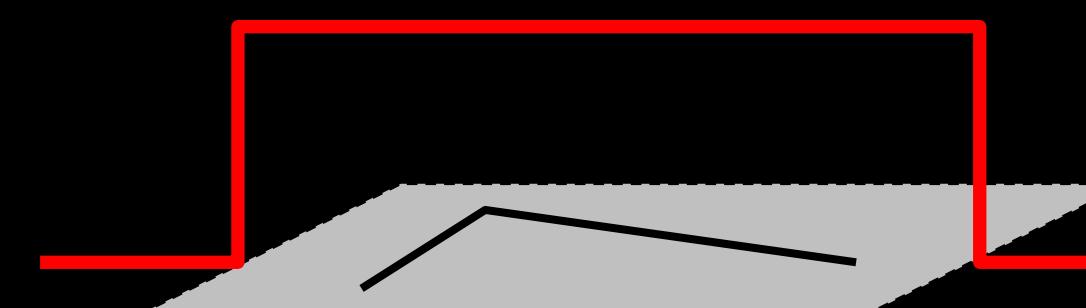
$$E(u, v) = \sum_{x,y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

Window
function

Shifted
intensity

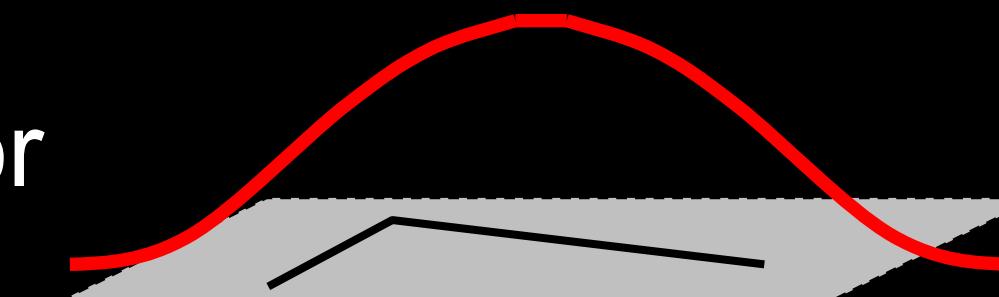
Intensity

Window function $w(x,y) =$



1 in window, 0 outside

or



Gaussian

CORNER DETECTION: MATHEMATICS

- ▶ Change in appearance of window $w(x,y)$ for the shift $[u,v]$:

$$E(u, v) = \sum_{x,y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

We want to find out how this function behaves for small shifts

But this is very slow to compute naively.

$O(\text{window_width}^2 * \text{shift_range}^2 * \text{image_width}^2)$

$O(112 * 112 * 600^2) = 5.2$ billion of these
14.6 thousand per pixel in your image

CORNER DETECTION: MATHEMATICS

- ▶ Change in appearance of window $w(x,y)$ for the shift $[u,v]$:

$$E(u, v) = \sum_{x,y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

We want to find out how this function behaves for small shifts

Recall Taylor series expansion. A function f can be approximated at point a as

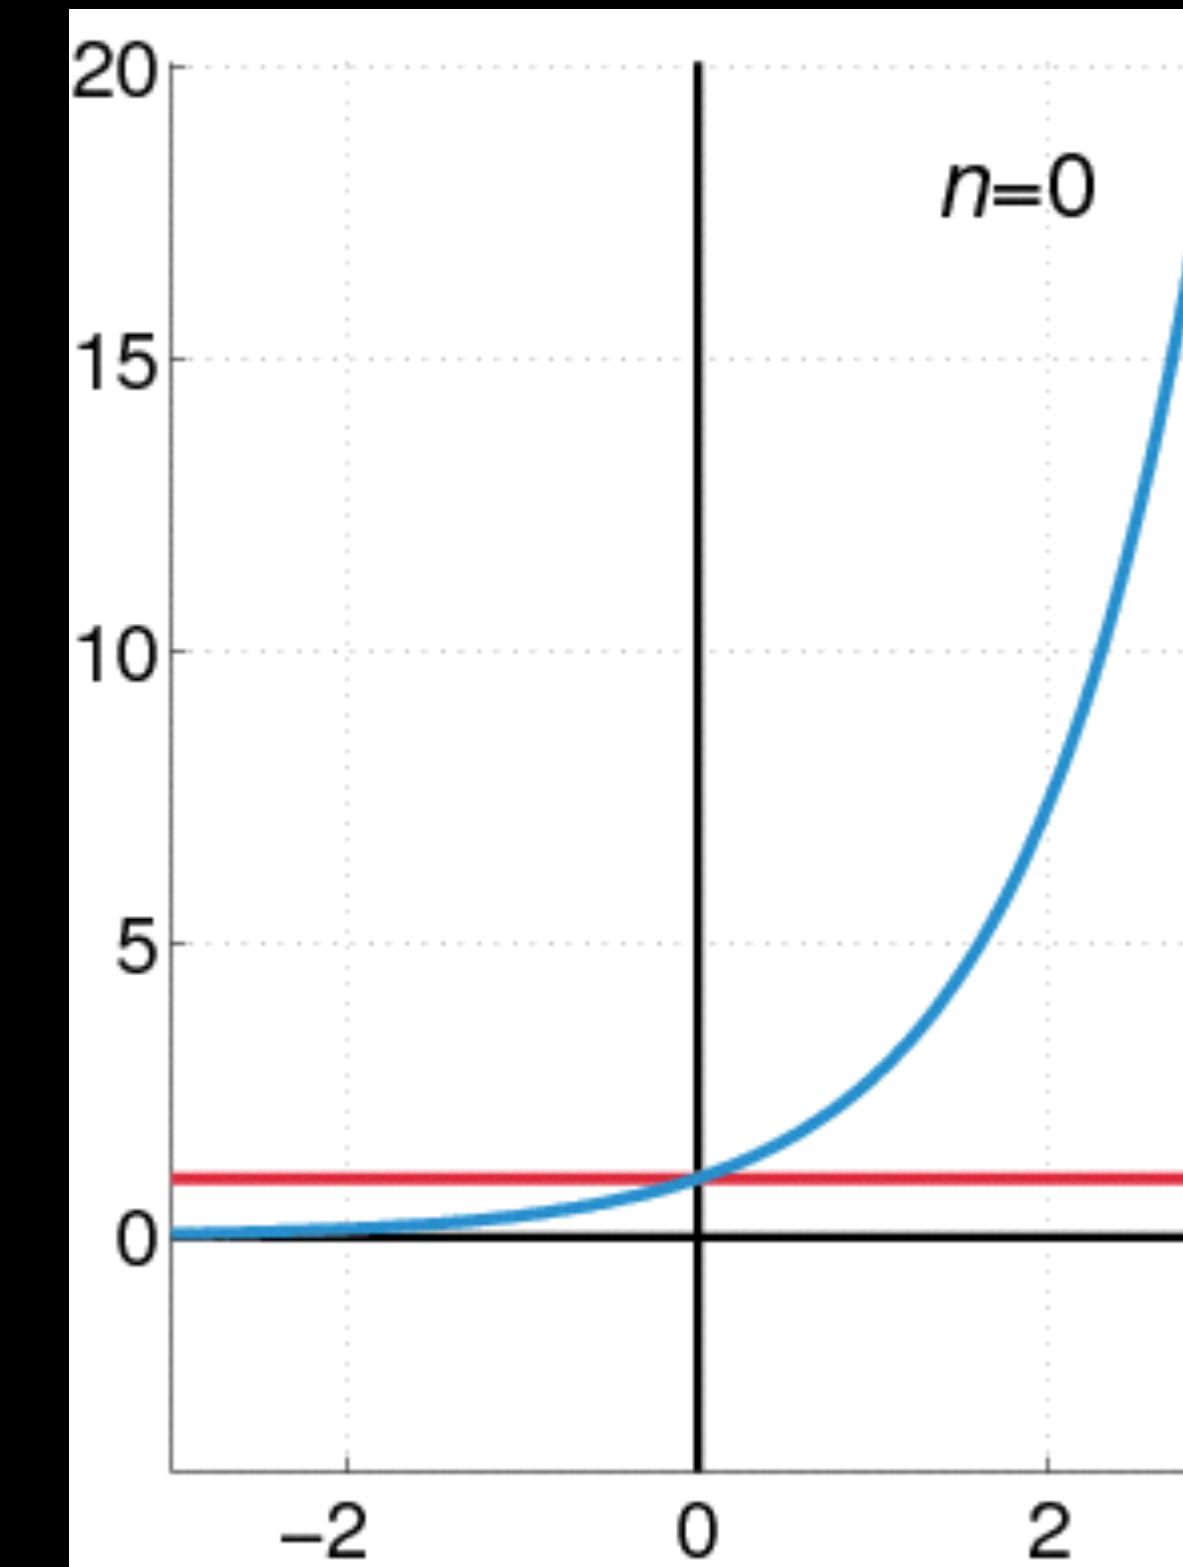
$$f(a) + \frac{f'(a)}{1!}(x - a) + \frac{f''(a)}{2!}(x - a)^2 + \frac{f'''(a)}{3!}(x - a)^3 + \dots$$

RECALL: TAYLOR SERIES EXPANSION

A FUNCTION F CAN BE APPROXIMATED AS

$$f(a) + \frac{f'(a)}{1!}(x - a) + \frac{f''(a)}{2!}(x - a)^2 + \frac{f'''(a)}{3!}(x - a)^3 + \dots$$

APPROXIMATION OF
 $F(X) = E^X$
CENTERED AT $F(0)$



HARRIS CORNER MATH

Taylor Series expansion of I:

$$I(x+u, y+v) = I(x, y) + \frac{\partial I}{\partial x}u + \frac{\partial I}{\partial y}v + \text{higher order terms}$$

If the motion (u, v) is small, then first order approx is good

$$I(x + u, y + v) \approx I(x, y) + \frac{\partial I}{\partial x}u + \frac{\partial I}{\partial y}v$$

$$\approx I(x, y) + [I_x \ I_y] \begin{bmatrix} u \\ v \end{bmatrix}$$

shorthand: $I_x = \frac{\partial I}{\partial x}$

HARRIS CORNER MATH

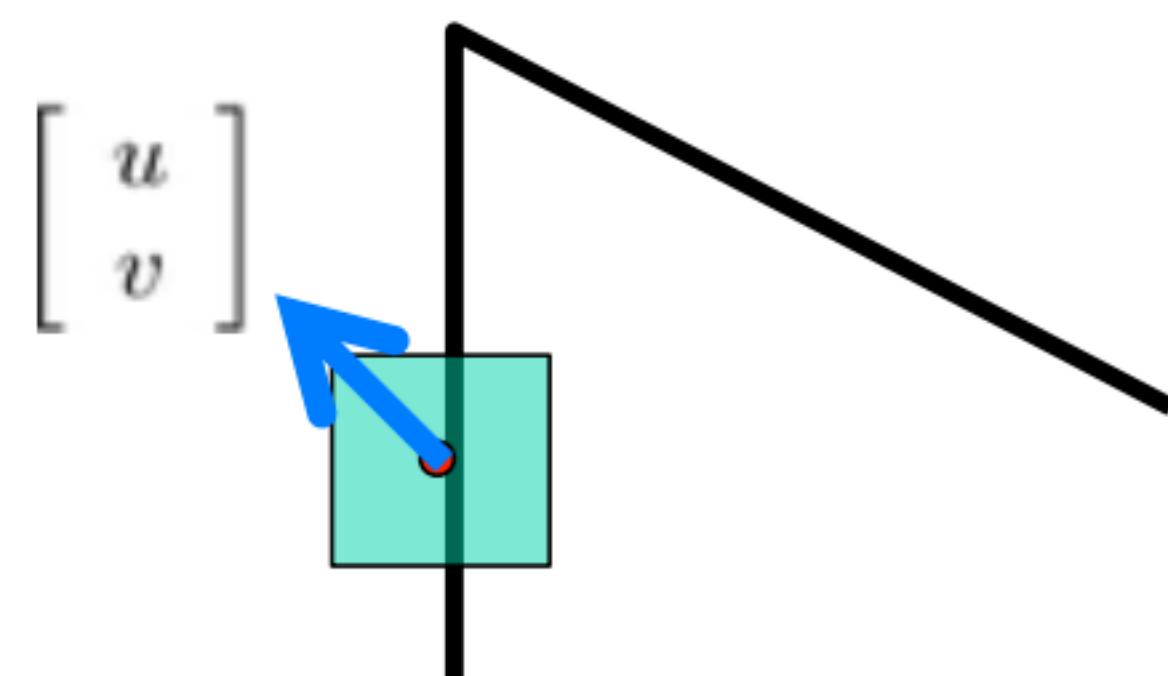
$$\begin{aligned} E(u, v) &= \sum_{(x,y) \in W} [I(x+u, y+v) - I(x, y)]^2 \\ &\approx \sum_{(x,y) \in W} [I(x, y) + [I_x \ I_y] \begin{bmatrix} u \\ v \end{bmatrix} - I(x, y)]^2 \\ &\approx \sum_{(x,y) \in W} \left[[I_x \ I_y] \begin{bmatrix} u \\ v \end{bmatrix} \right]^2 \end{aligned}$$

HARRIS CORNER MATH

This can be rewritten:

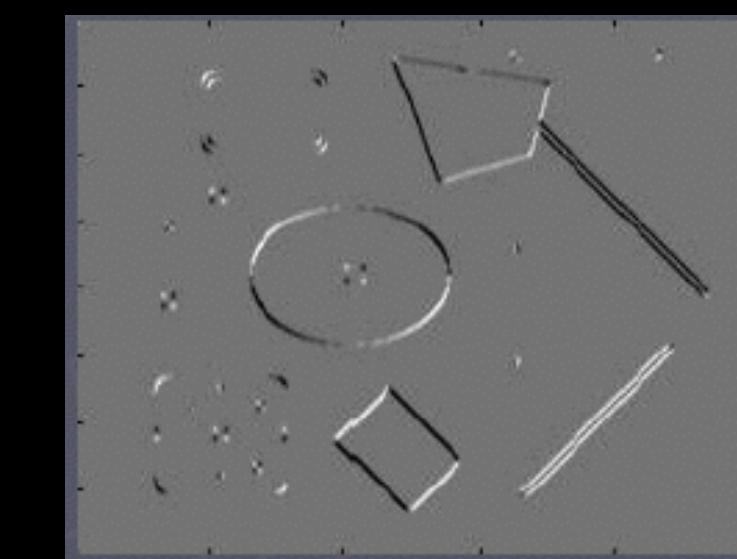
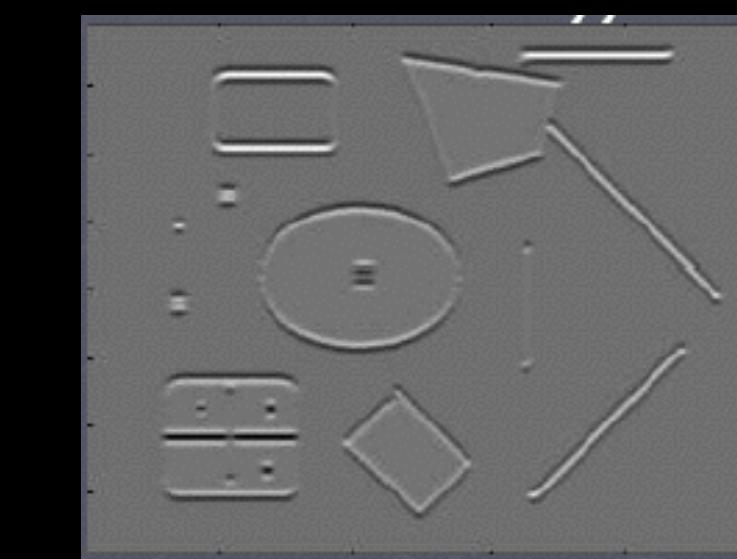
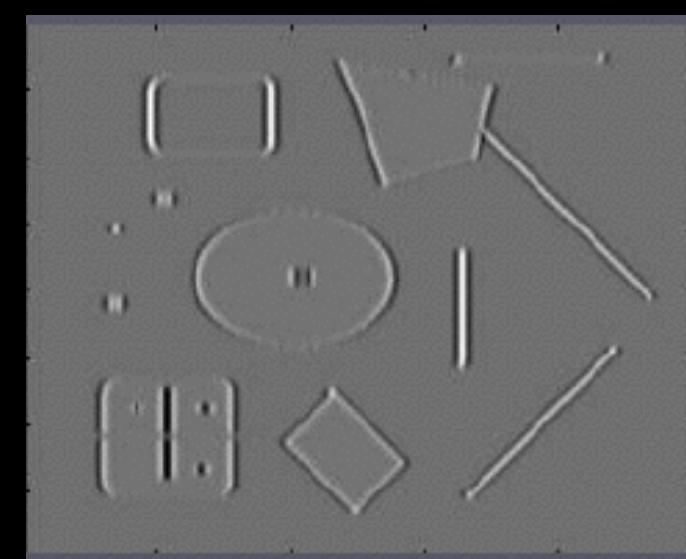
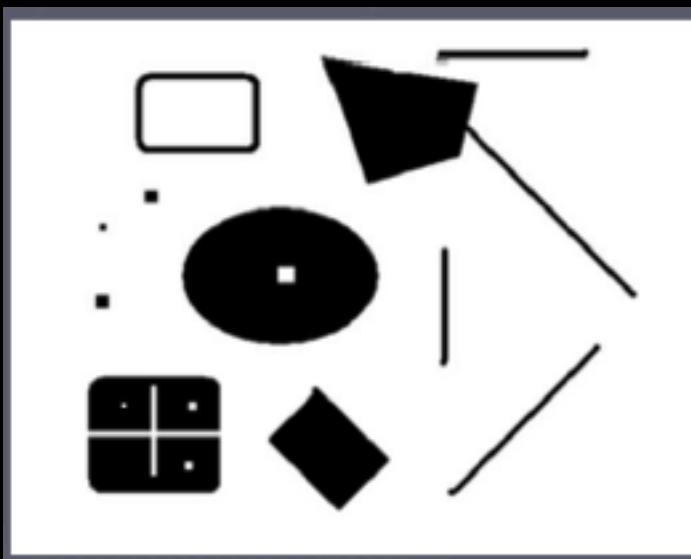
$$E(u, v) = [u \ v] \left(\sum_{(x,y) \in W} \begin{bmatrix} I_x^2 & I_x I_y \\ I_y I_x & I_y^2 \end{bmatrix} \right) \begin{bmatrix} u \\ v \end{bmatrix}$$

H



$$M = \sum w(x, y) \begin{bmatrix} I_x I_x & I_x I_y \\ I_x I_y & I_y I_y \end{bmatrix}$$

2 x 2 matrix of image derivatives (averaged in neighborhood of a point).



NOTATION:

$$I_x \Leftrightarrow \frac{\partial I}{\partial x}$$

$$I_y \Leftrightarrow \frac{\partial I}{\partial y}$$

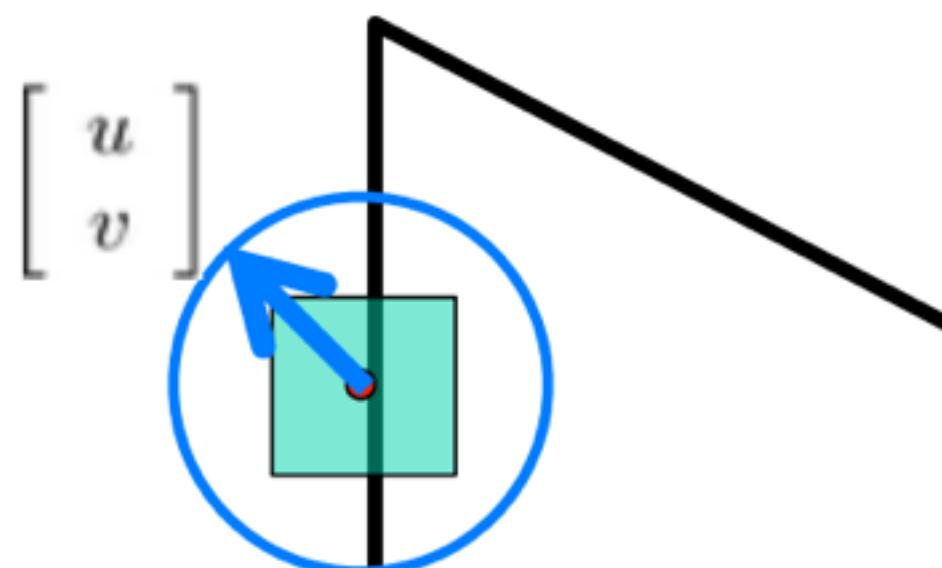
$$I_x I_y \Leftrightarrow \frac{\partial I}{\partial x} \frac{\partial I}{\partial y}$$

HARRIS CORNER MATH

This can be rewritten:

$$E(u, v) = [u \ v] \left(\sum_{(x,y) \in W} \begin{bmatrix} I_x^2 & I_x I_y \\ I_y I_x & I_y^2 \end{bmatrix} \right) \begin{bmatrix} u \\ v \end{bmatrix}$$

H



Which $[u \ v]$ maximizes $E(u,v)$?

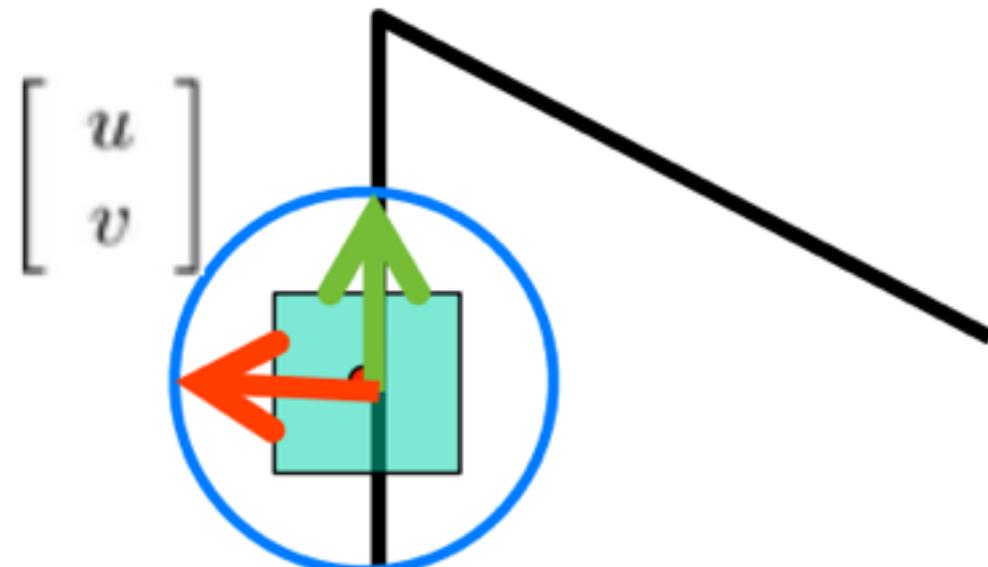
Which $[u \ v]$ minimizes $E(u,v)$?

HARRIS CORNER MATH

This can be rewritten:

$$E(u, v) = [u \ v] \left(\sum_{(x,y) \in W} \begin{bmatrix} I_x^2 & I_x I_y \\ I_y I_x & I_y^2 \end{bmatrix} \right) \begin{bmatrix} u \\ v \end{bmatrix}$$

H



Which $[u \ v]$ maximizes $E(u,v)$?

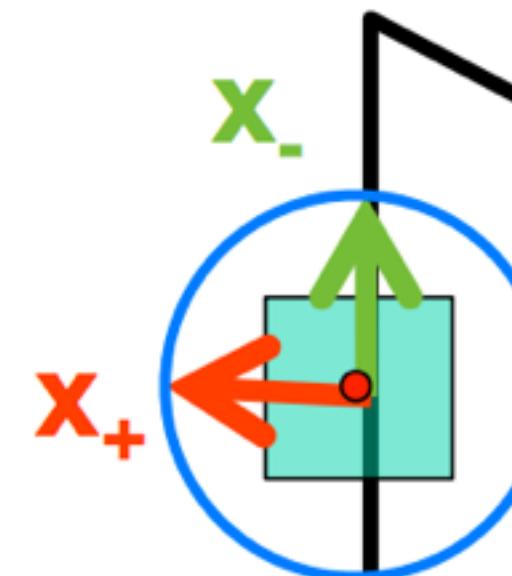
Which $[u \ v]$ minimizes $E(u,v)$?

HARRIS CORNER MATH

This can be rewritten:

$$E(u, v) = [u \ v] \left(\sum_{(x,y) \in W} \begin{bmatrix} I_x^2 & I_x I_y \\ I_y I_x & I_y^2 \end{bmatrix} \right) \begin{bmatrix} u \\ v \end{bmatrix}$$

H



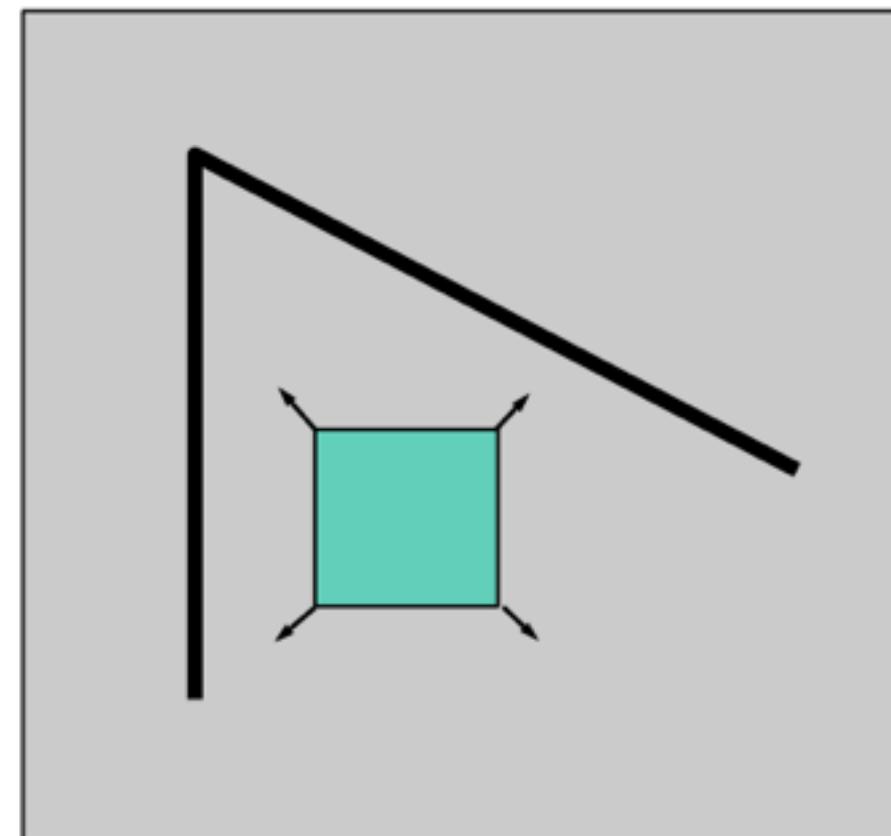
Eigenvector \mathbf{x}_+ with the largest eigen value?

Eigenvector \mathbf{x}_- with the smallest eigen value?

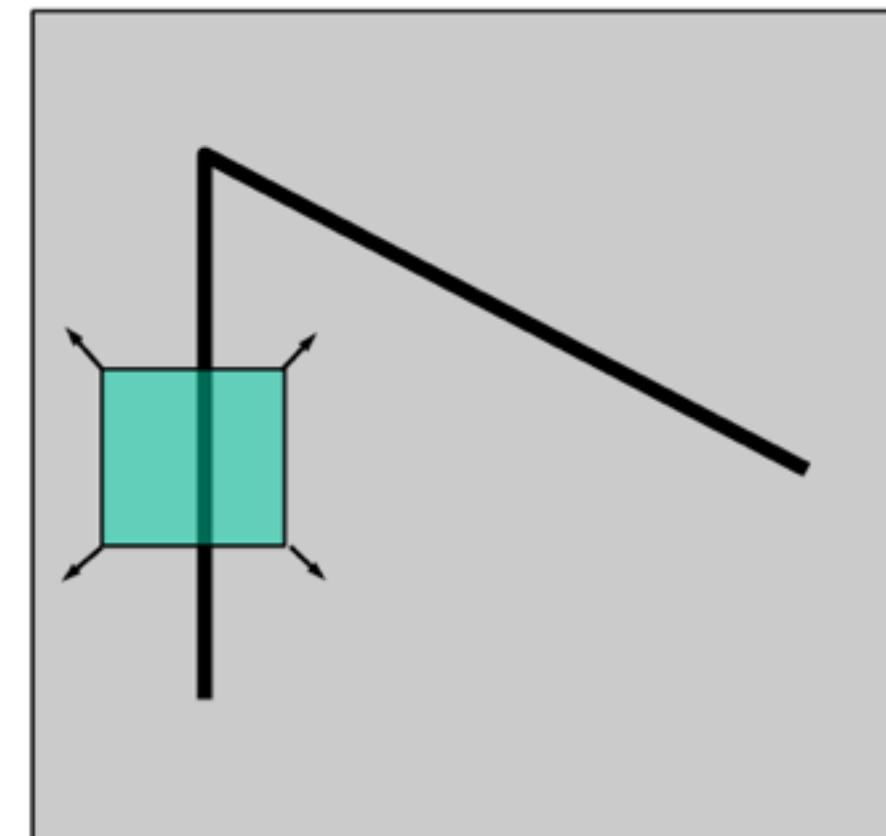
HARRIS CORNER MATH

Local measure of feature uniqueness

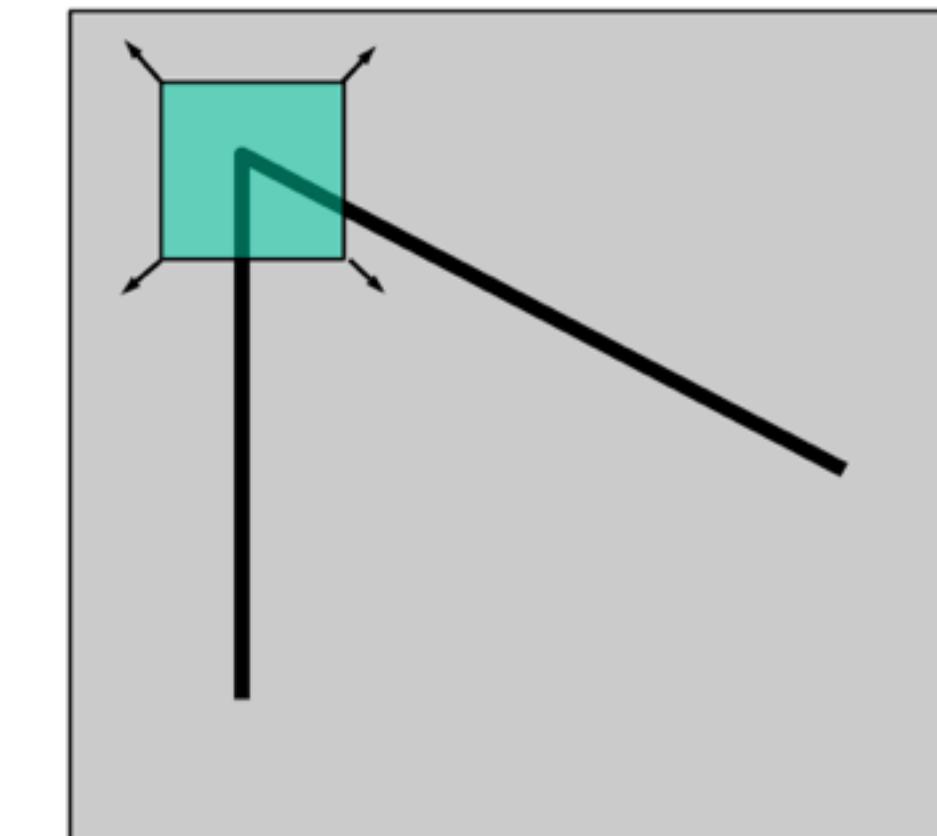
- $E(u,v)$ = amount of change when you shift the window by (u,v)



$E(u,v)$ is small
for **all** shifts



$E(u,v)$ is small
for **some** shifts

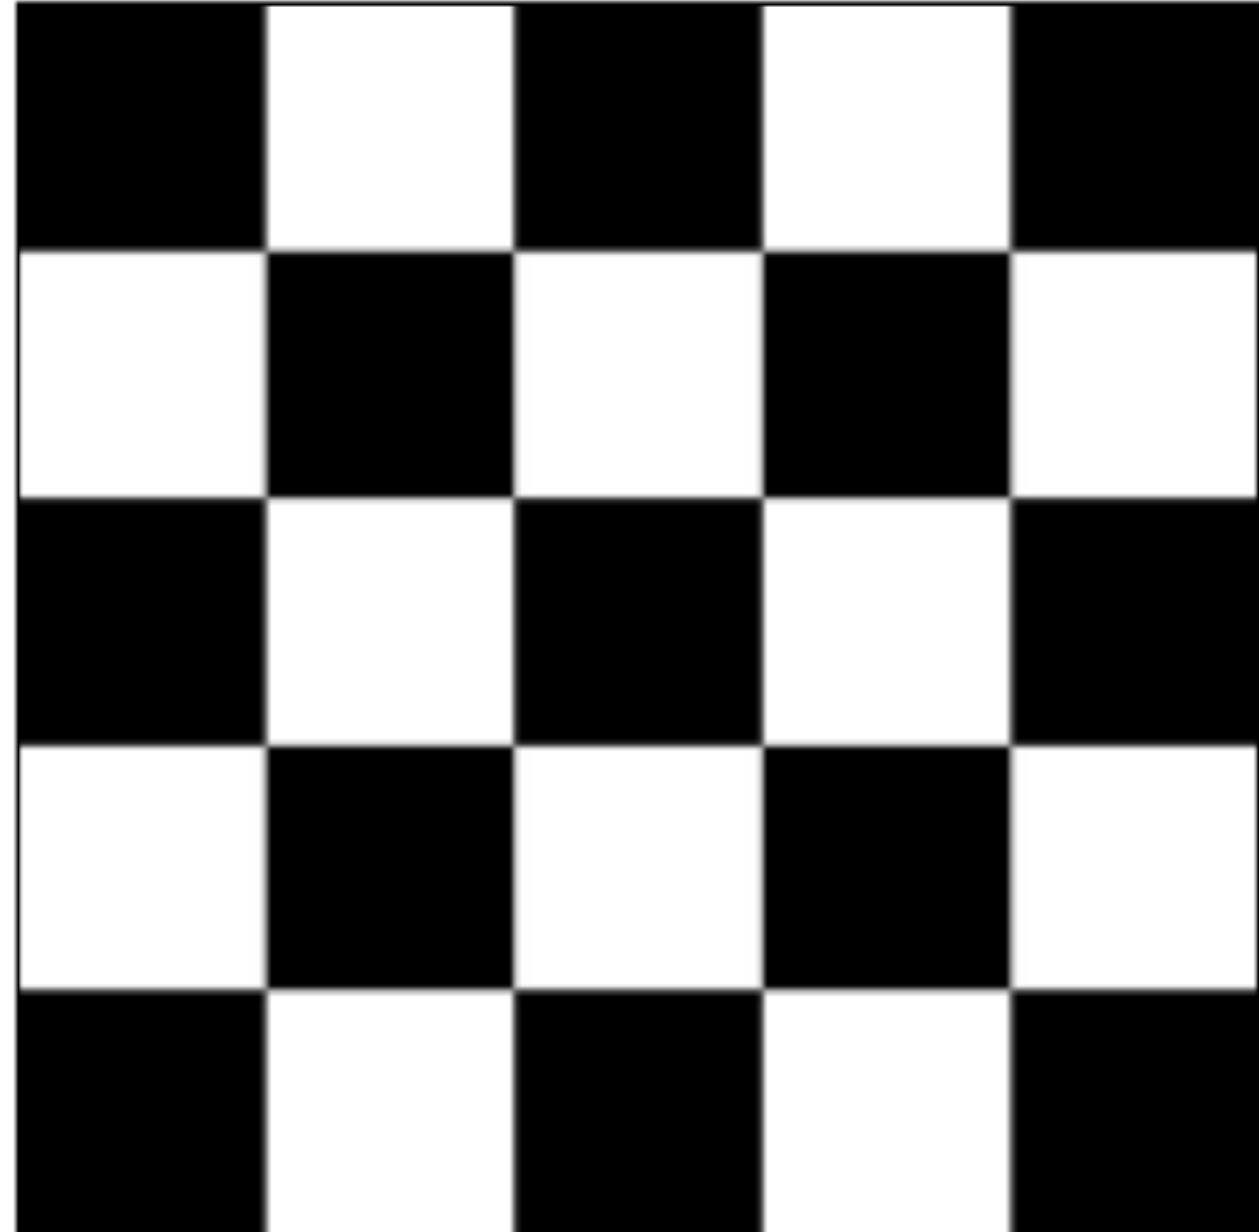


$E(u,v)$ is small
for **no** shifts

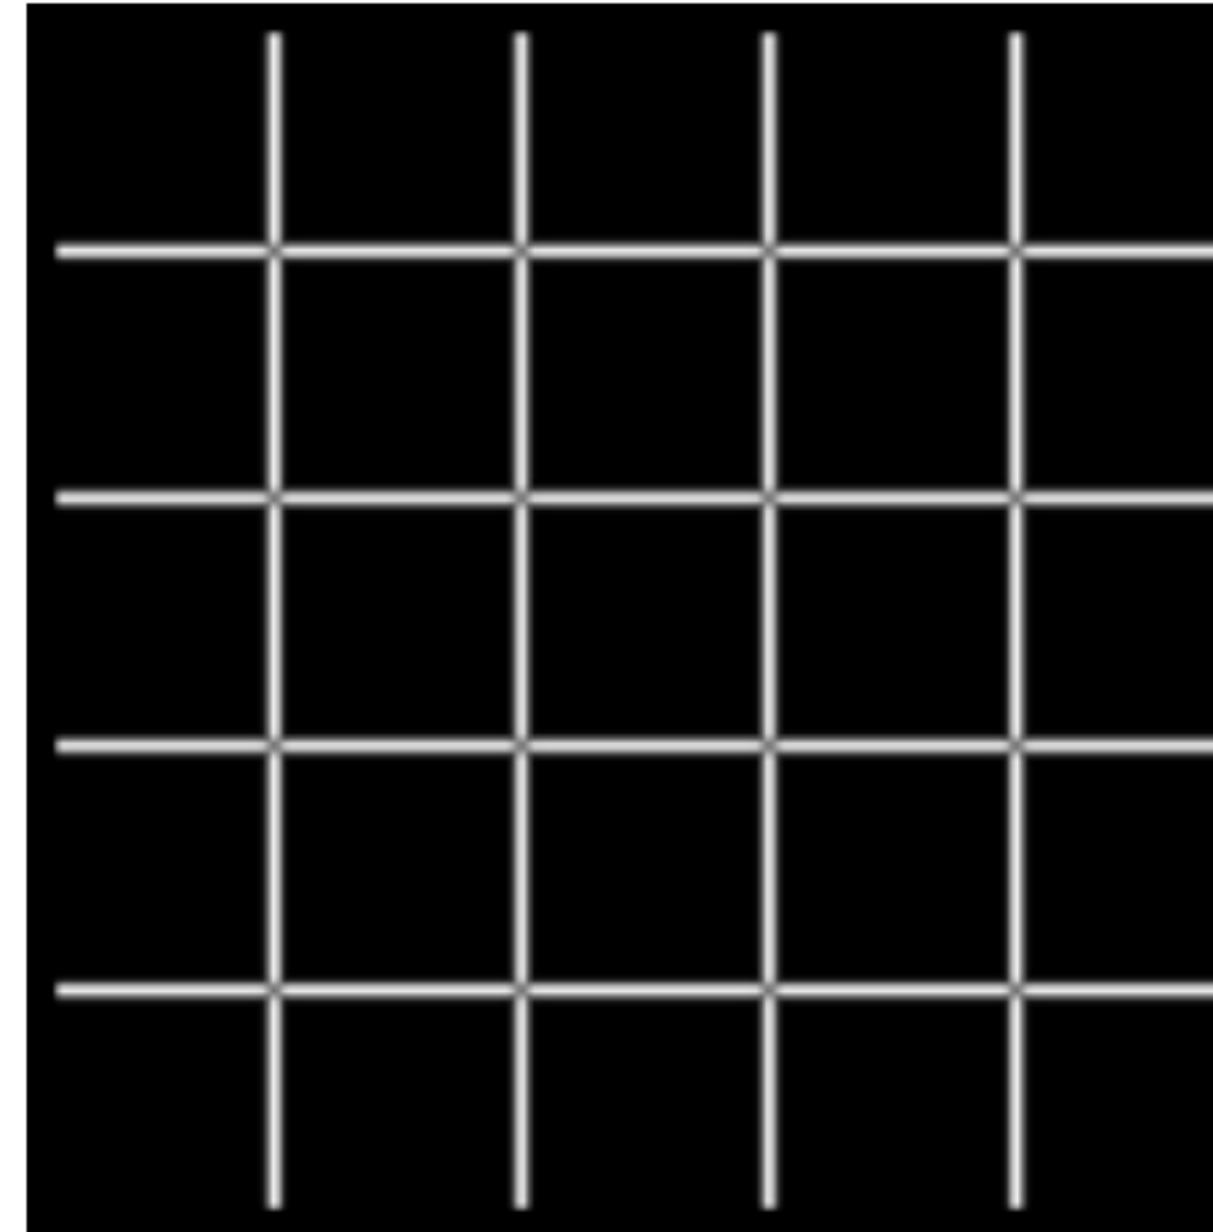
We want $\min_{(u,v)} E(u, v)$ to be large

TEXT

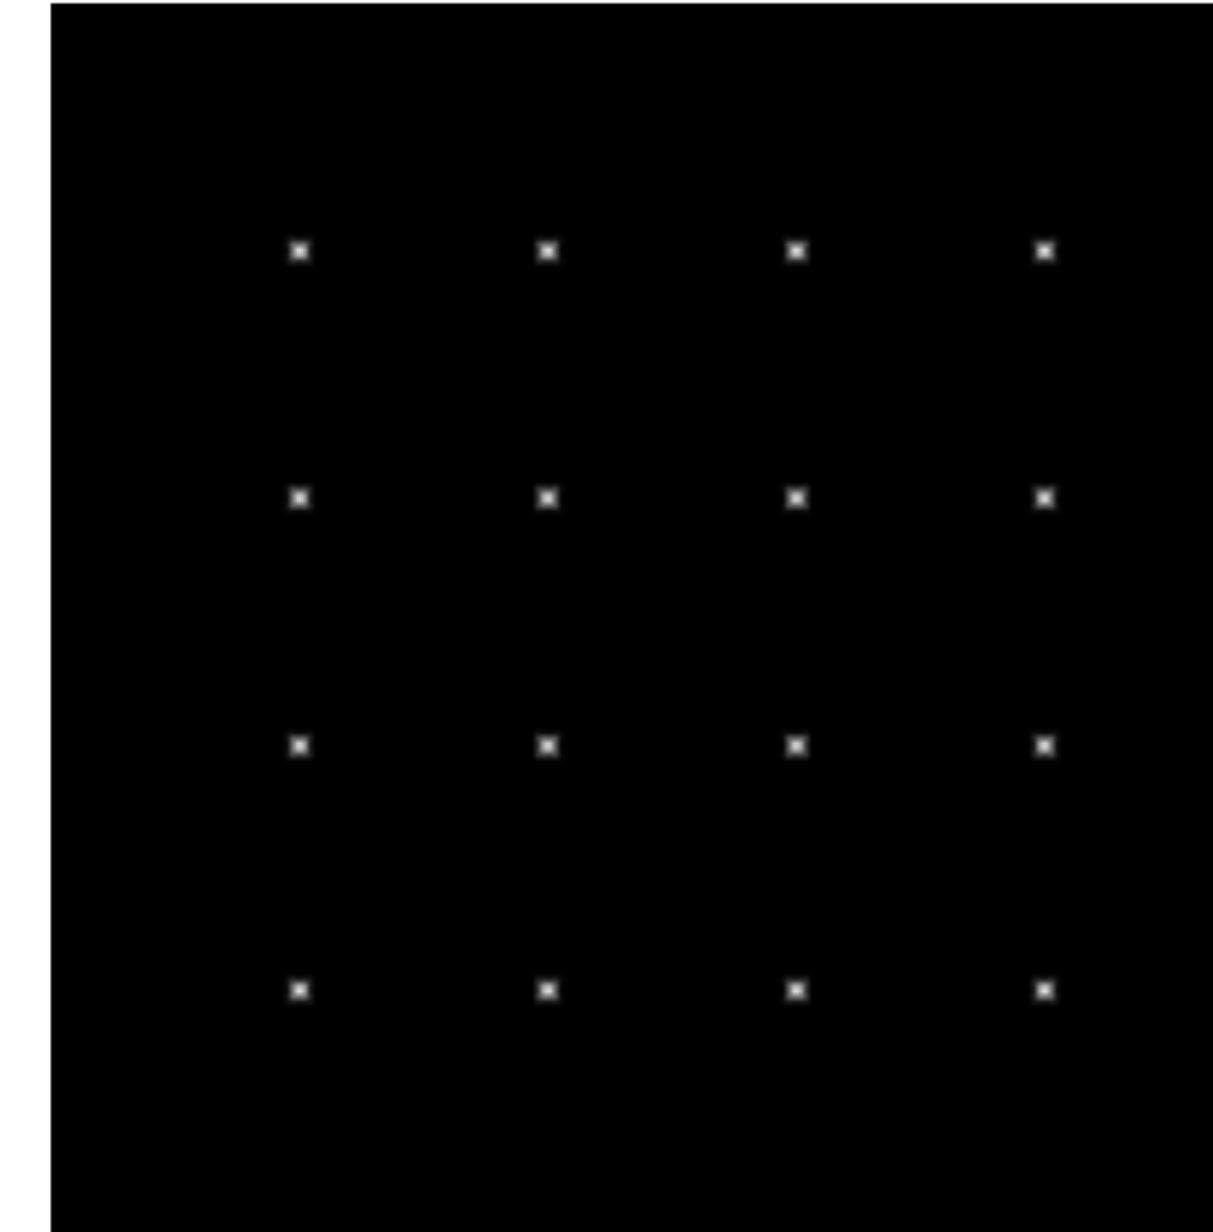
HARRIS CORNER MATH



I



λ_+

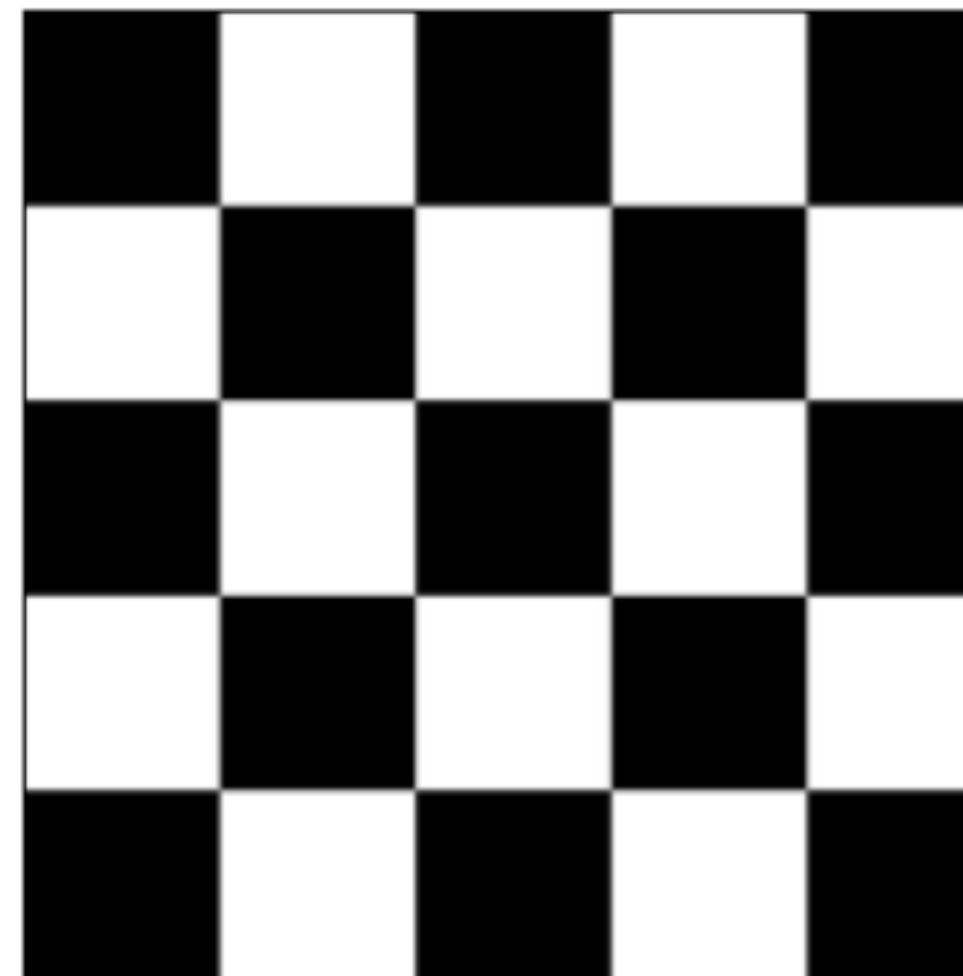


λ_-

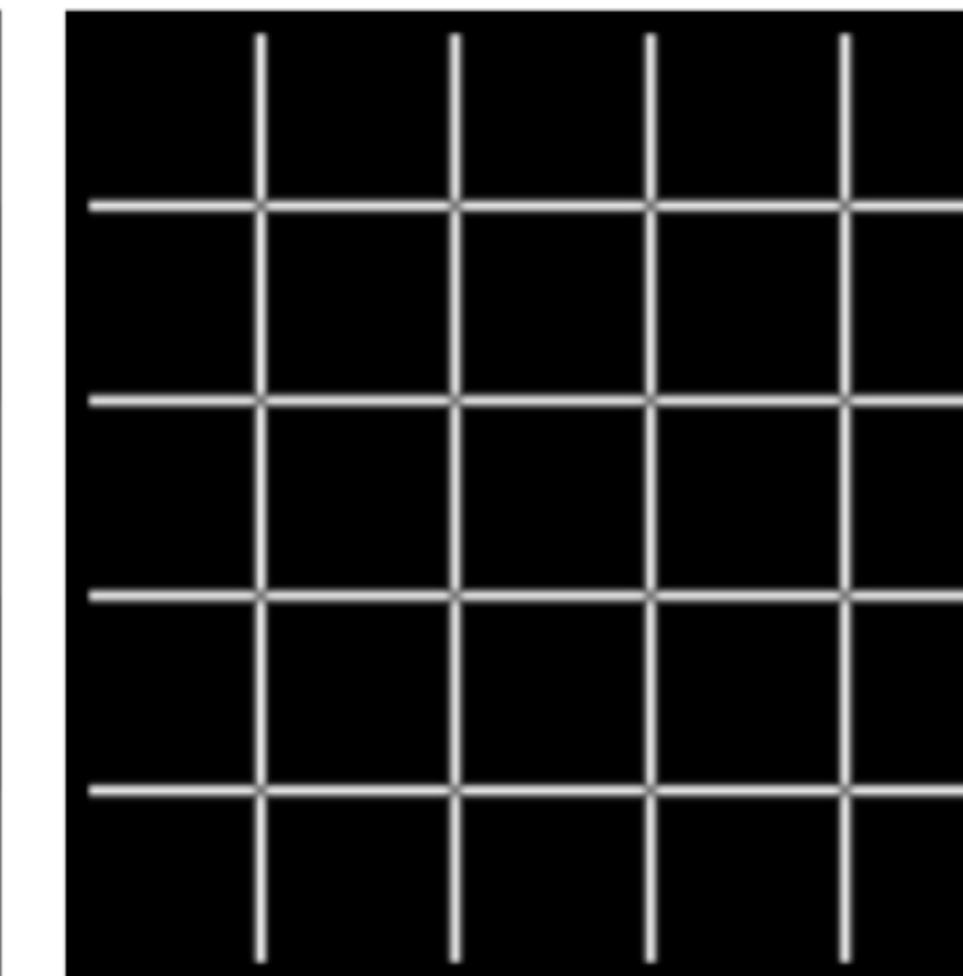
HARRIS CORNER MATH

Here's what you do

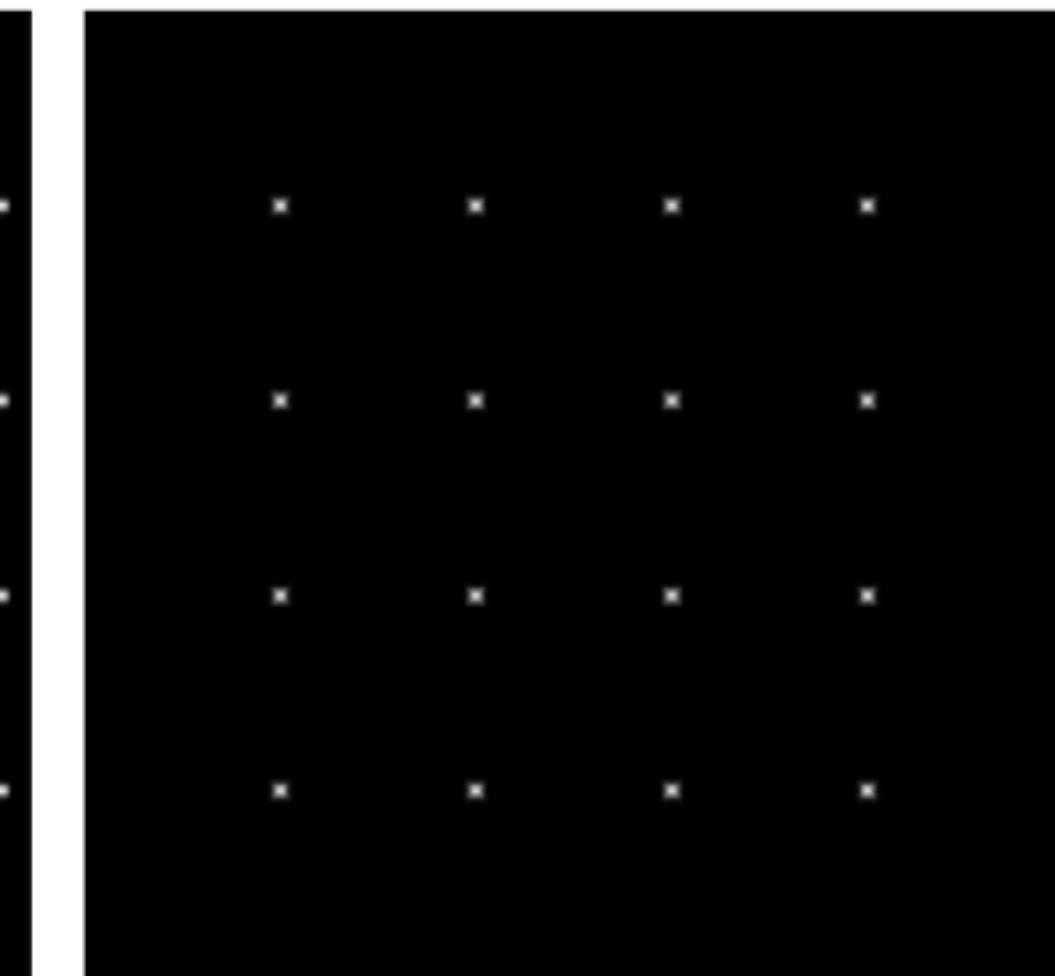
- Compute the gradient at each point in the image
- Create the H matrix from the entries in the gradient
- Compute the eigenvalues.
- Find points with large response ($\lambda_- > \text{threshold}$)
- Choose those points where λ_- is a local maximum as features



I



λ_+



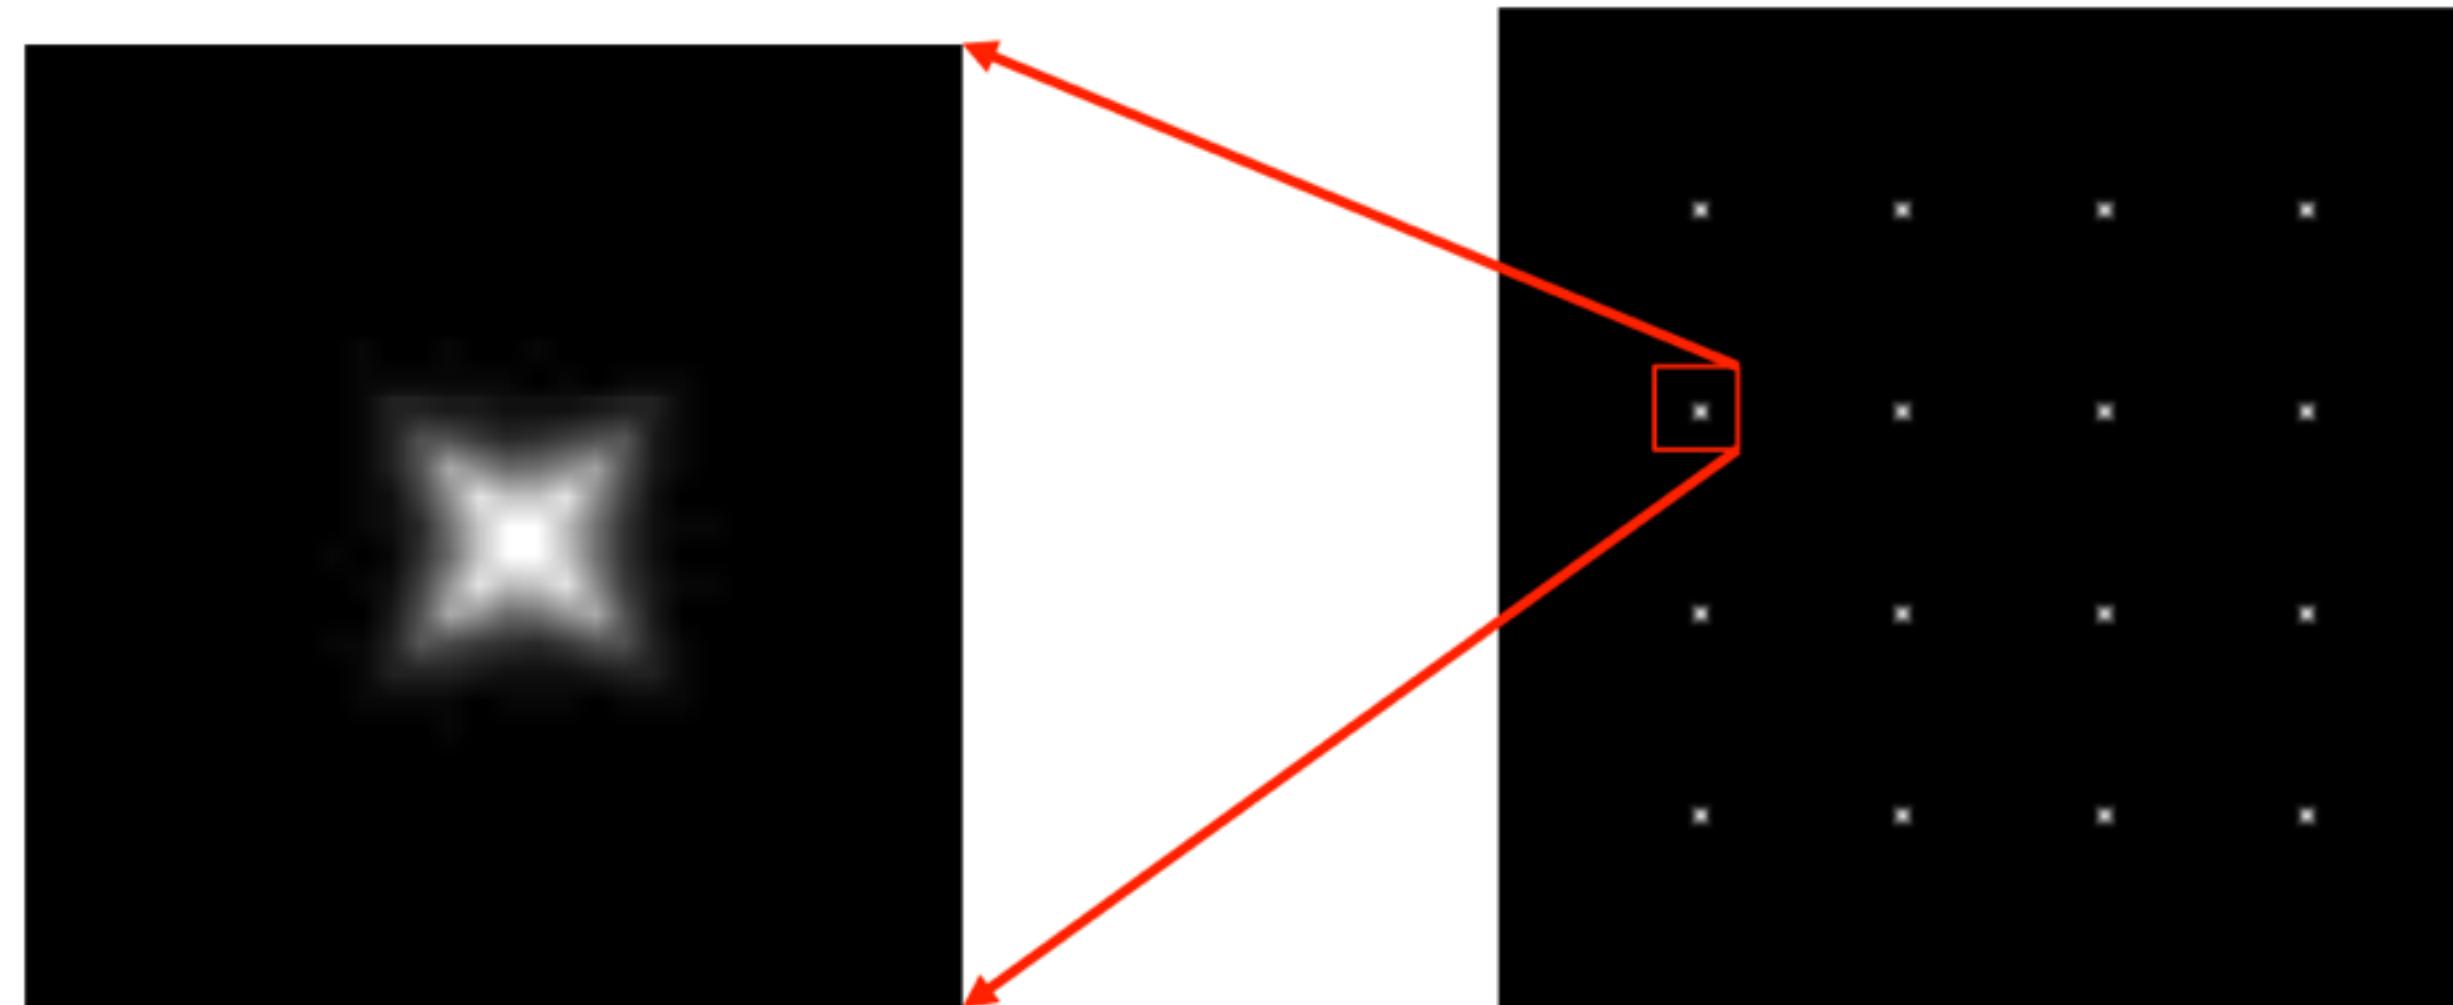
λ_-

HARRIS CORNER MATH

Here's what you do

- Compute the gradient at each point in the image
- Create the H matrix from the entries in the gradient
- Compute the eigenvalues.
- Find points with large response ($\lambda_- > \text{threshold}$)
- Choose those points where λ_- is a local maximum as features

Called “non-local max suppression”



λ_-

INTERPRETING THE SECOND MOMENT MATRIX

Consider a horizontal “slice” of $E(u, v)$:

$$[u \ v] M \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$$

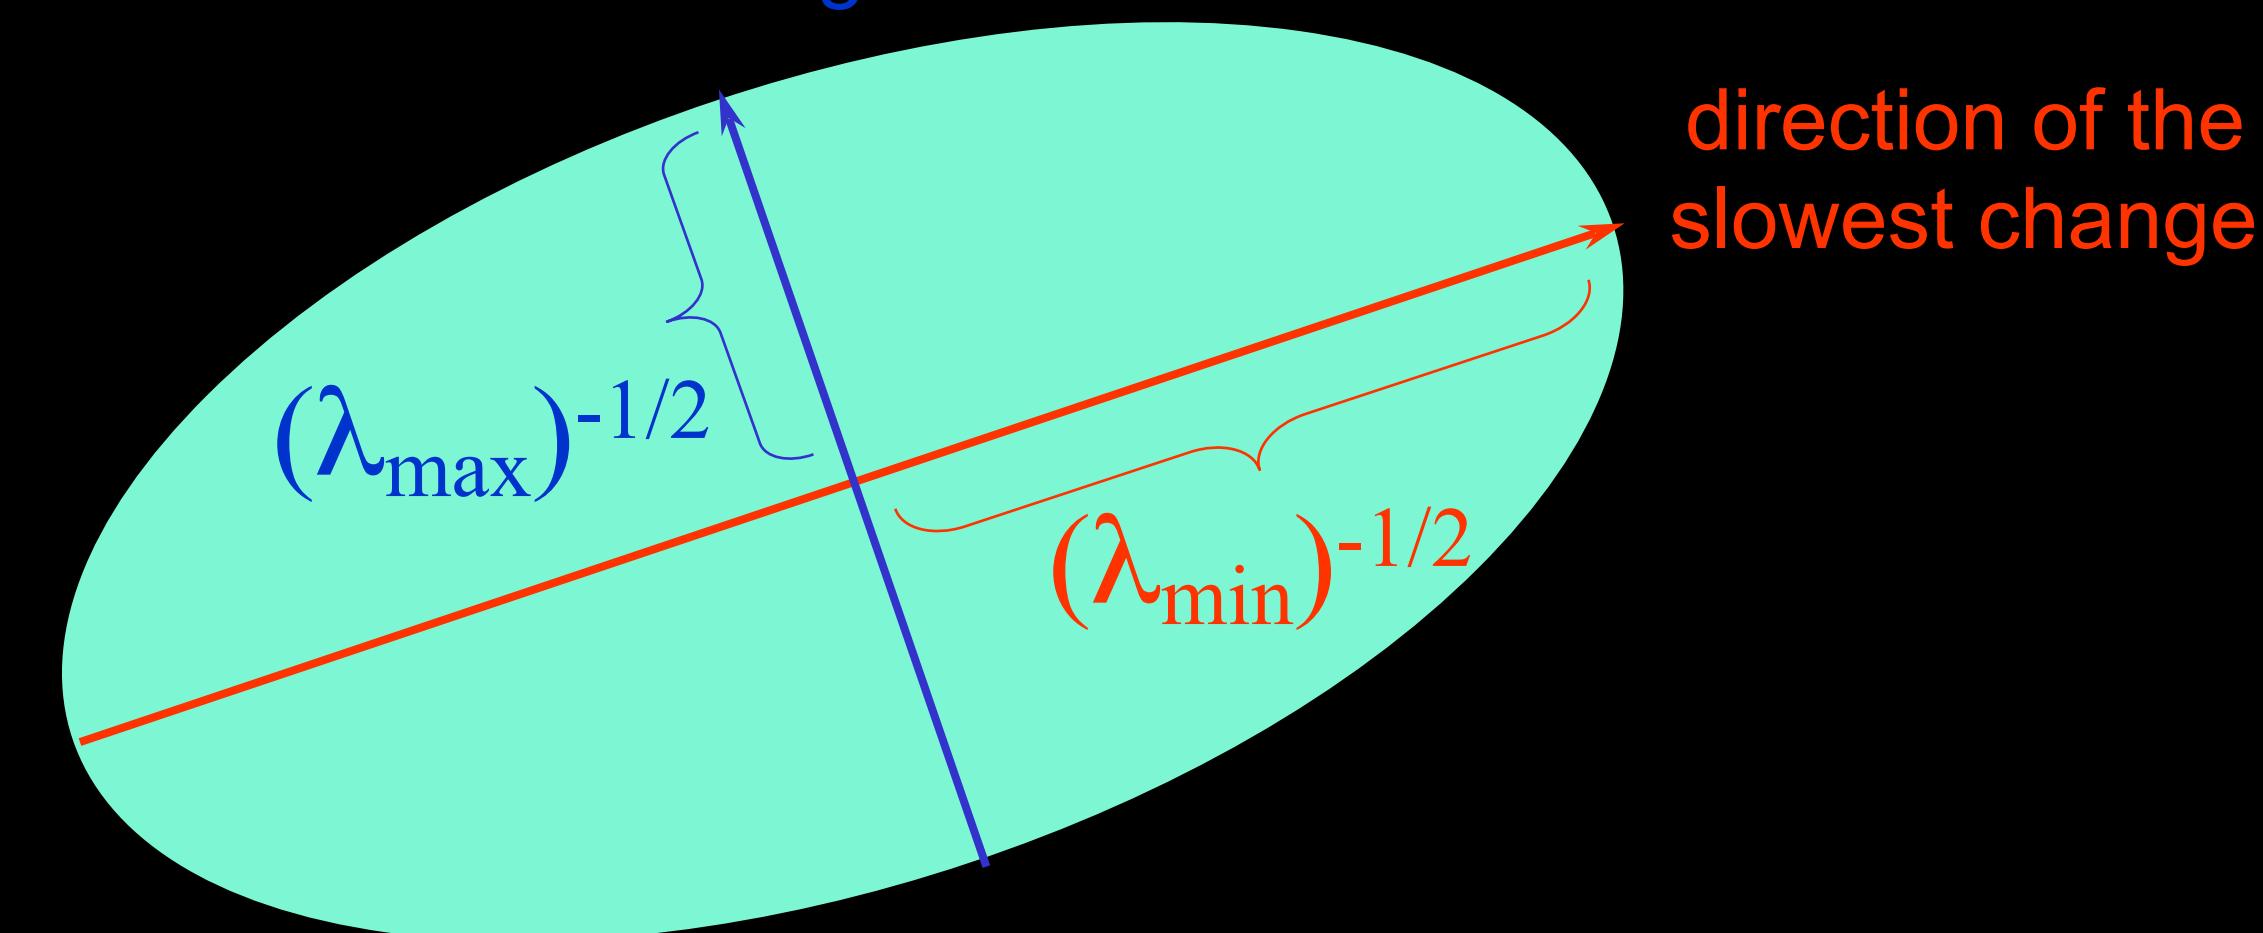
This is the equation of an ellipse.

The axis lengths of the ellipse are determined by the eigenvalues and the orientation is determined by R

direction of the
fastest change

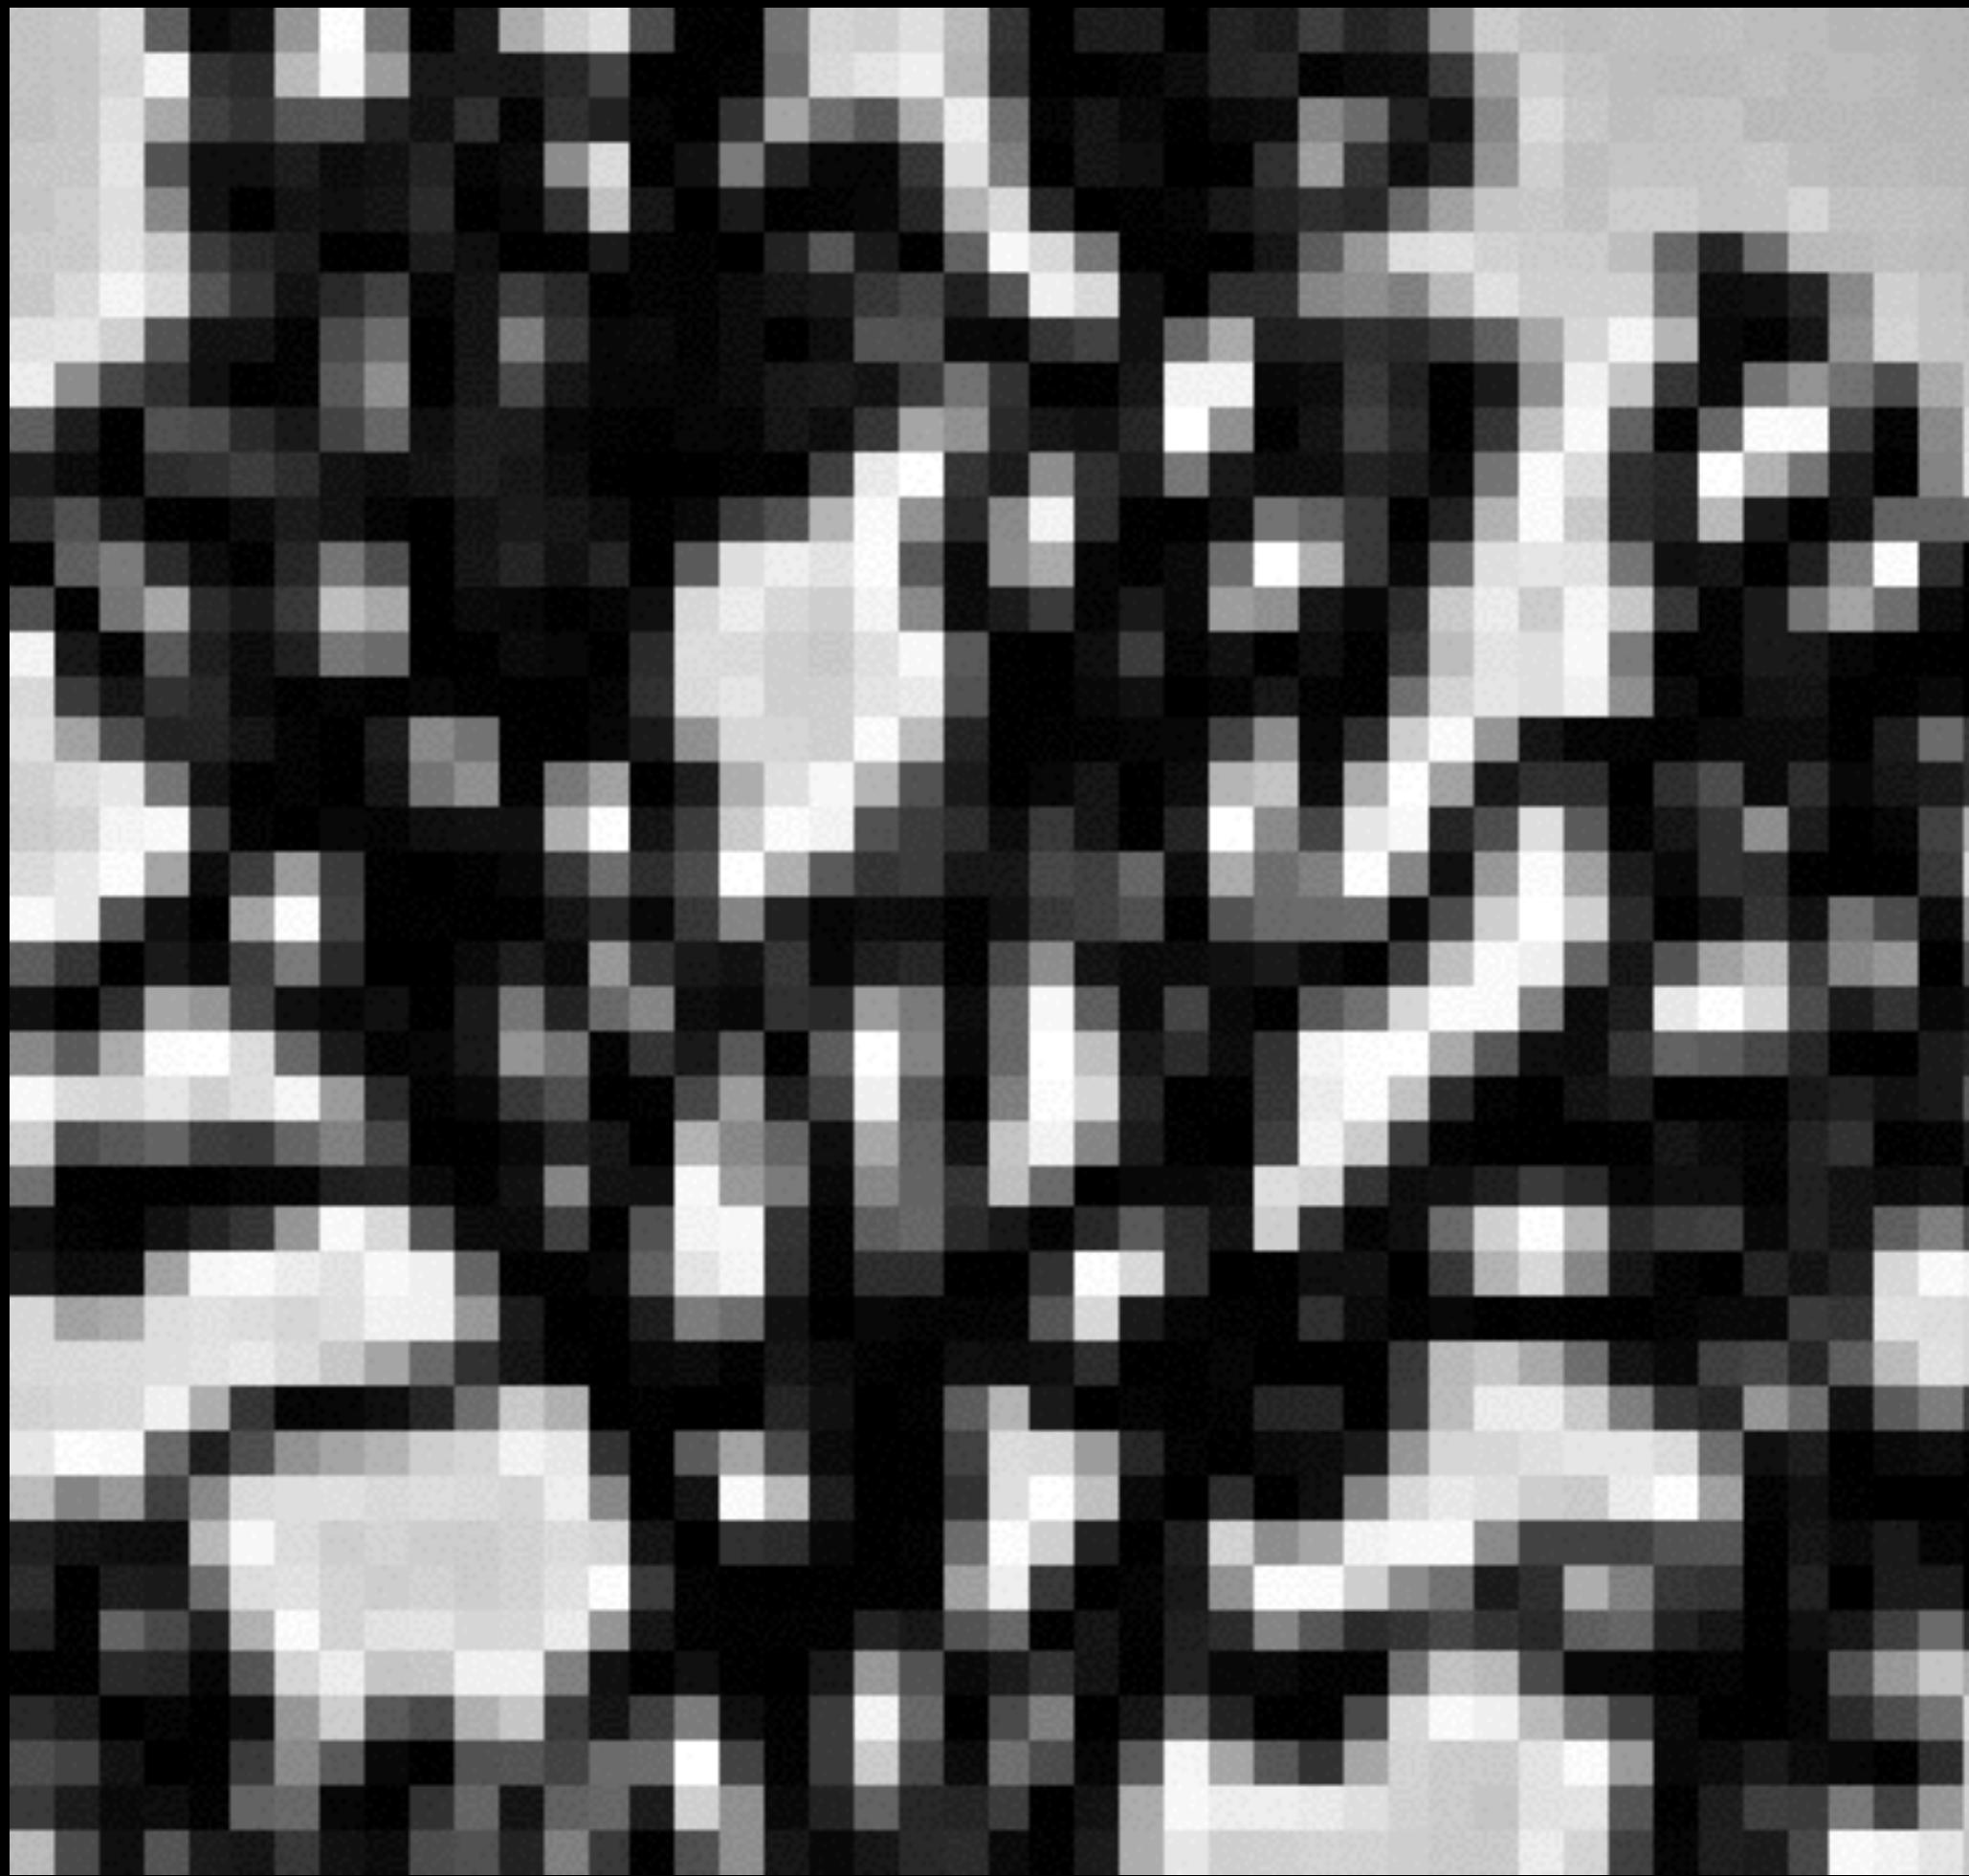
Diagonalization of M :

$$M = R^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} R$$



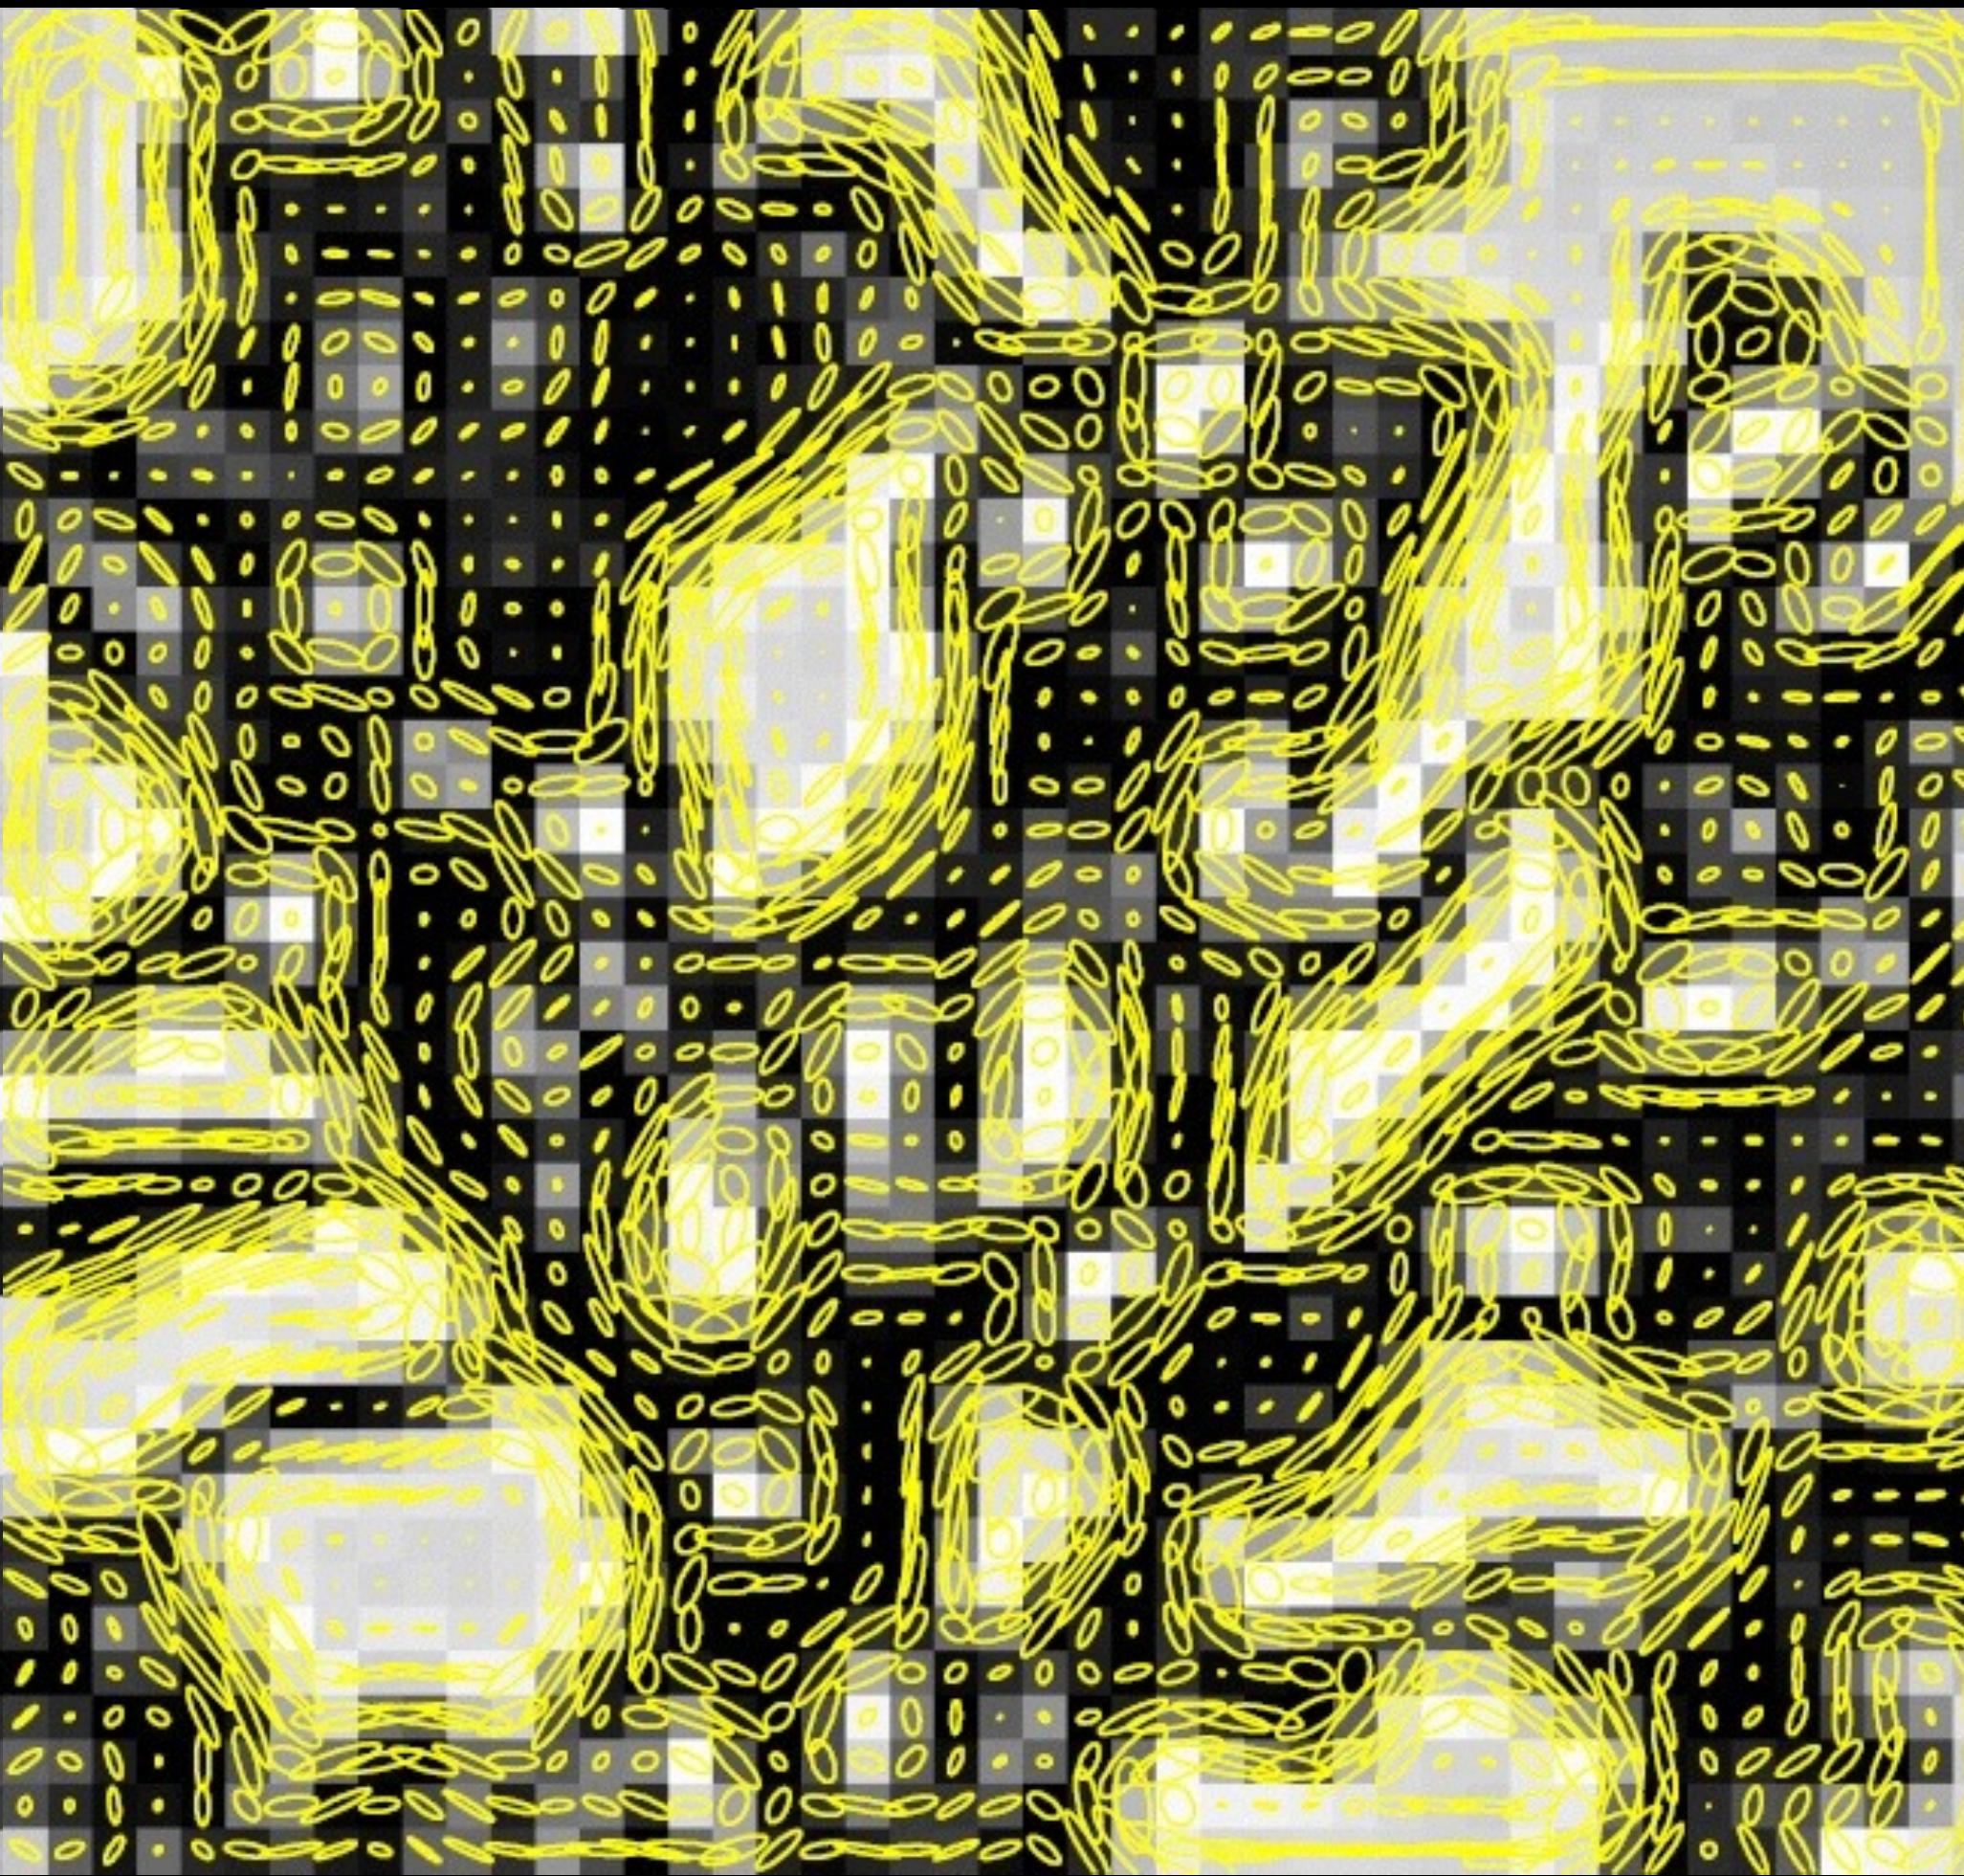
TEXT

VISUALIZATION OF SECOND MOMENT MATRICES



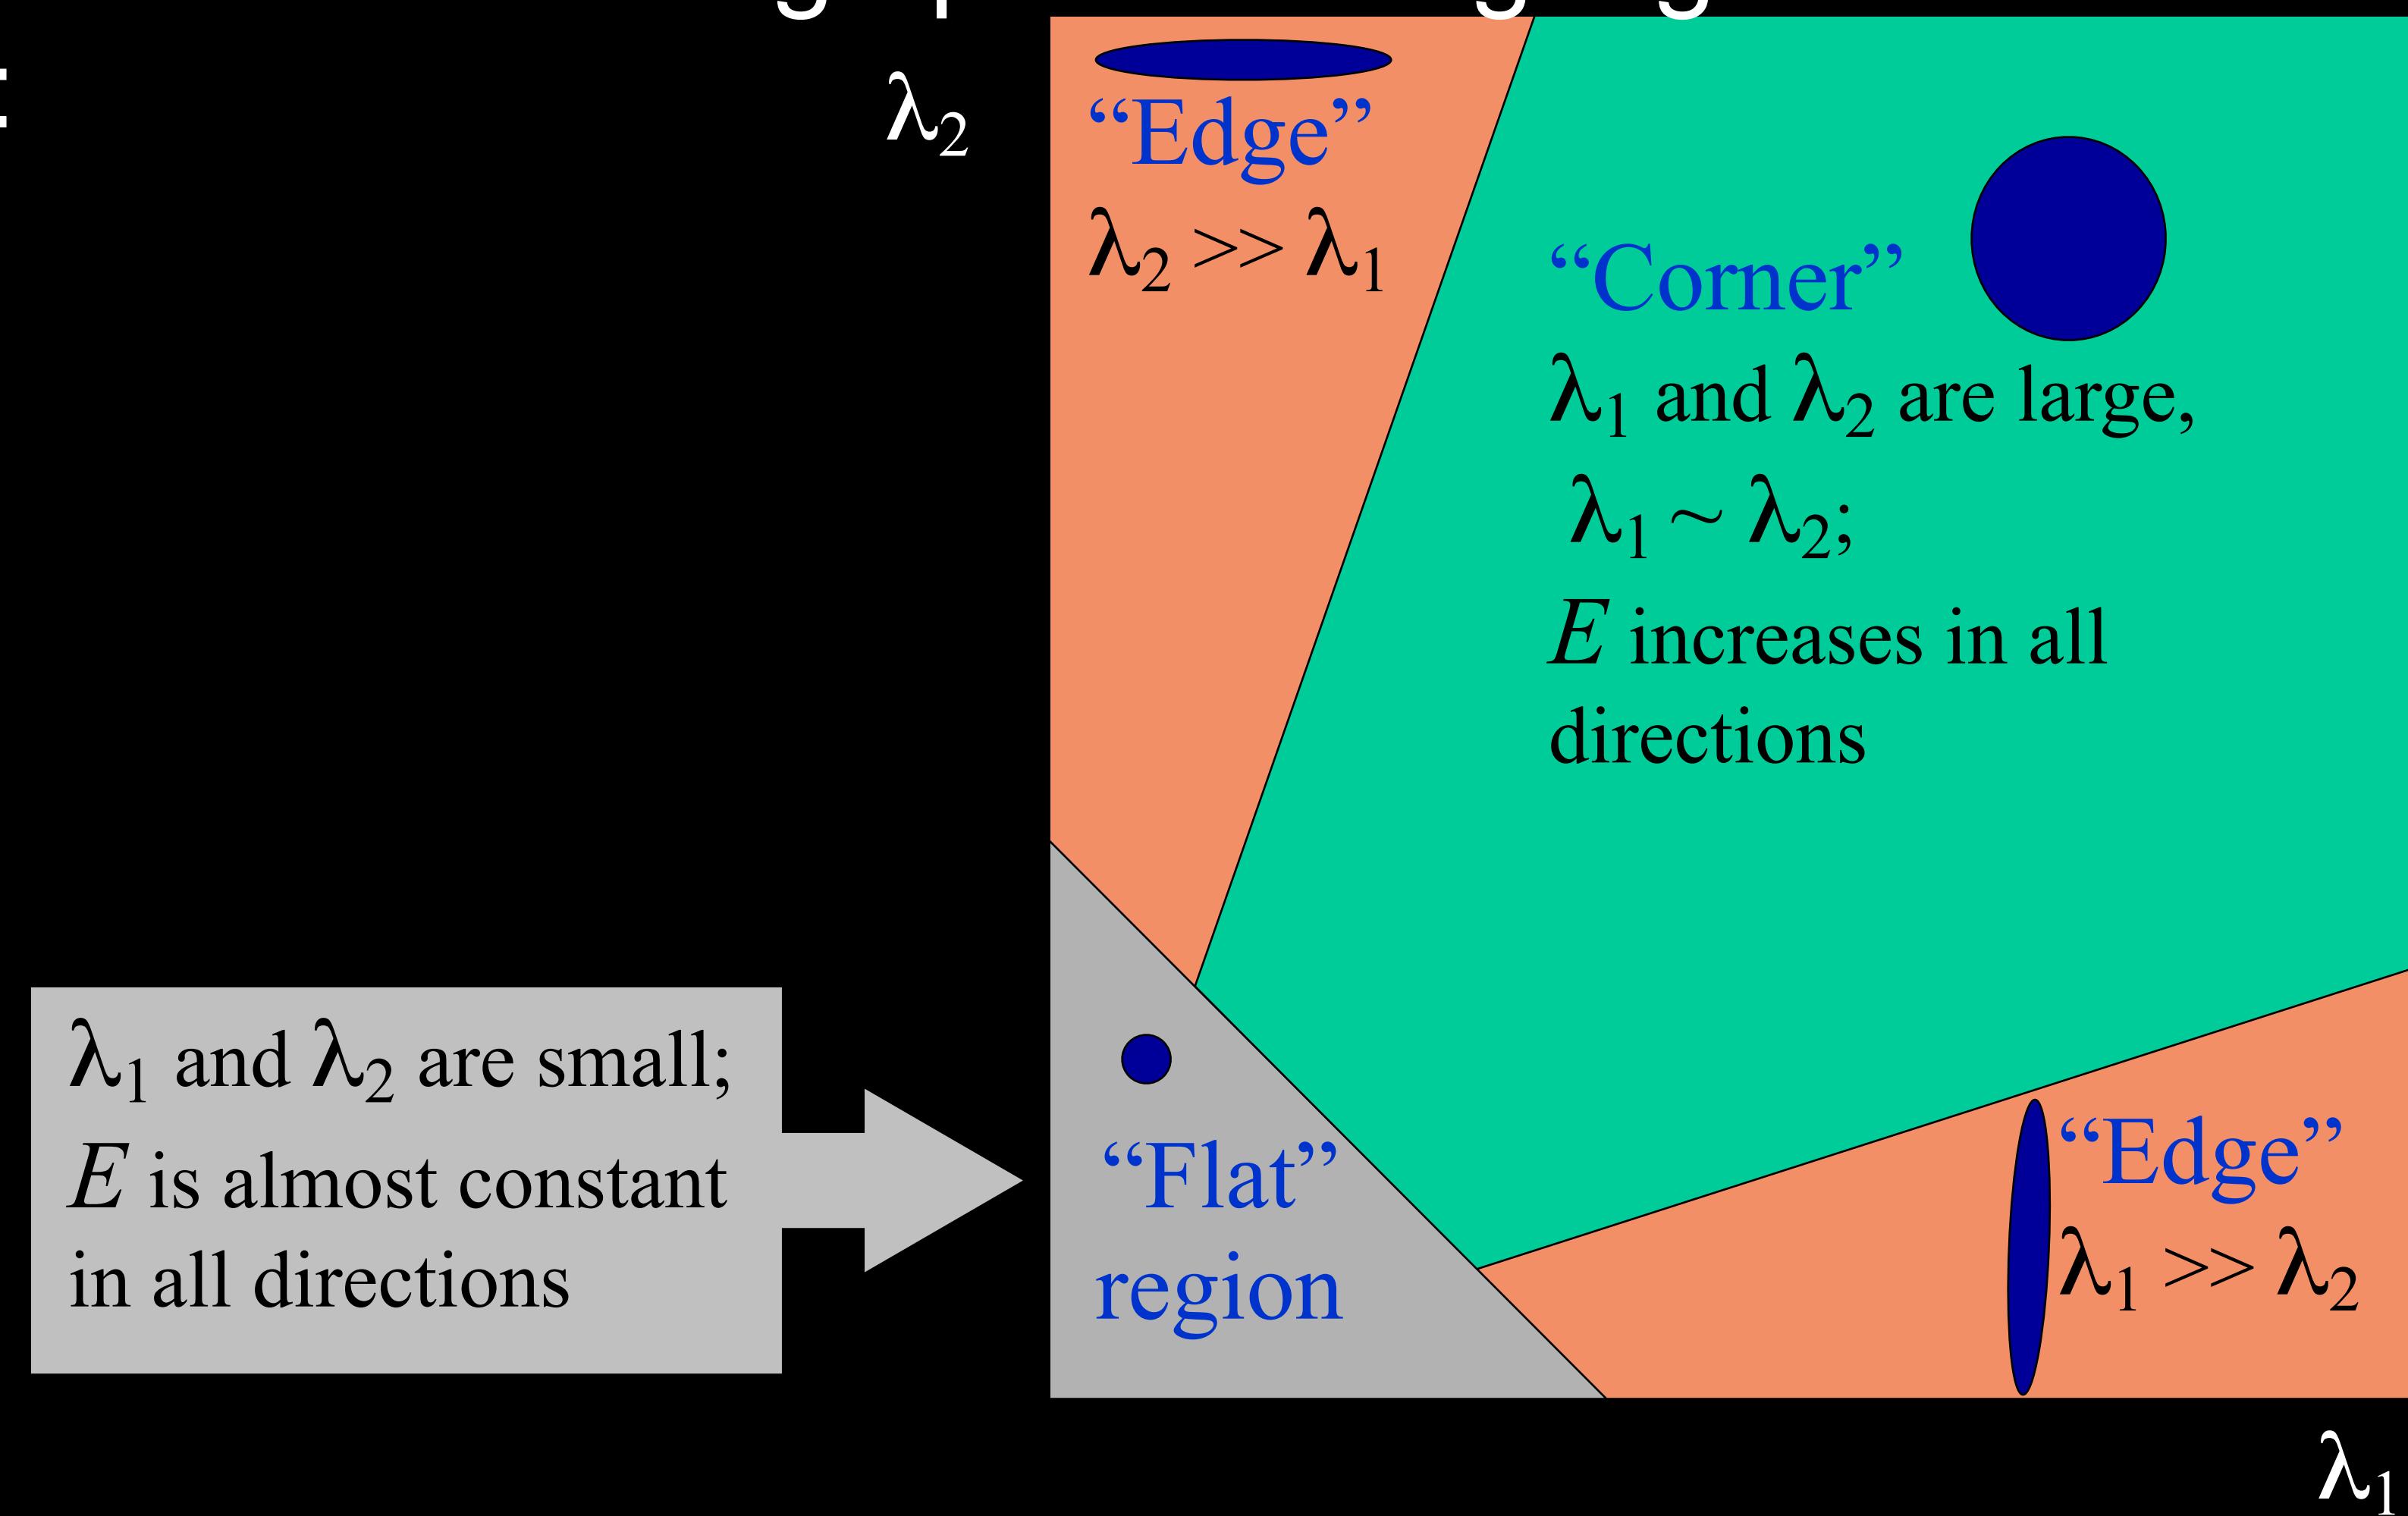
TEXT

VISUALIZATION OF SECOND MOMENT MATRICES



INTERPRETING THE EIGENVALUES

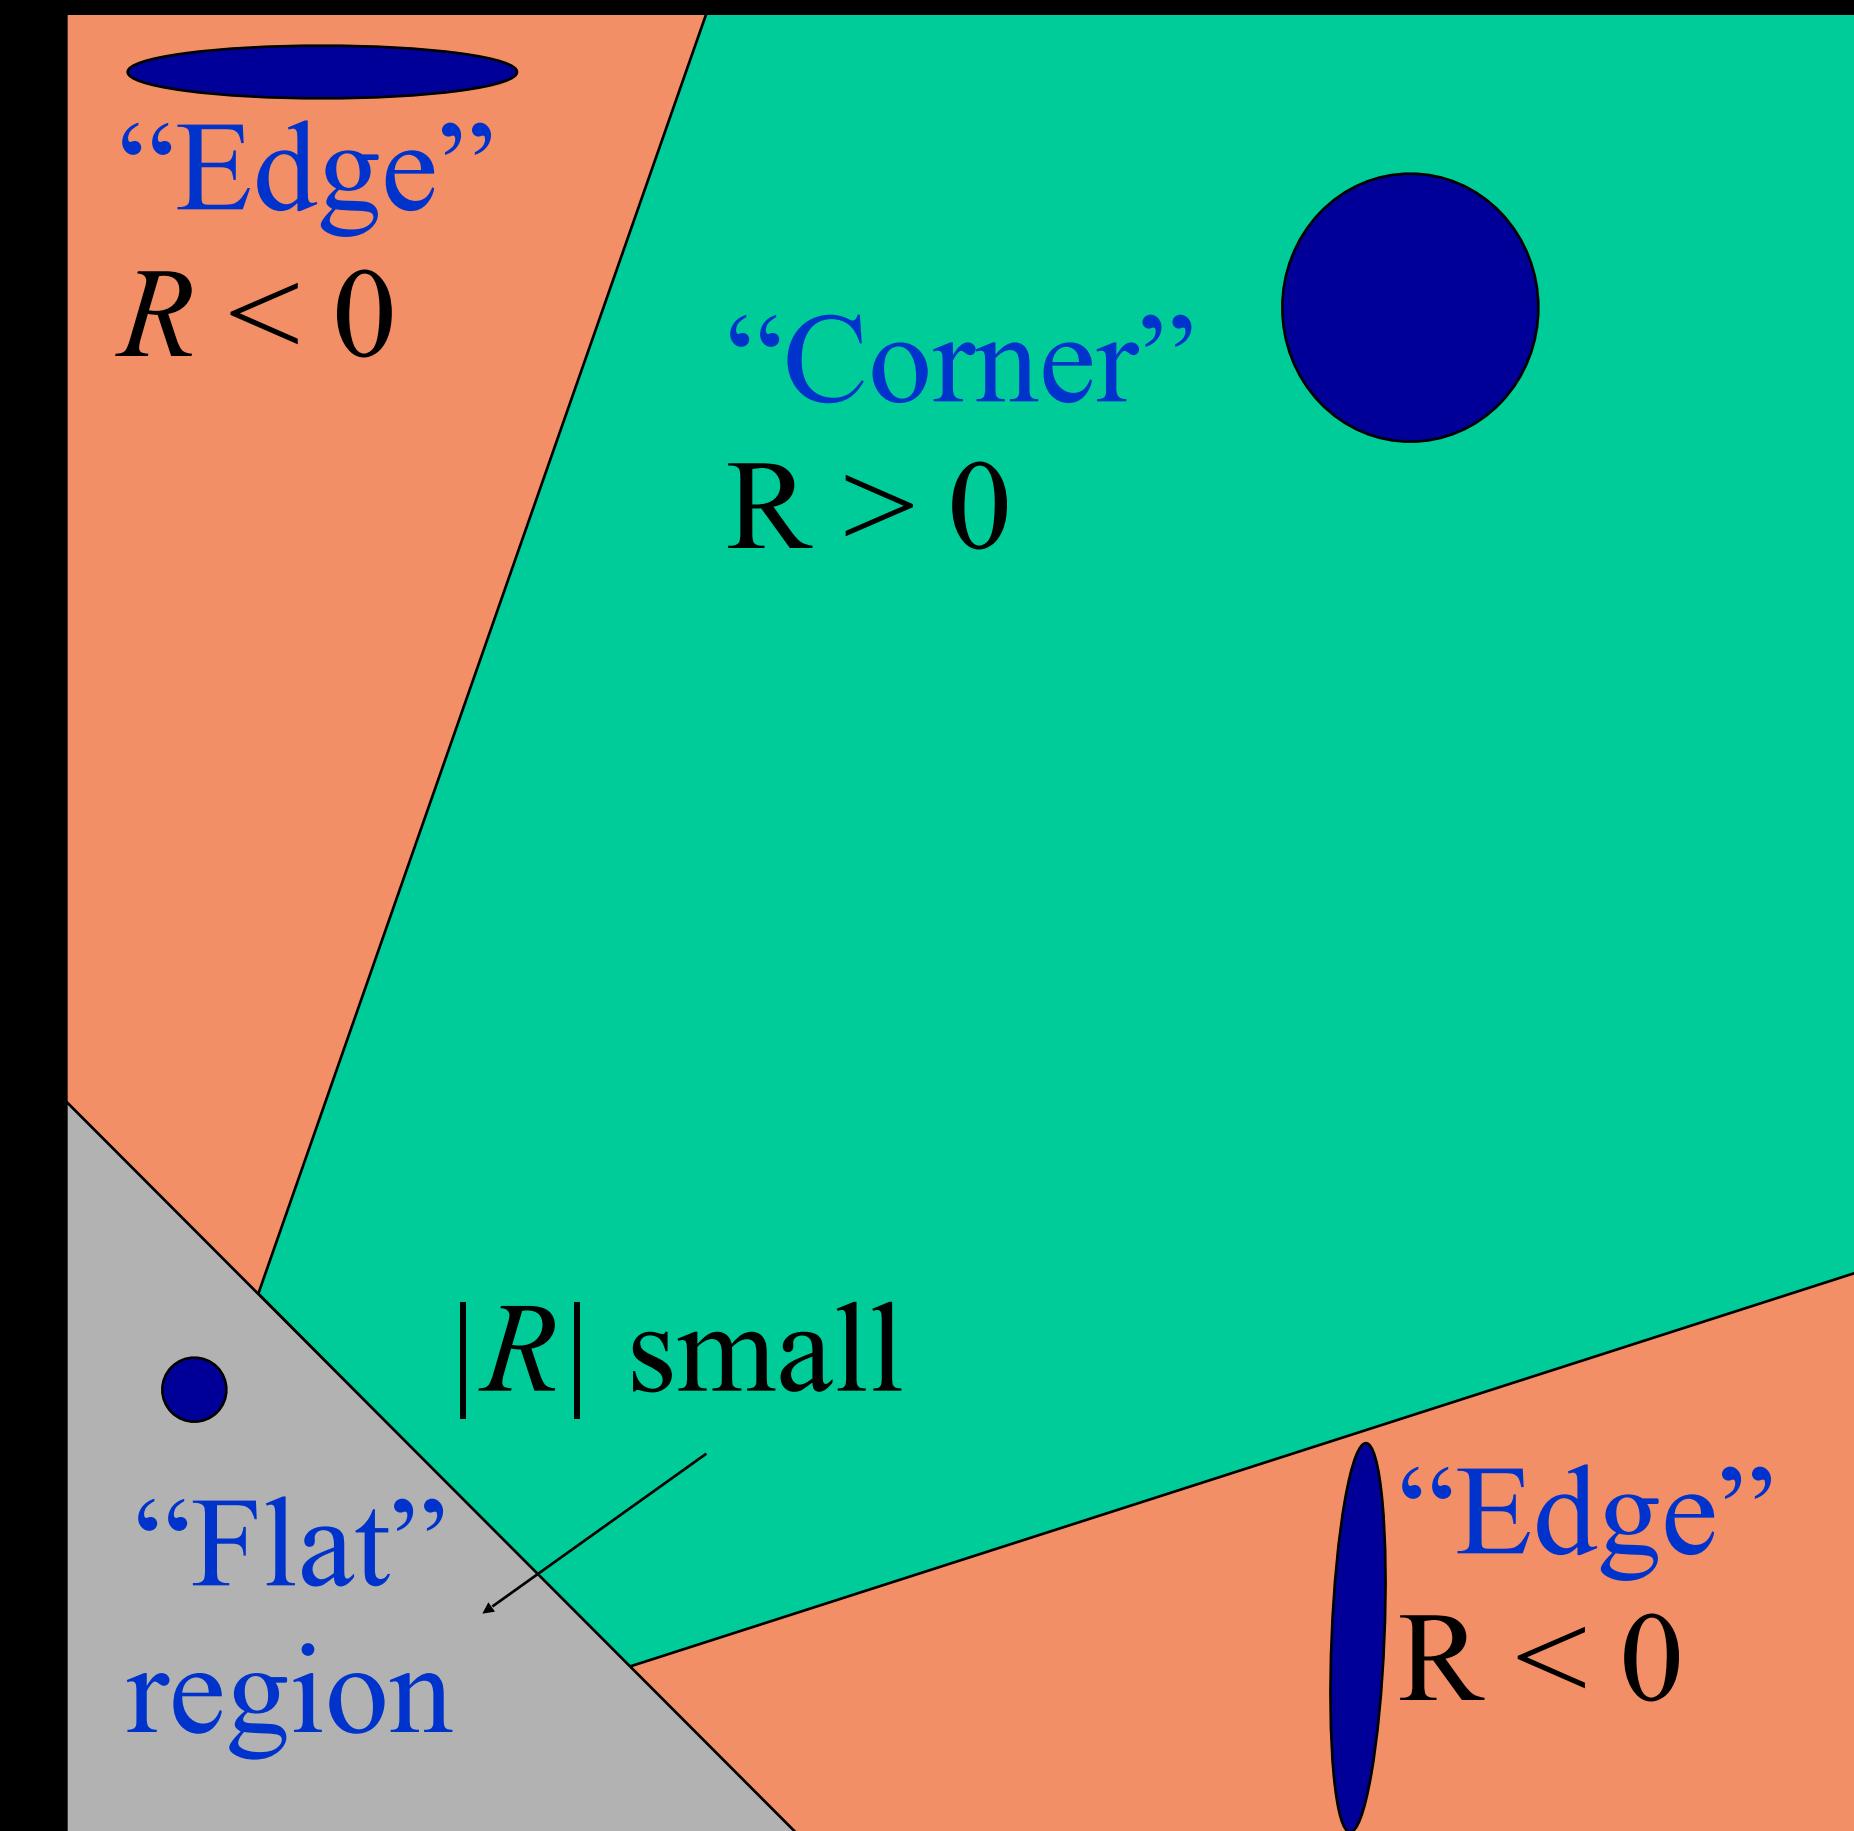
Classification of image points using eigenvalues of M :



CORNER RESPONSE FUNCTION

$$R = \det(M) - \alpha \operatorname{trace}(M)^2 = \lambda_1 \lambda_2 - \alpha (\lambda_1 + \lambda_2)^2$$

α : constant (0.04 to 0.06)



HARRIS CORNER DETECTOR

- 1) Compute M matrix for each image window to get their *cornerness scores*.
- 2) Find points whose surrounding window gave large corner response ($f > \text{threshold}$)
- 3) Take the points of local maxima, i.e., perform non-maximum suppression

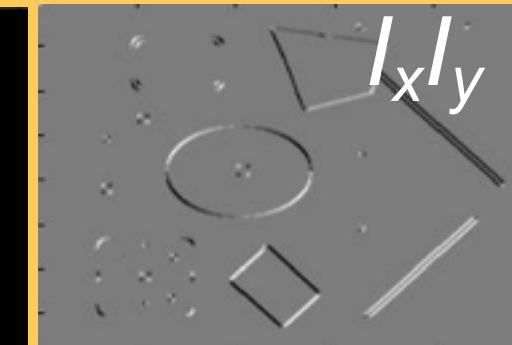
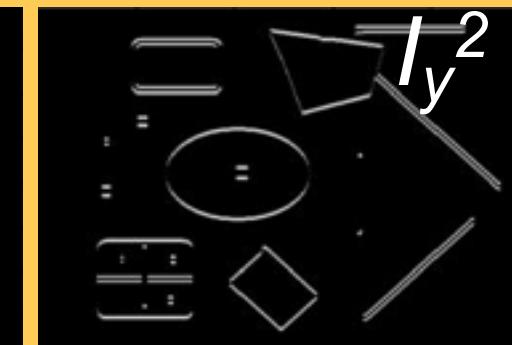
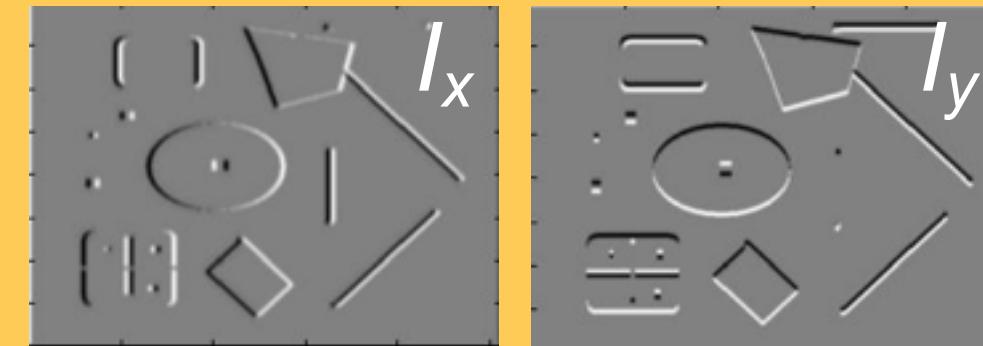
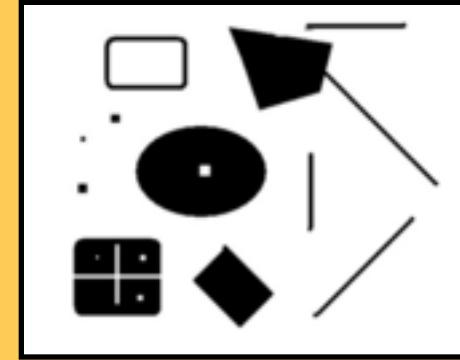
C.Harris and M.Stephens. “A Combined Corner and Edge Detector.” *Proceedings of the 4th Alvey Vision Conference*: pages 147—151, 1988.

Harris Detector [Harris88]

- Second moment matrix

$$\mu(\sigma_I, \sigma_D) = g(\sigma_I) * \begin{bmatrix} I_x^2(\sigma_D) & I_x I_y(\sigma_D) \\ I_x I_y(\sigma_D) & I_y^2(\sigma_D) \end{bmatrix}$$

1. Image derivatives
(optionally, blur first)



$$\det M = \lambda_1 \lambda_2$$

$$\text{trace } M = \lambda_1 + \lambda_2$$

2. Square of derivatives

3. Gaussian filter $g(\sigma_I)$

4. Cornerness function – both eigenvalues are strong

$$har = \det[\mu(\sigma_I, \sigma_D)] - \alpha [\text{trace}(\mu(\sigma_I, \sigma_D))^2] =$$

$$g(I_x^2)g(I_y^2) - [g(I_x I_y)]^2 - \alpha[g(I_x^2) + g(I_y^2)]^2$$

5. Non-maxima suppression



TEXT

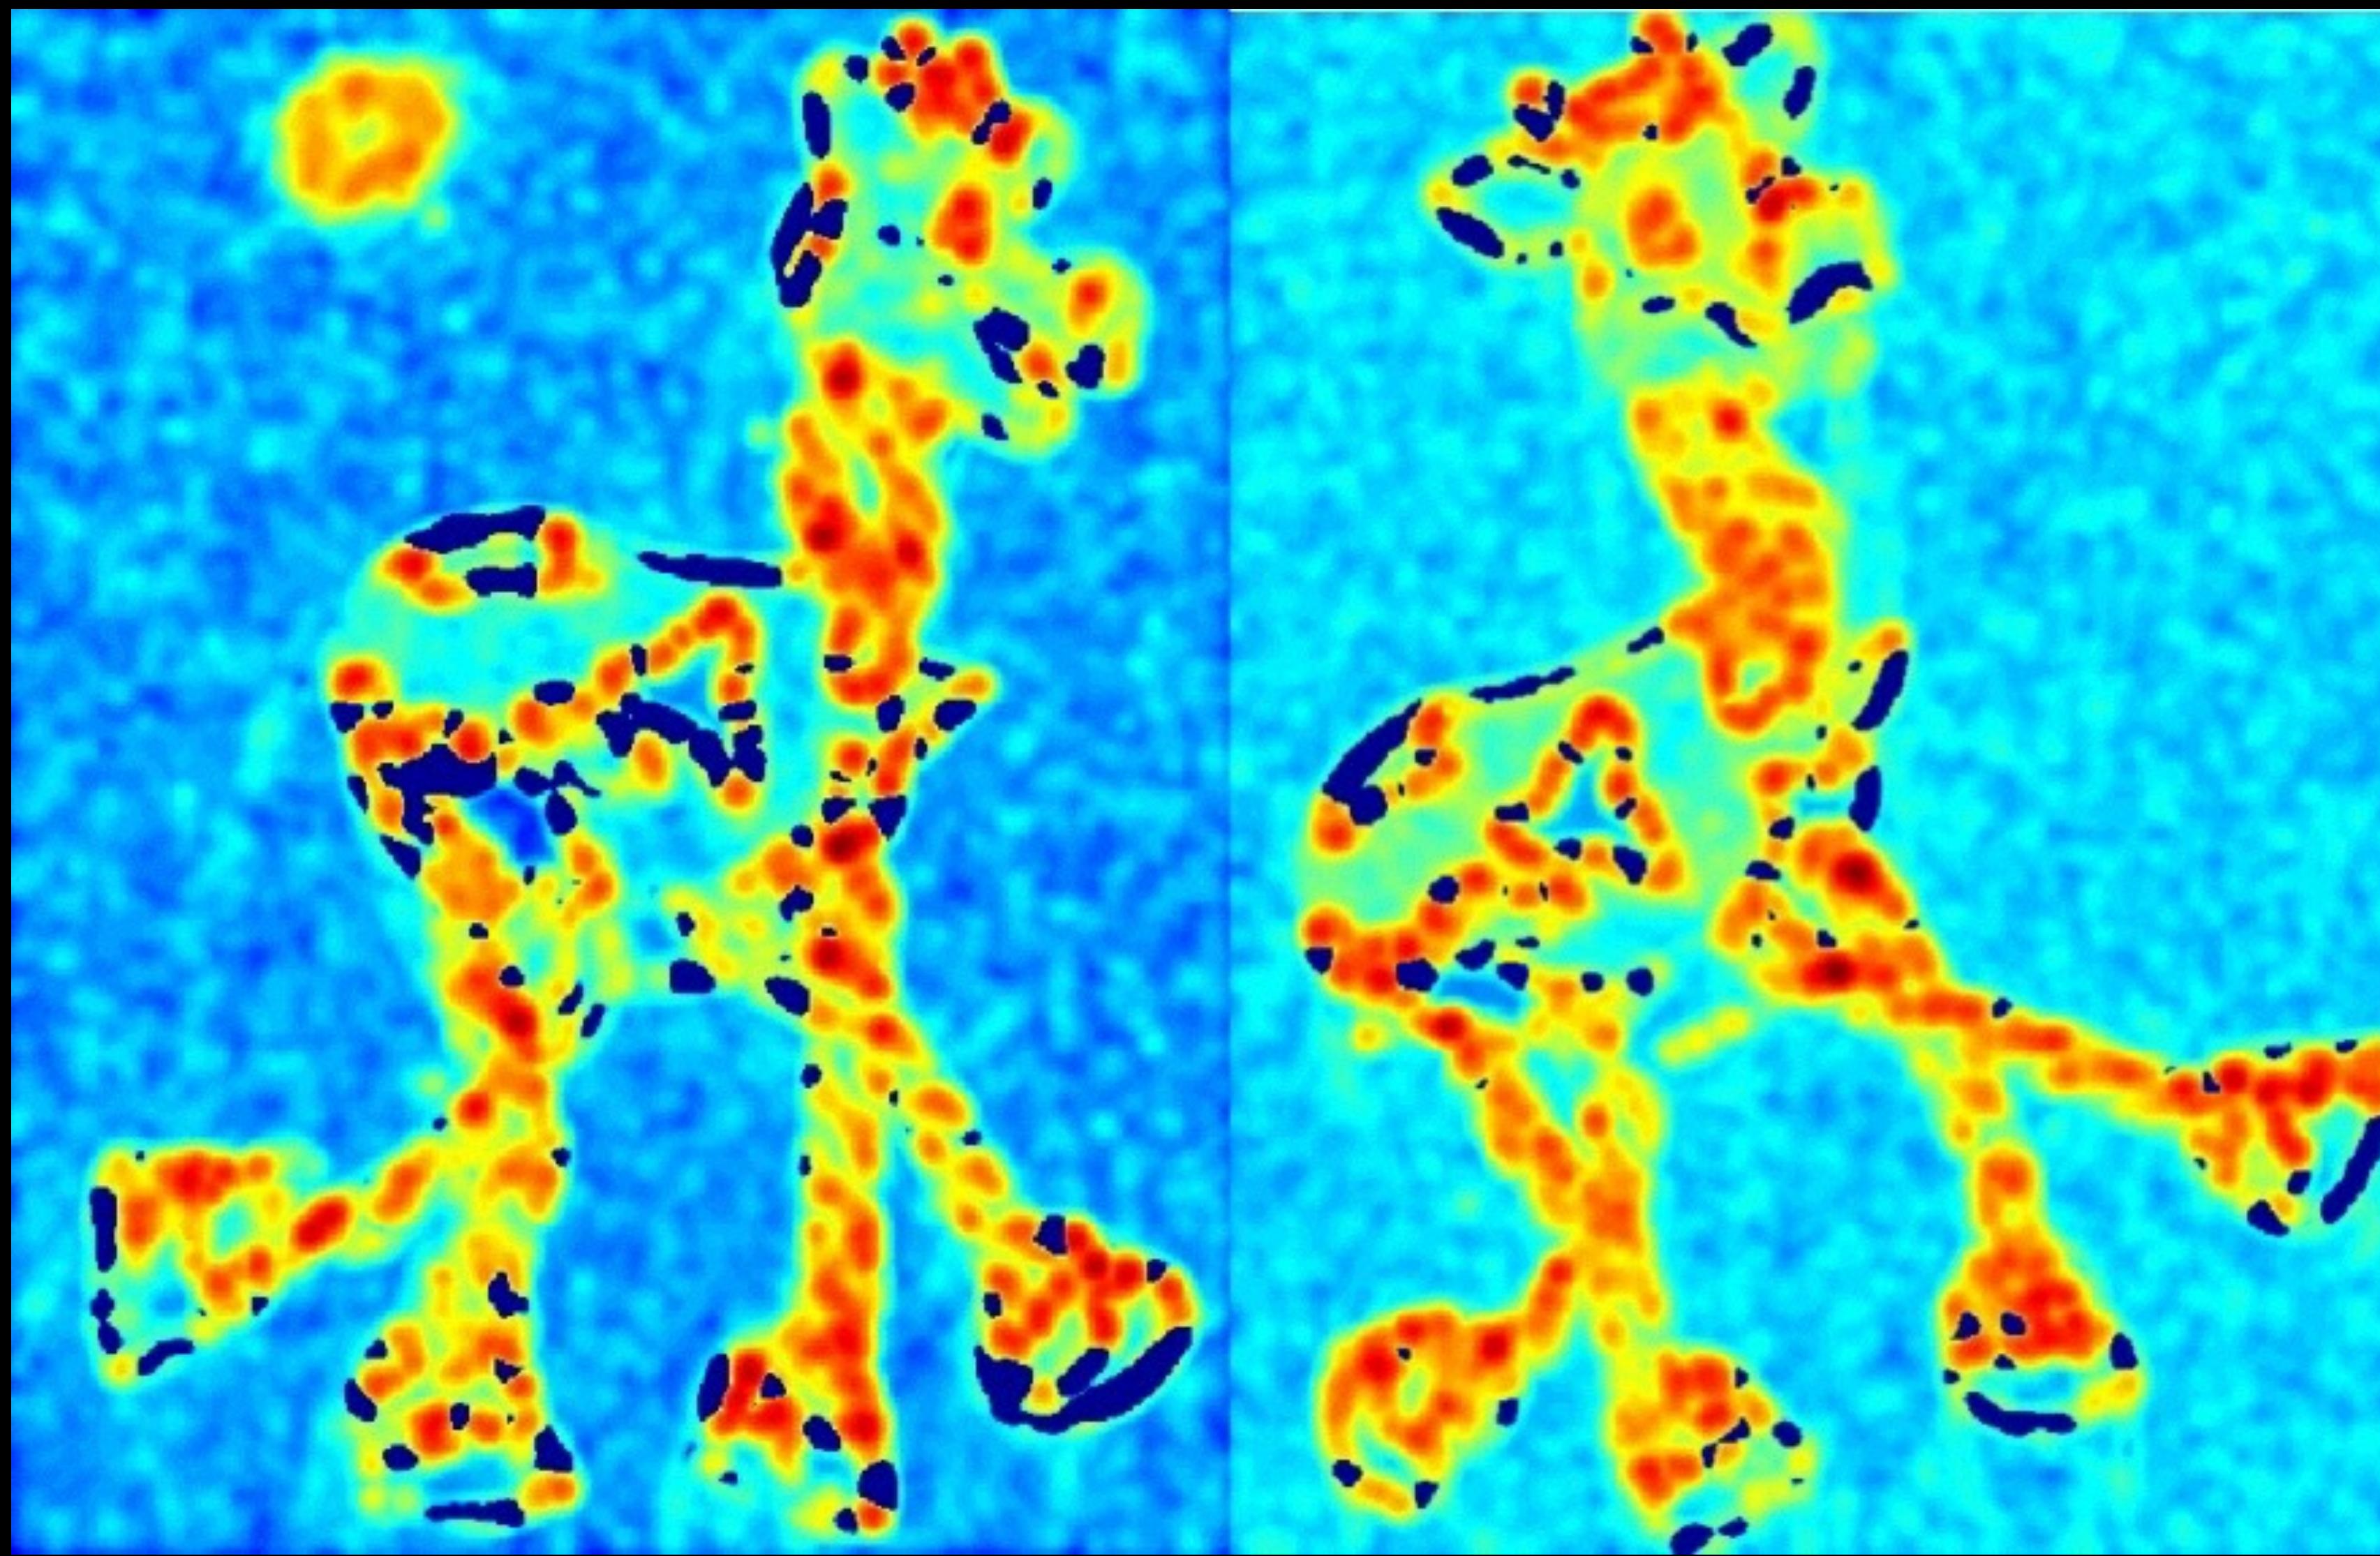
HARRIS DETECTOR: STEPS



TEXT

HARRIS DETECTOR: STEPS

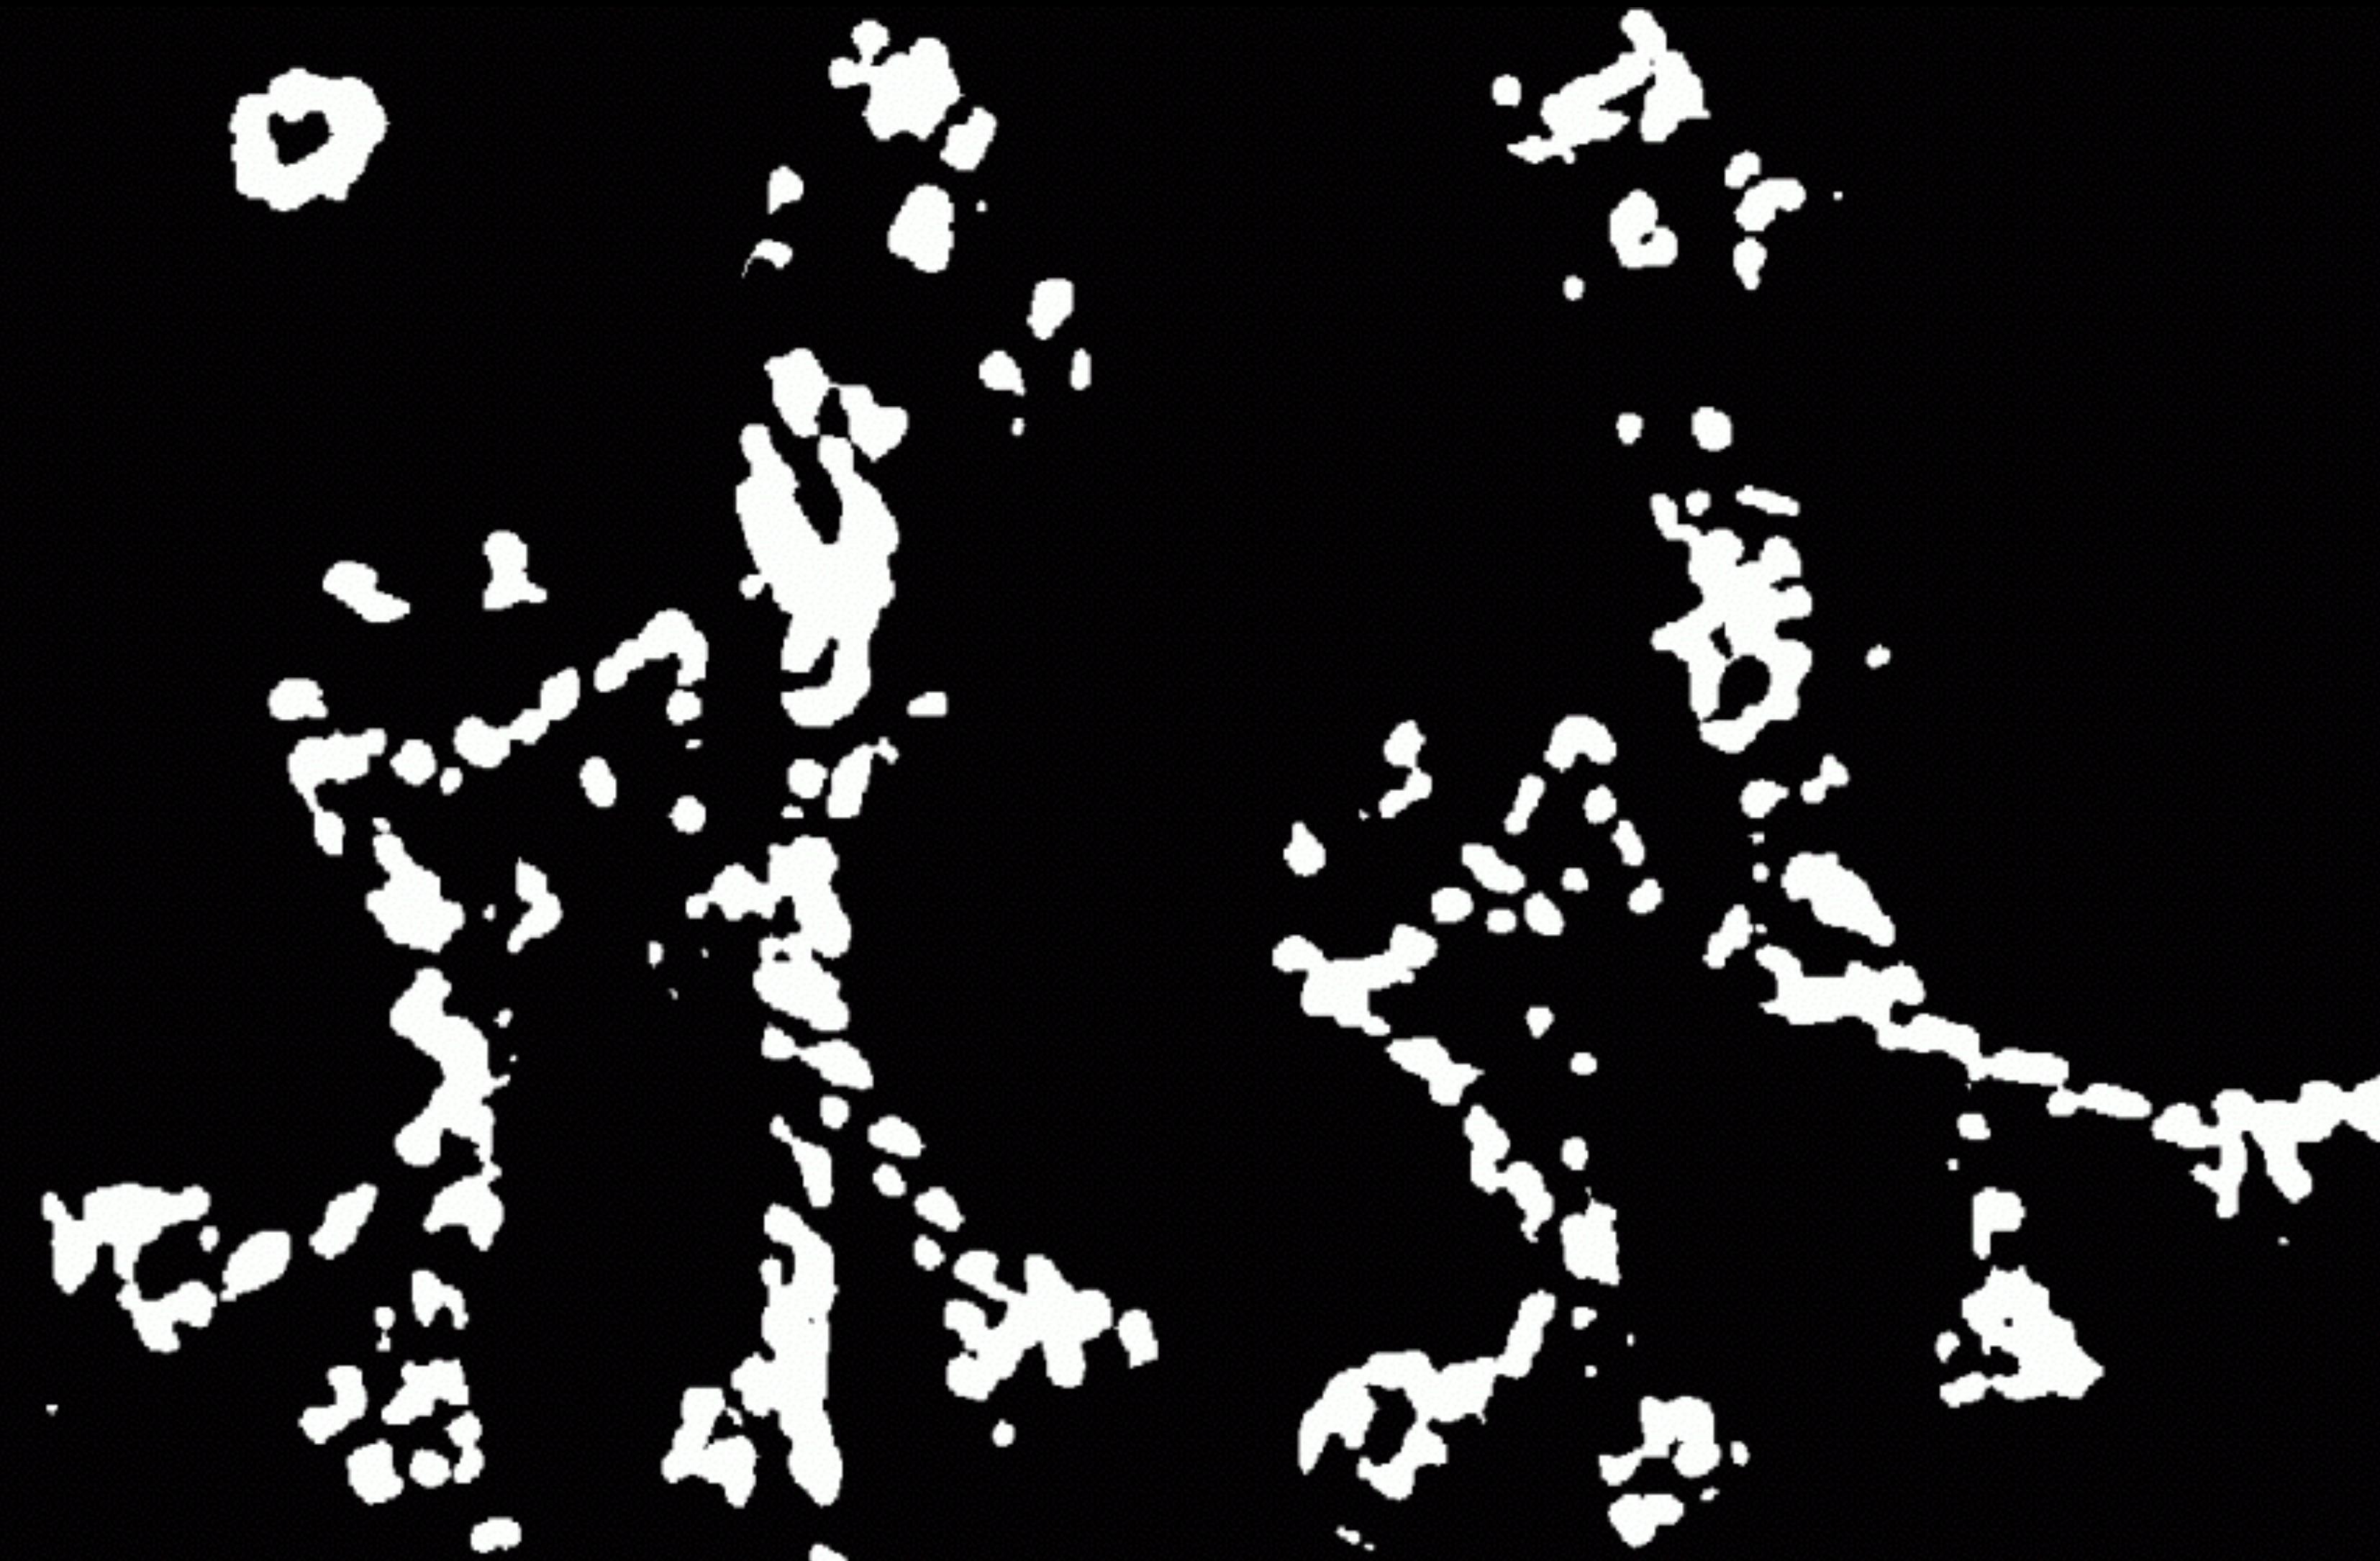
Compute corner response R



TEXT

HARRIS DETECTOR: STEPS

Find points with large corner response: $R > \text{threshold}$



TEXT

HARRIS DETECTOR: STEPS

Take only the points of local maxima of R



TEXT

HARRIS DETECTOR: STEPS

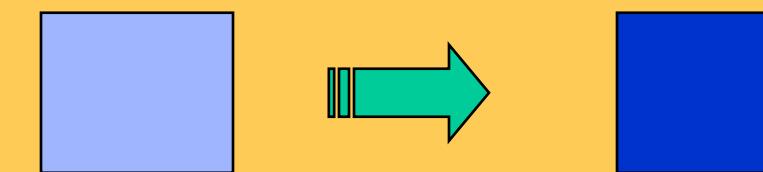


INvariance AND COVARIANCE

- ▶ We want corner locations to be *invariant* to photometric transformations and *covariant* to geometric transformations
- ▶ **Invariance:** image is transformed and corner locations do not change
- ▶ **Covariance:** if we have two transformed versions of the same image, features should be detected in corresponding locations

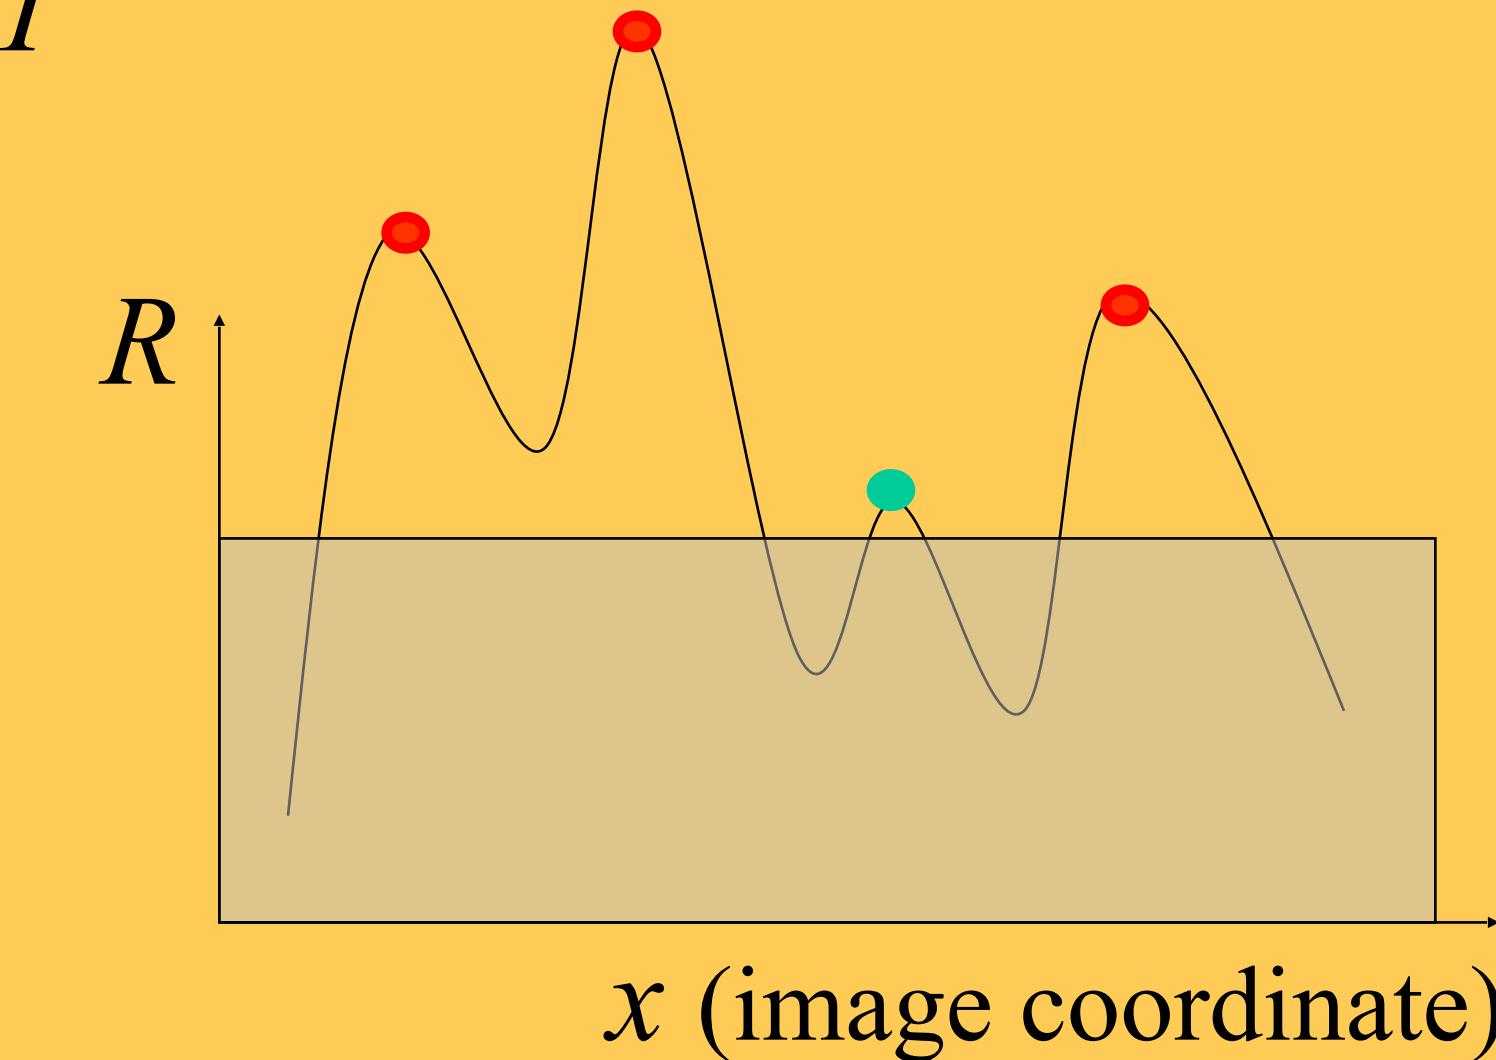
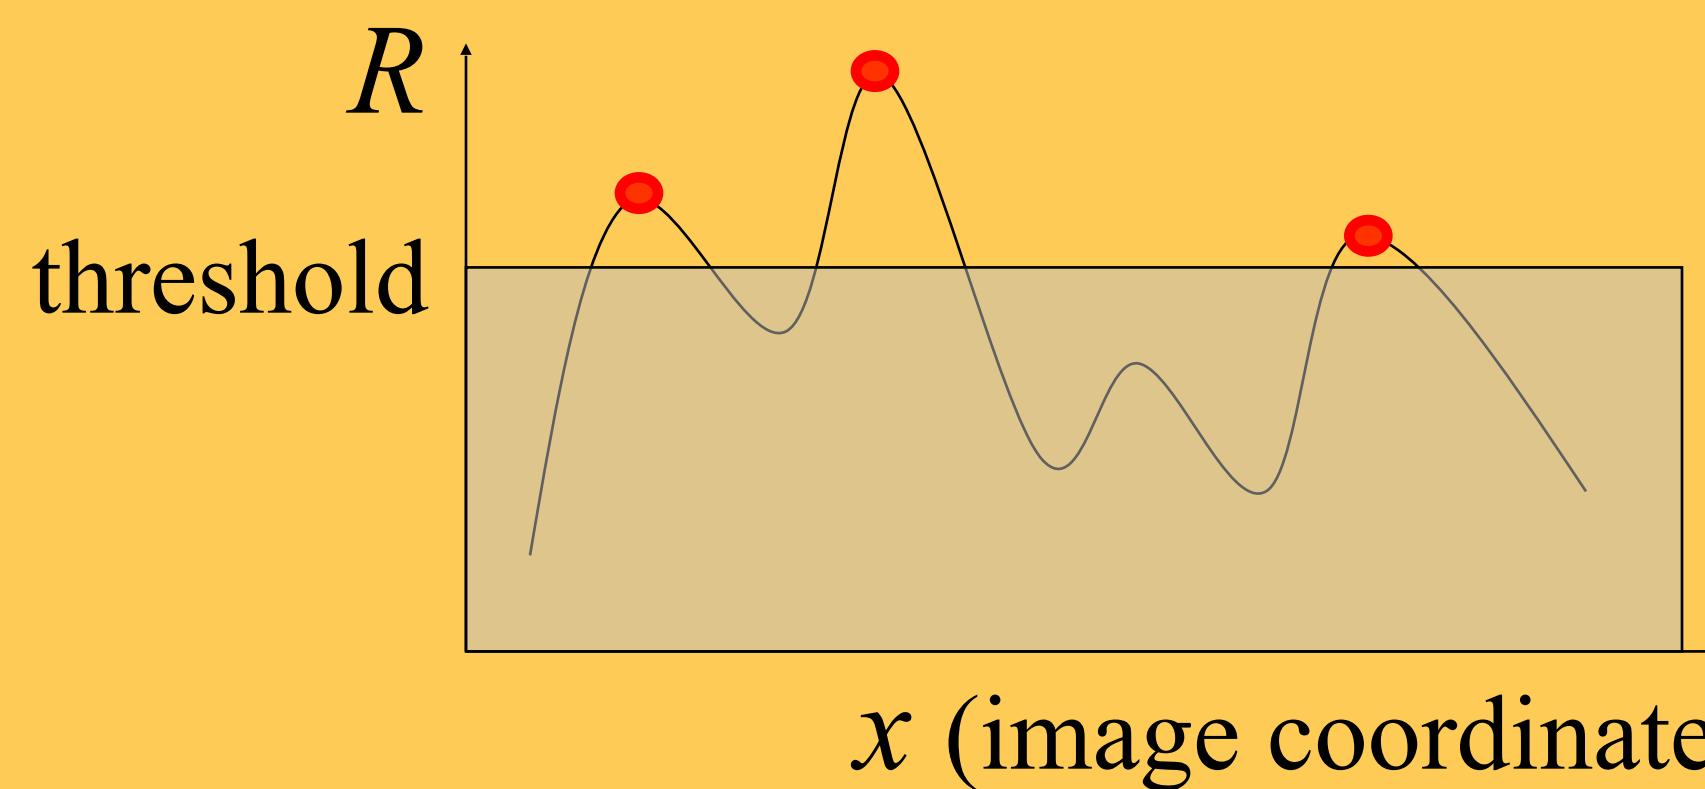


AFFINE INTENSITY CHANGE



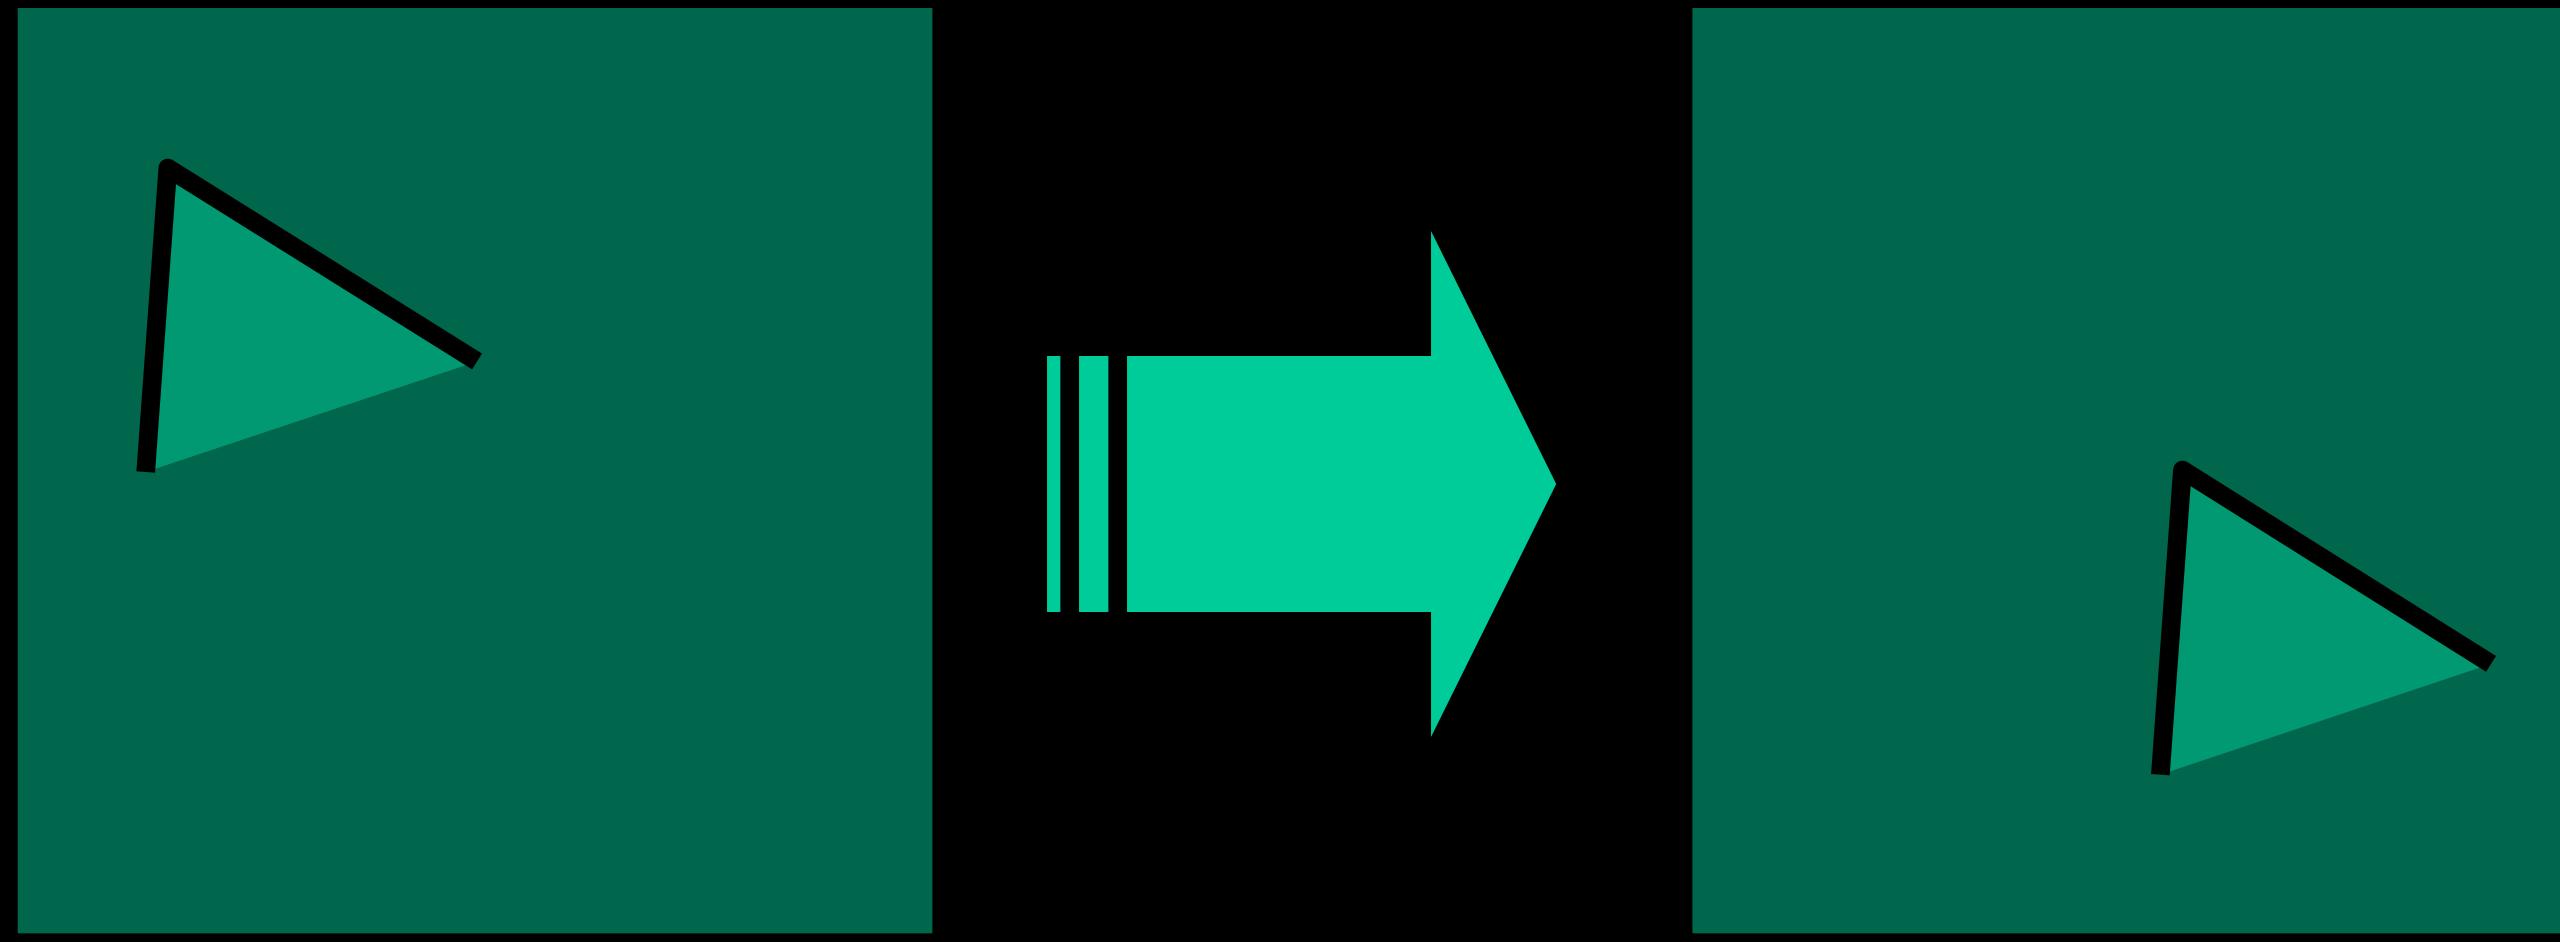
$$I \rightarrow a I + b$$

- Only derivatives are used => invariance to intensity shift $I \rightarrow I + b$
- Intensity scaling: $I \rightarrow a I$



Partially invariant to affine intensity change

IMAGE TRANSLATION

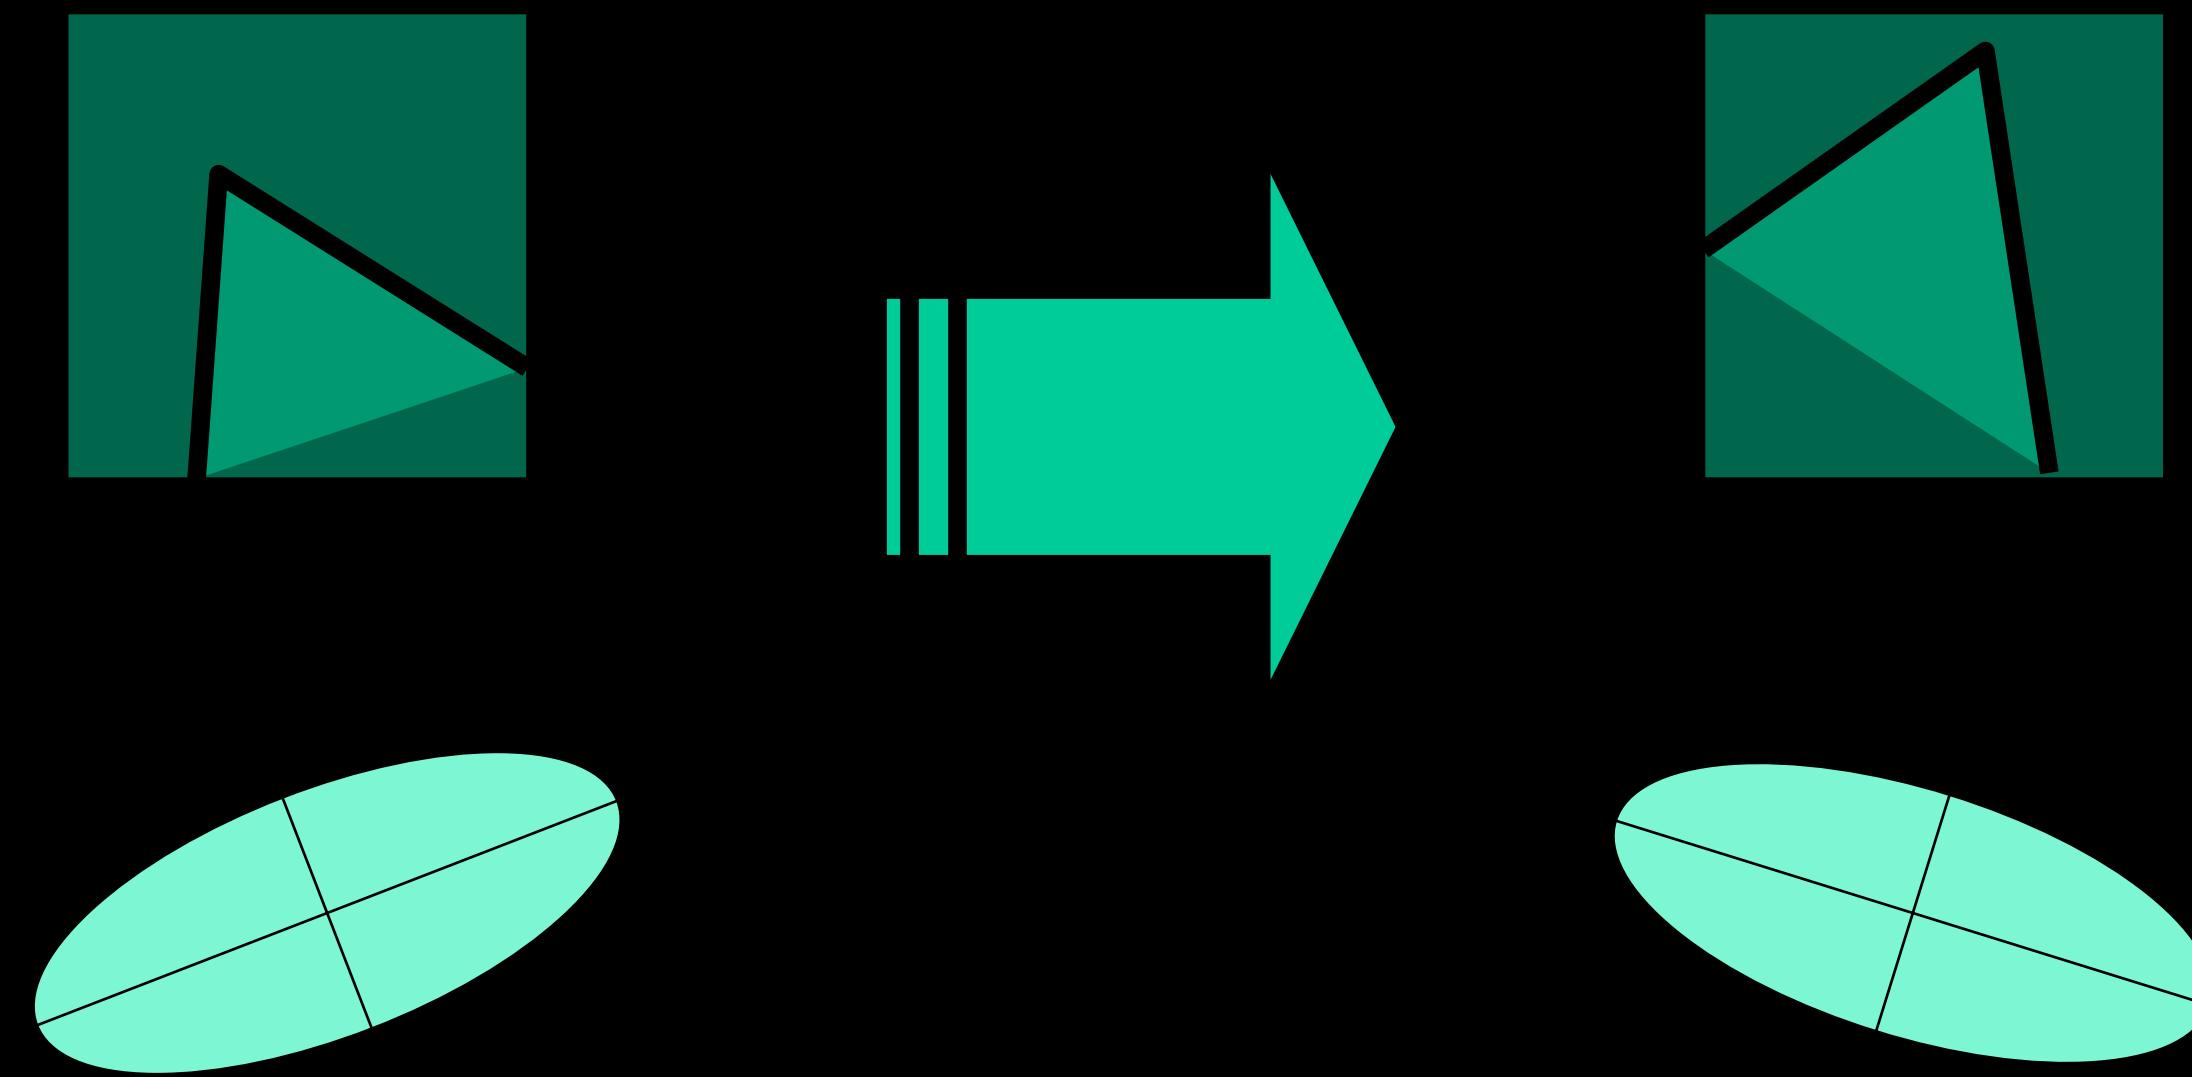


- Derivatives and window function are shift-invariant

Corner location is covariant w.r.t. translation

TEXT

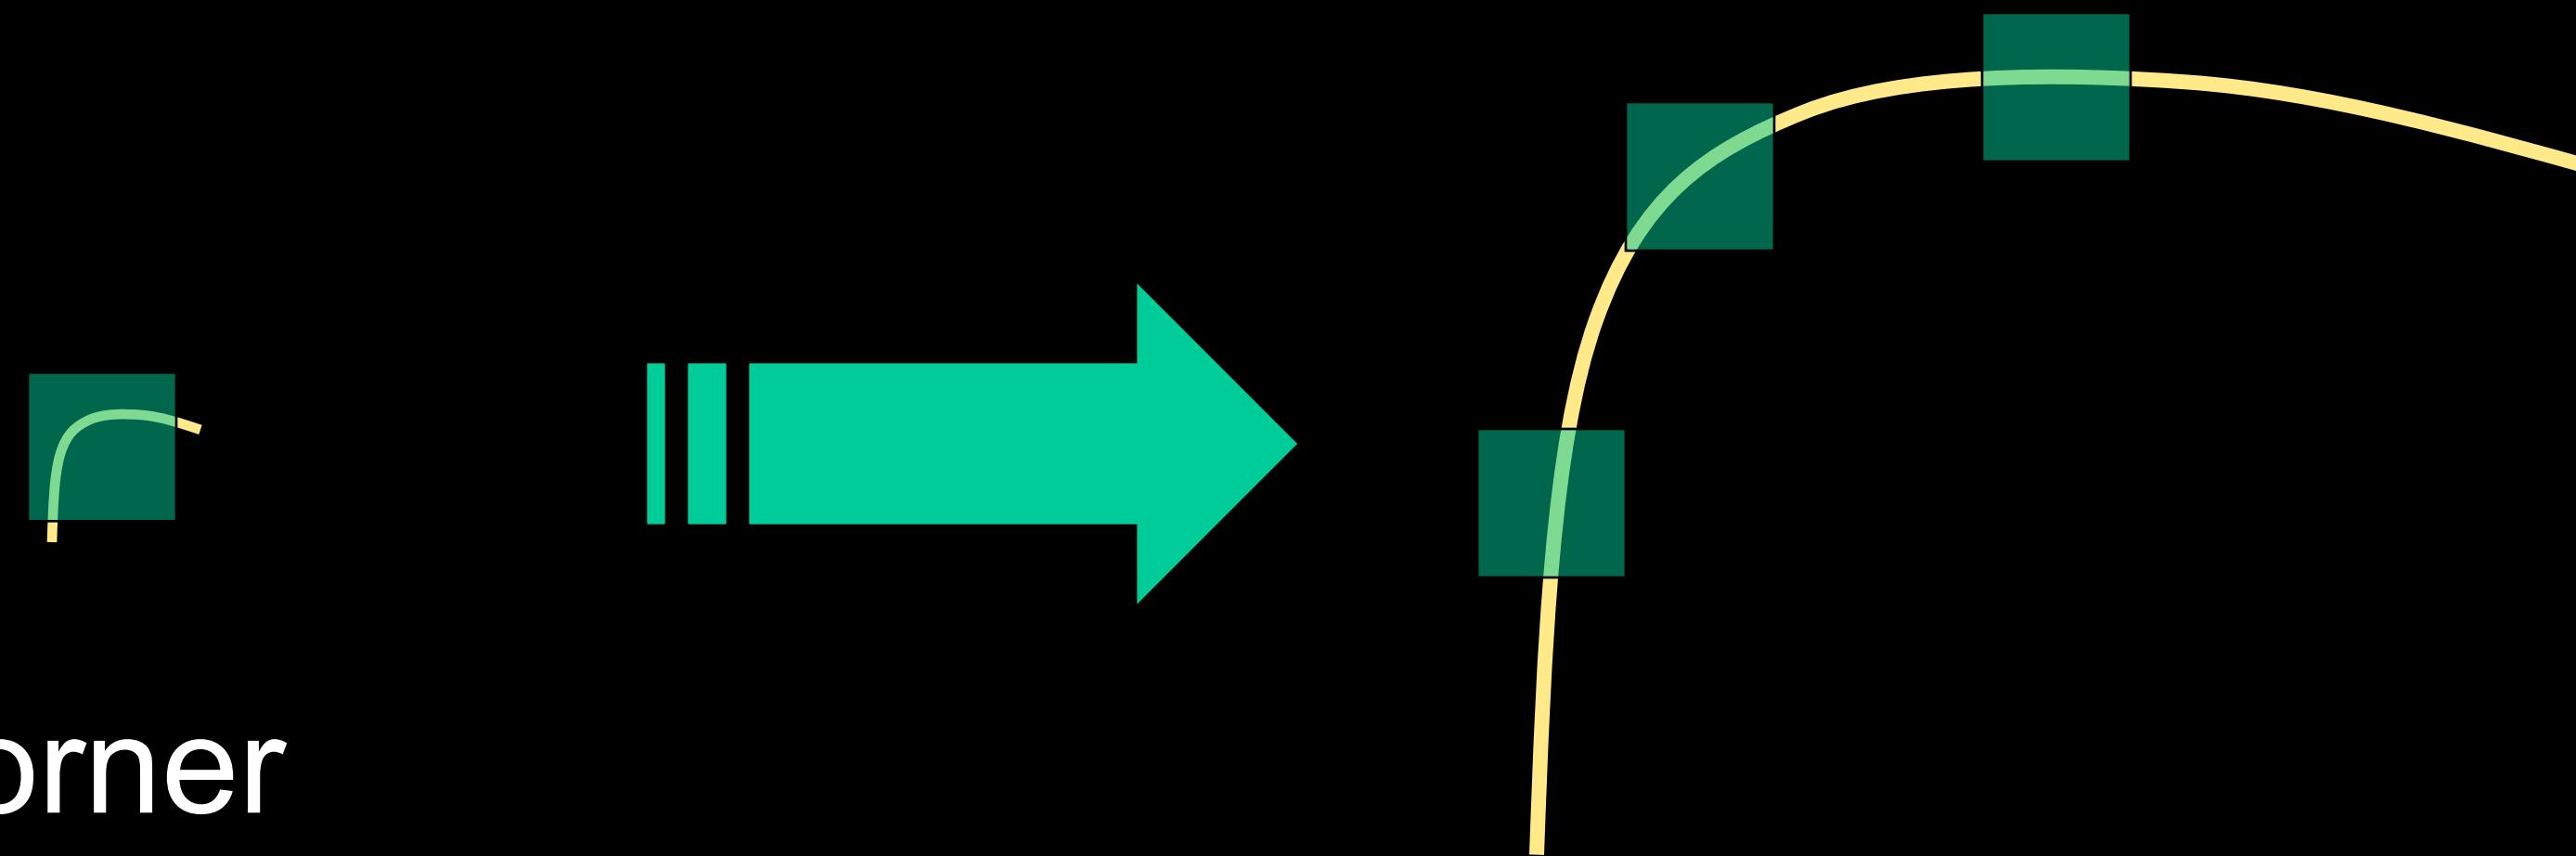
IMAGE ROTATION



Second moment ellipse rotates but its shape (i.e. eigenvalues) remains the same

Corner location is covariant w.r.t. rotation

SCALING



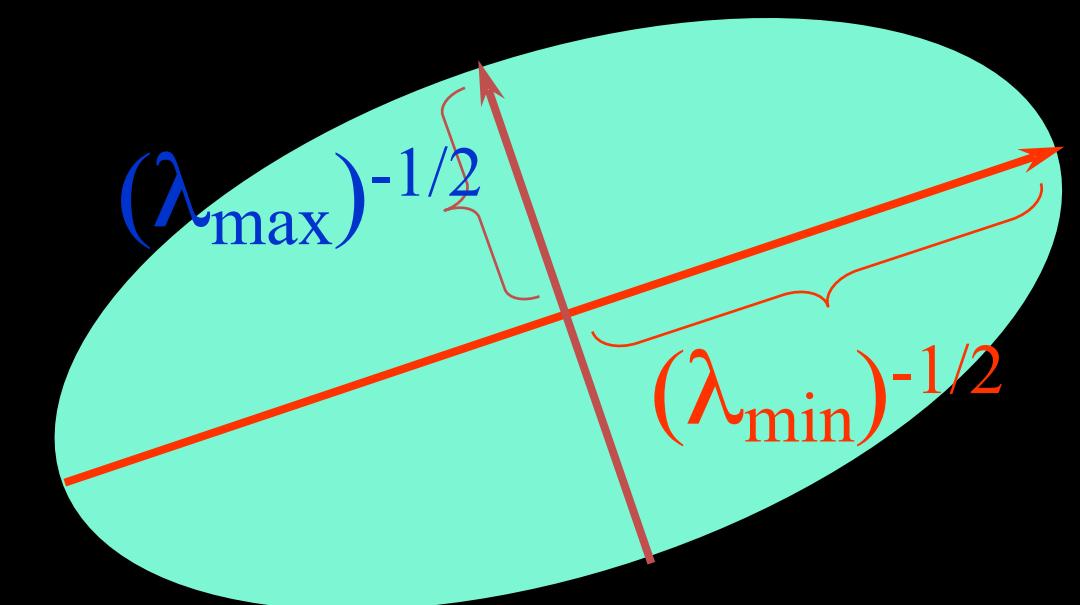
Corner

All points will
be classified as
edges

Corner location is not covariant to scaling!

REVIEW: HARRIS CORNER DETECTOR

- ▶ Approximate distinctiveness by local auto-correlation.
- ▶ Approximate local auto-correlation by second moment matrix
- ▶ Quantify distinctiveness (or cornerness) as function of the eigenvalues of the second moment matrix.
- ▶ But we don't actually need to compute the eigenvalues by using the determinant and trace of the second moment matrix.



TEXT

SO FAR: CAN LOCALIZE IN X-Y, BUT NOT SCALE



AUTOMATIC SCALE SELECTION

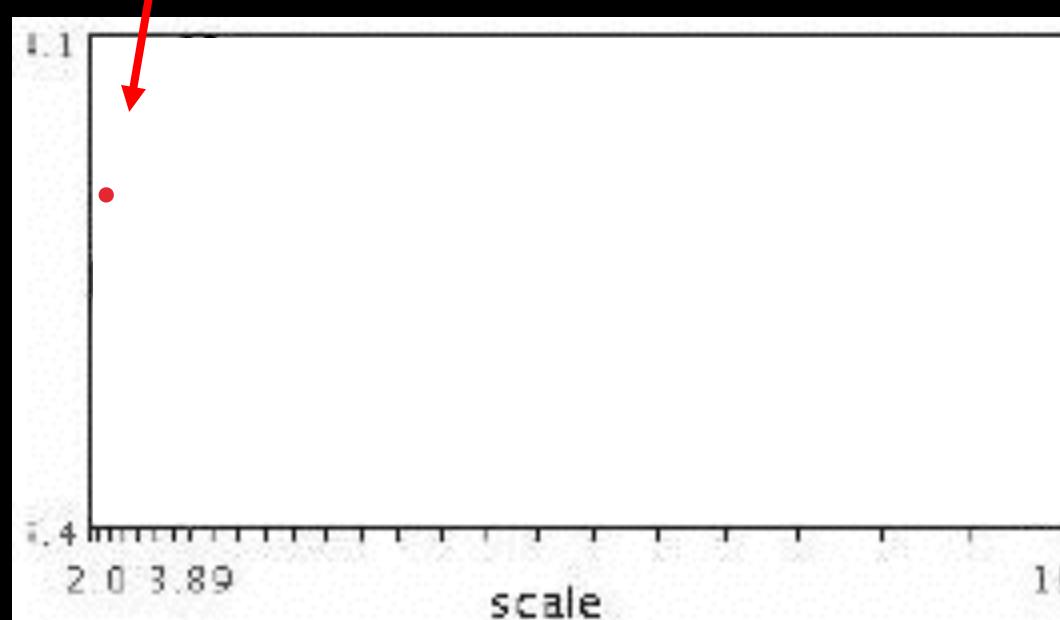


$$f(I_{i_1 \sqcup i_m}(x, \sigma)) = f(I_{i_1 \sqcup i_m}(x', \sigma'))$$

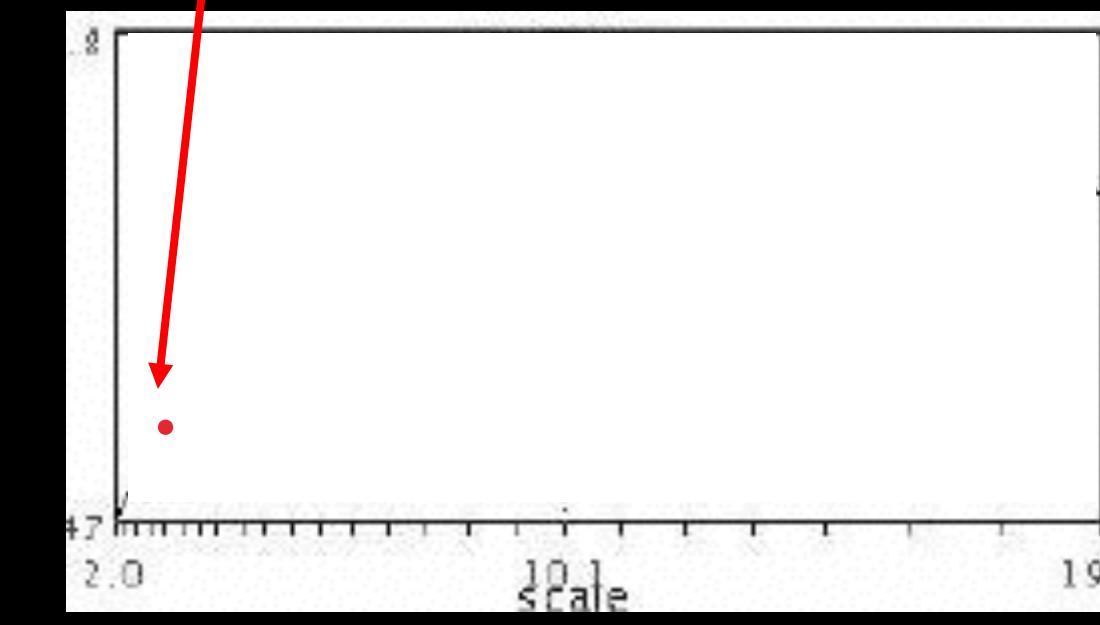
How to find corresponding patch sizes?

AUTOMATIC SCALE SELECTION

- ▶ Function responses for increasing scale (scale signature)

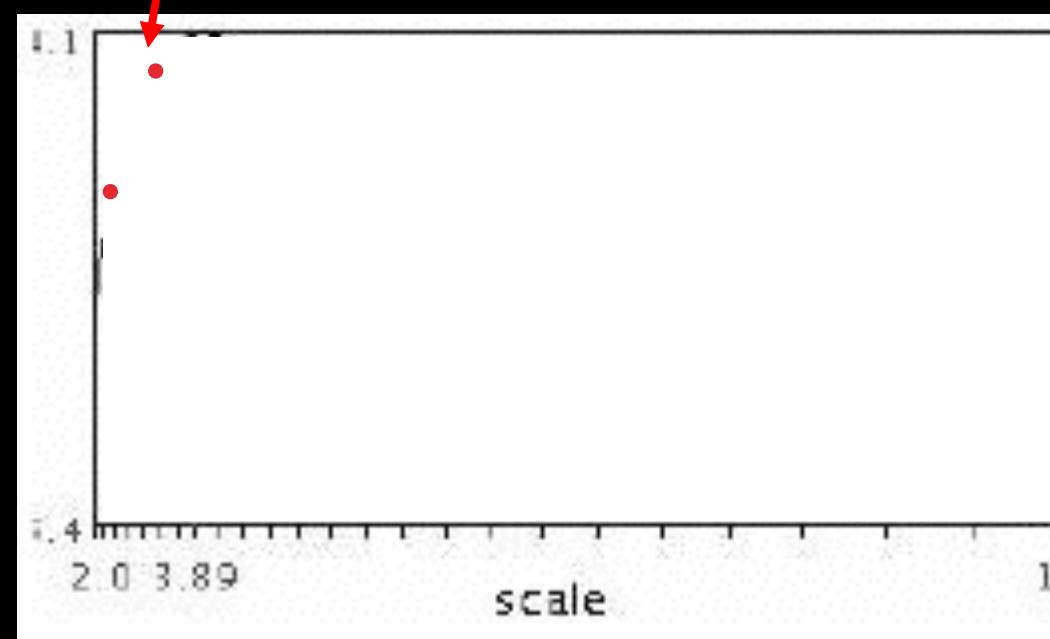


$$f(I_{i_1 \sqcup i_m}(x, \sigma))$$

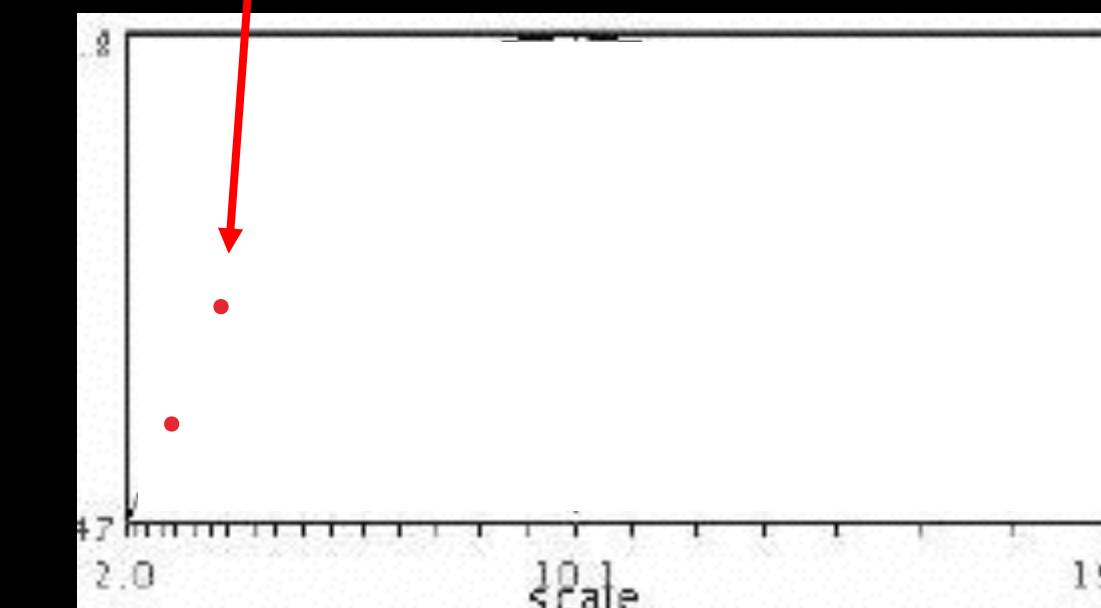


$$f(I_{i_1 \sqcup i_m}(x', \sigma'))$$

AUTOMATIC SCALE SELECTION



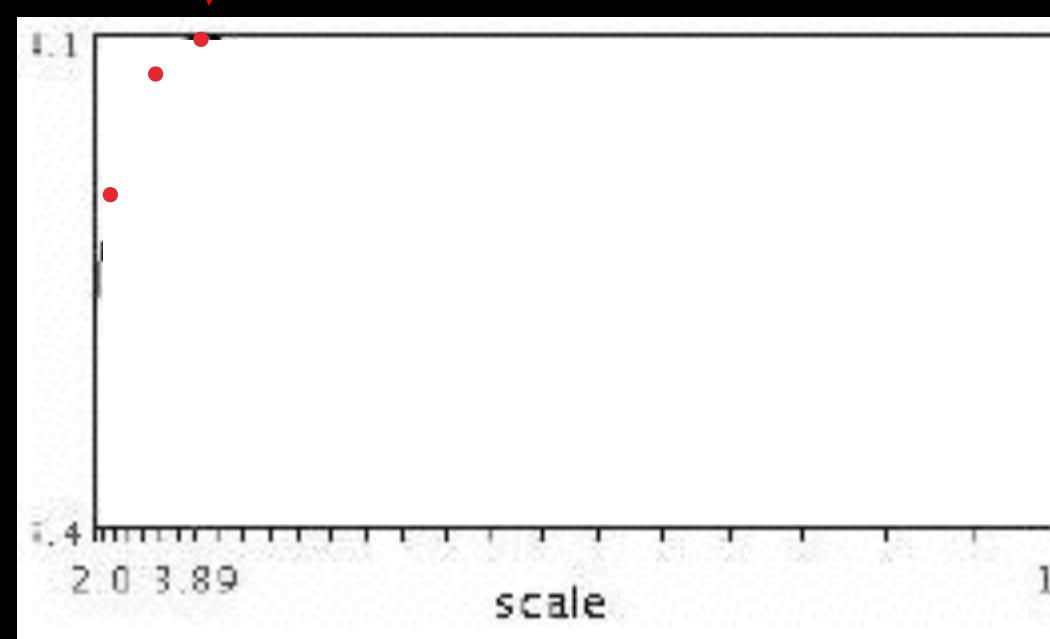
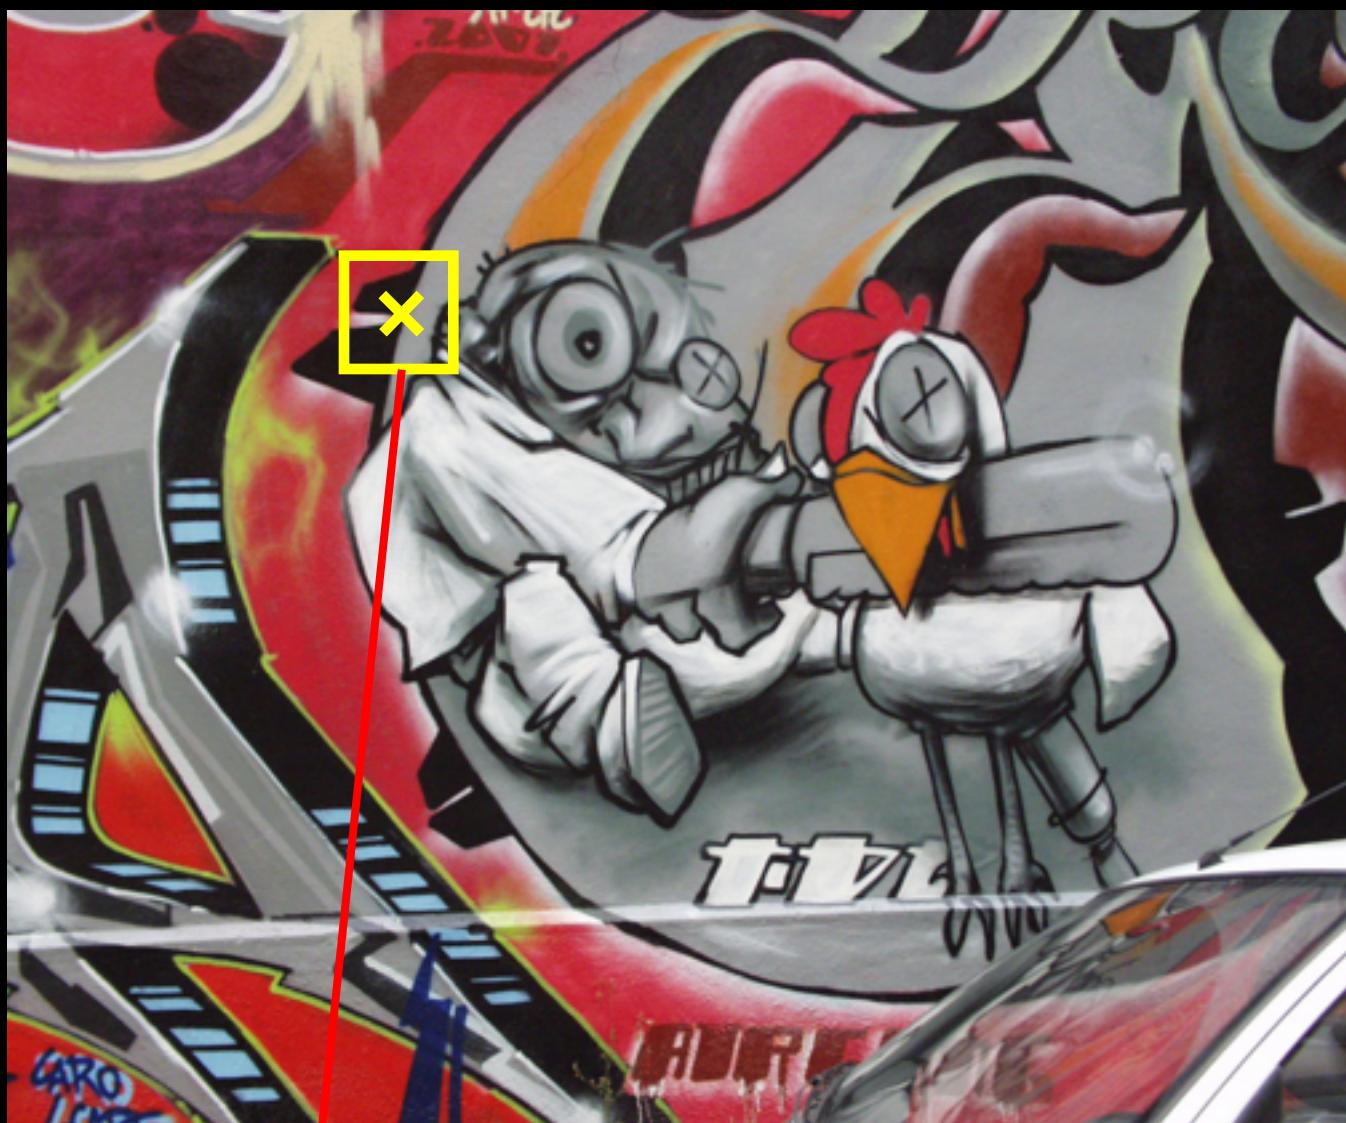
$$f(I_{i_1 \square j_m}(x, \sigma))$$



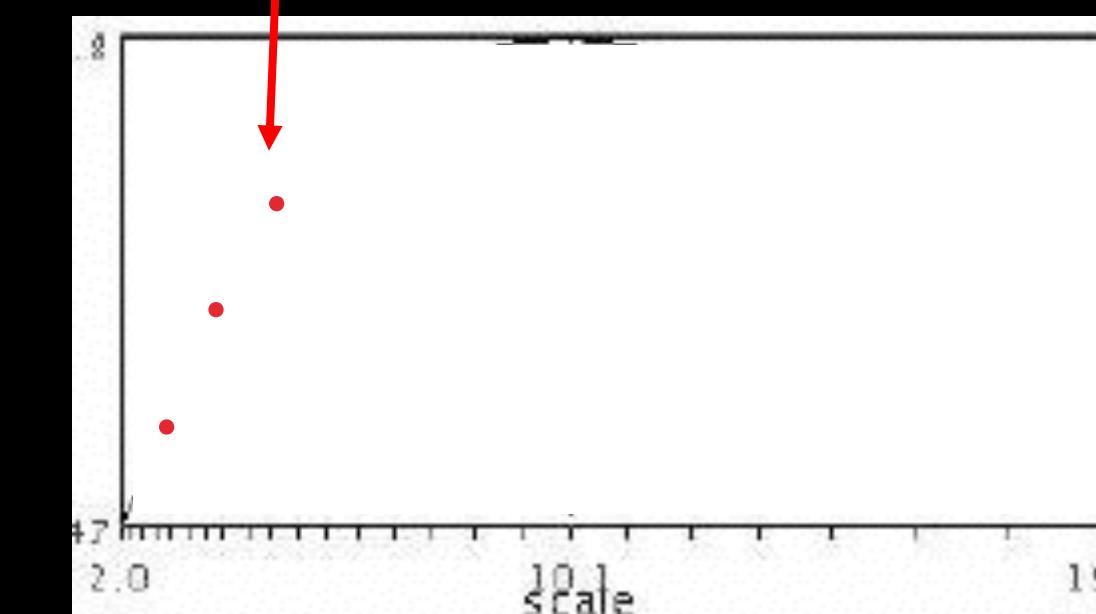
$$f(I_{i_1 \square j_m}(x', \sigma'))$$

TEXT

AUTOMATIC SCALE SELECTION

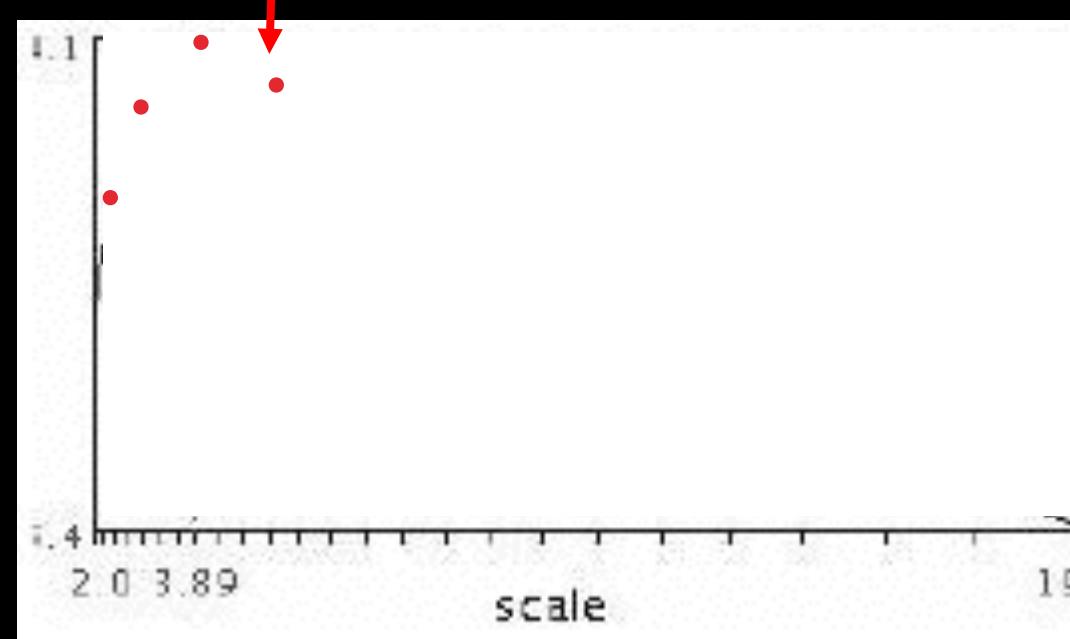
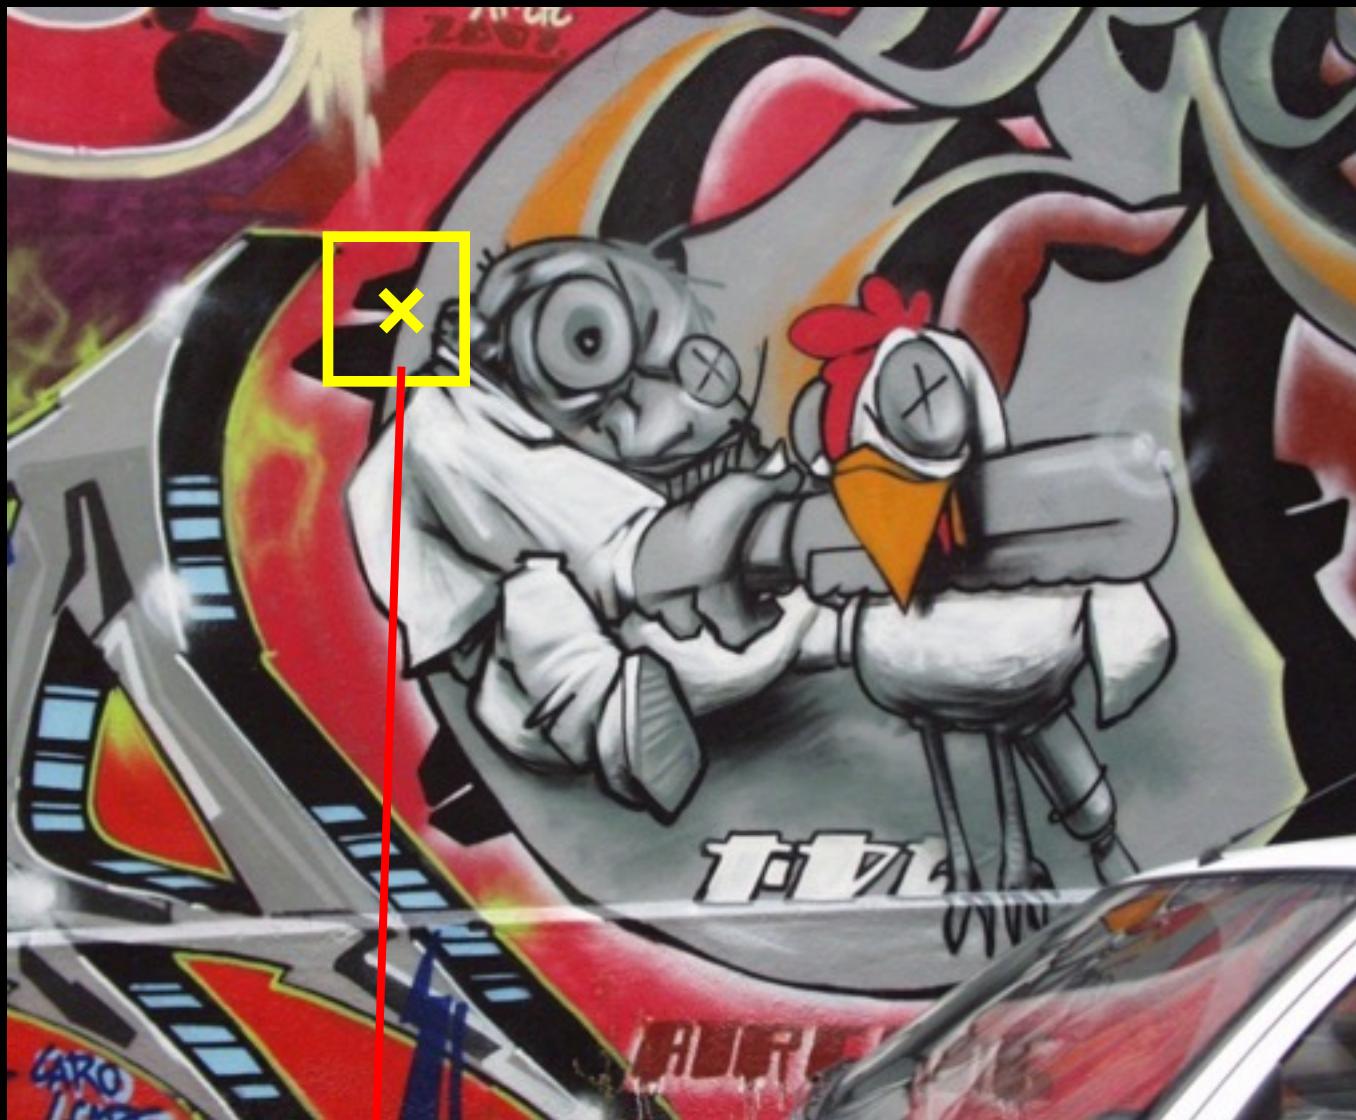


$$f(I_{i_1 \sqcap \cdots \sqcap i_m}(x,\sigma))$$

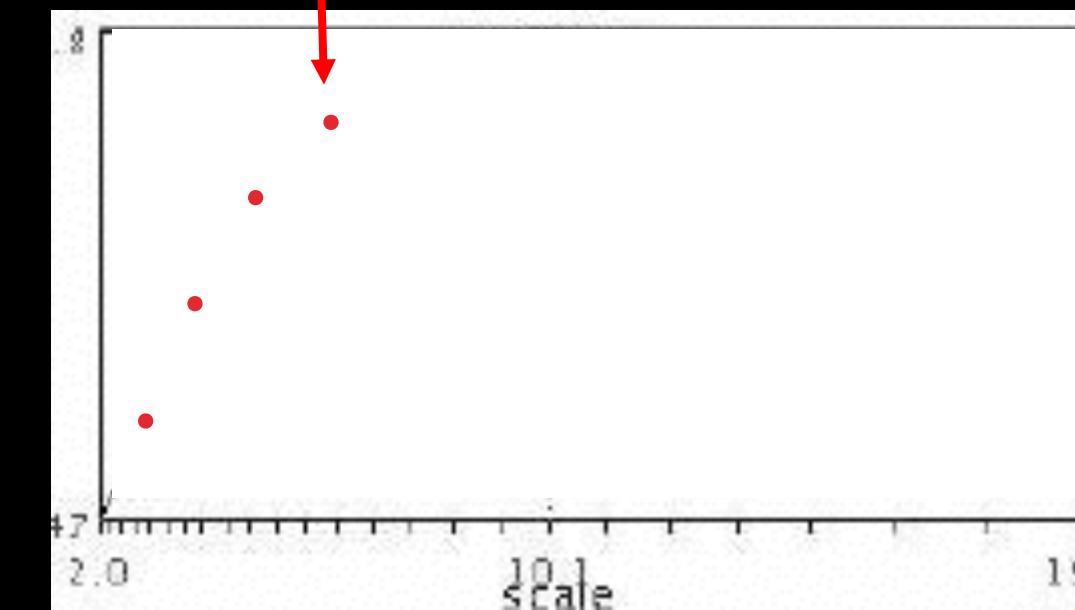


$$f(I_{i_1 \blacksquare i_m}(x', \sigma'))$$

AUTOMATIC SCALE SELECTION

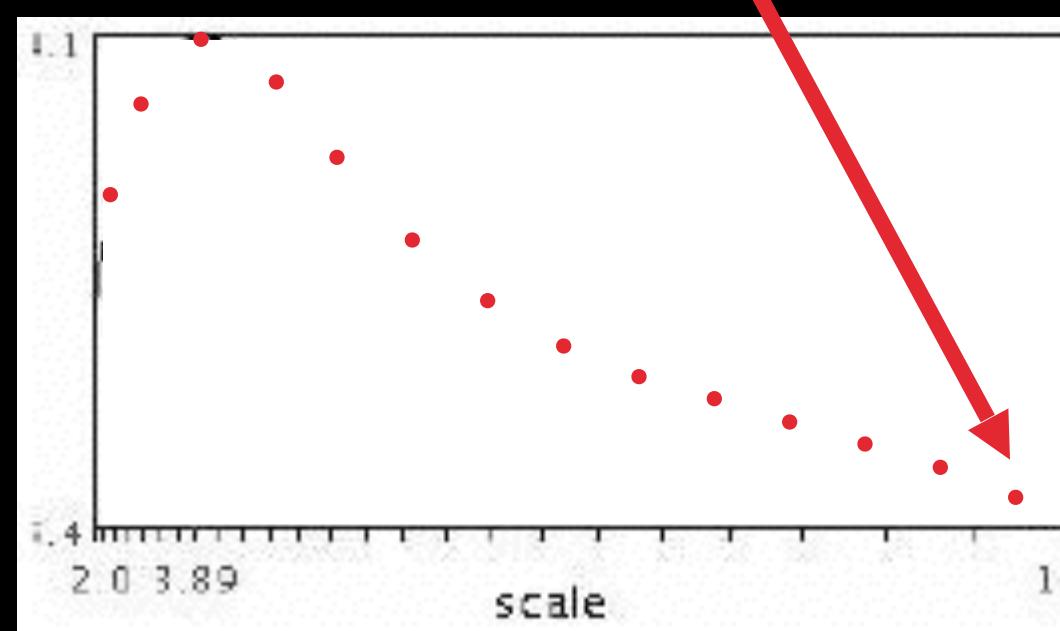
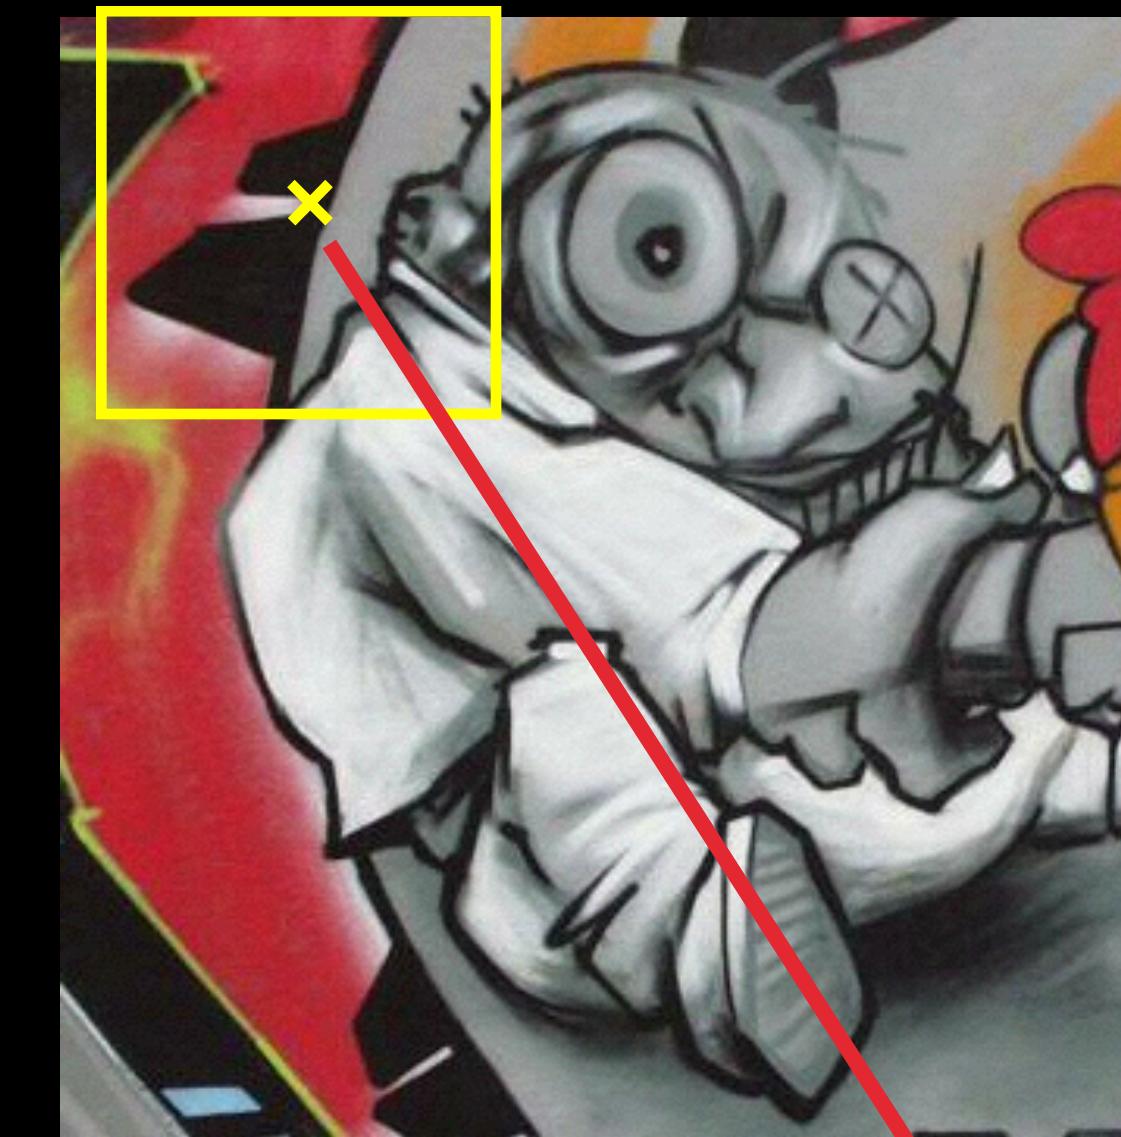


$$f(I_{i_1 \oplus i_m}(x, \sigma))$$

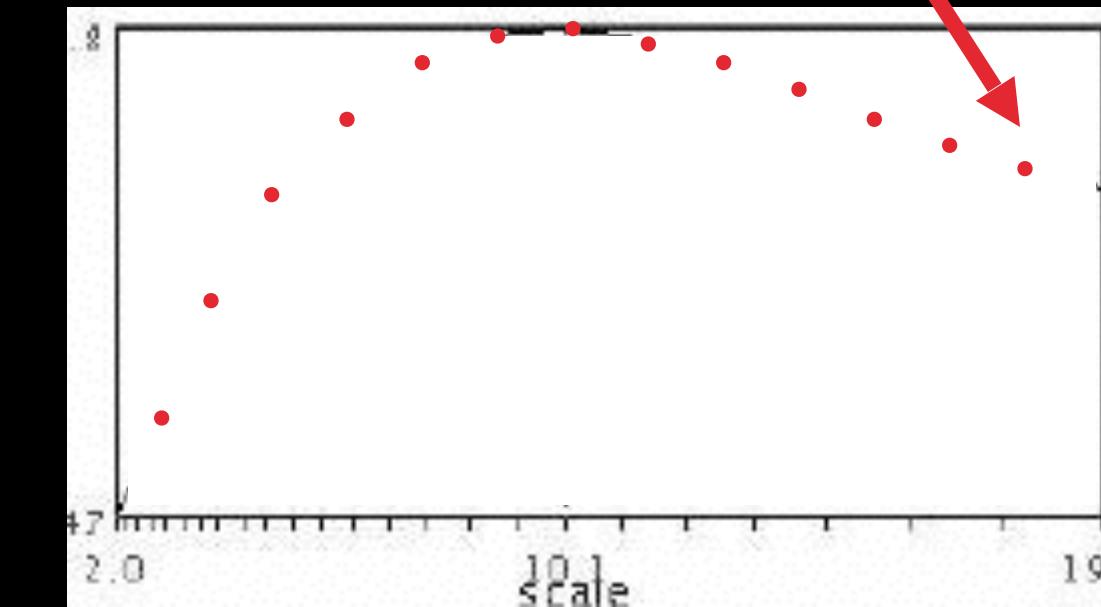


$$f(I_{i_1 \oplus i_m}(x', \sigma'))$$

AUTOMATIC SCALE SELECTION



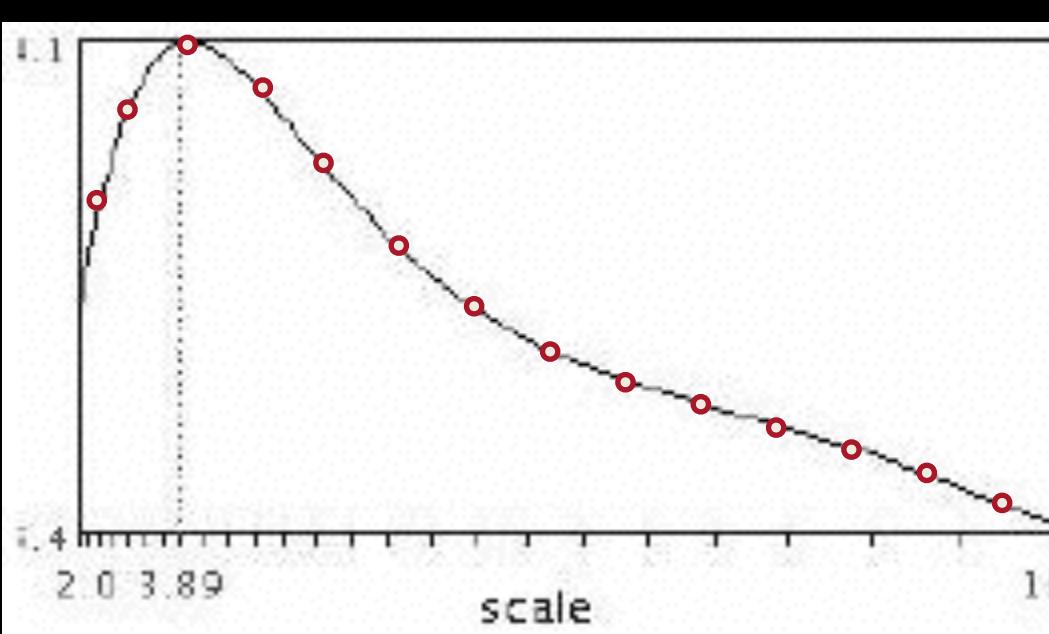
$$f(I_{i_1 \square j_m}(x, \sigma))$$



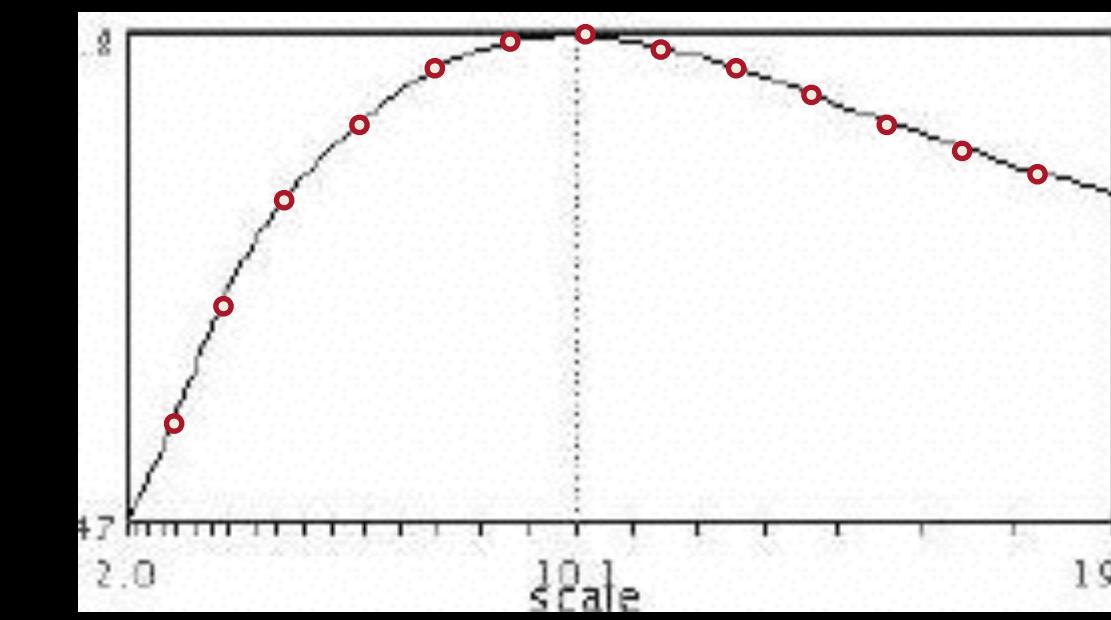
$$f(I_{i_1 \square j_m}(x', \sigma'))$$

AUTOMATIC SCALE SELECTION

- ▶ Function responses for increasing scale (scale signature)



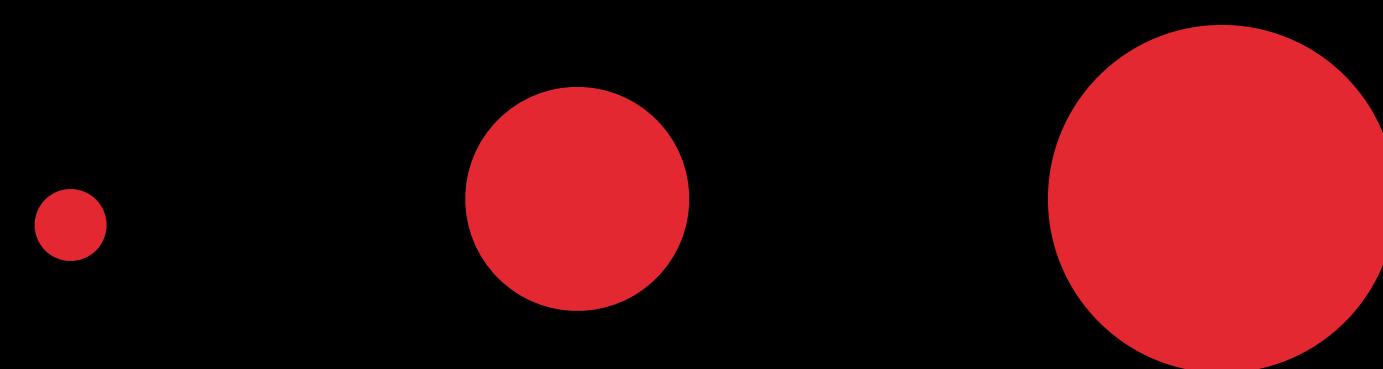
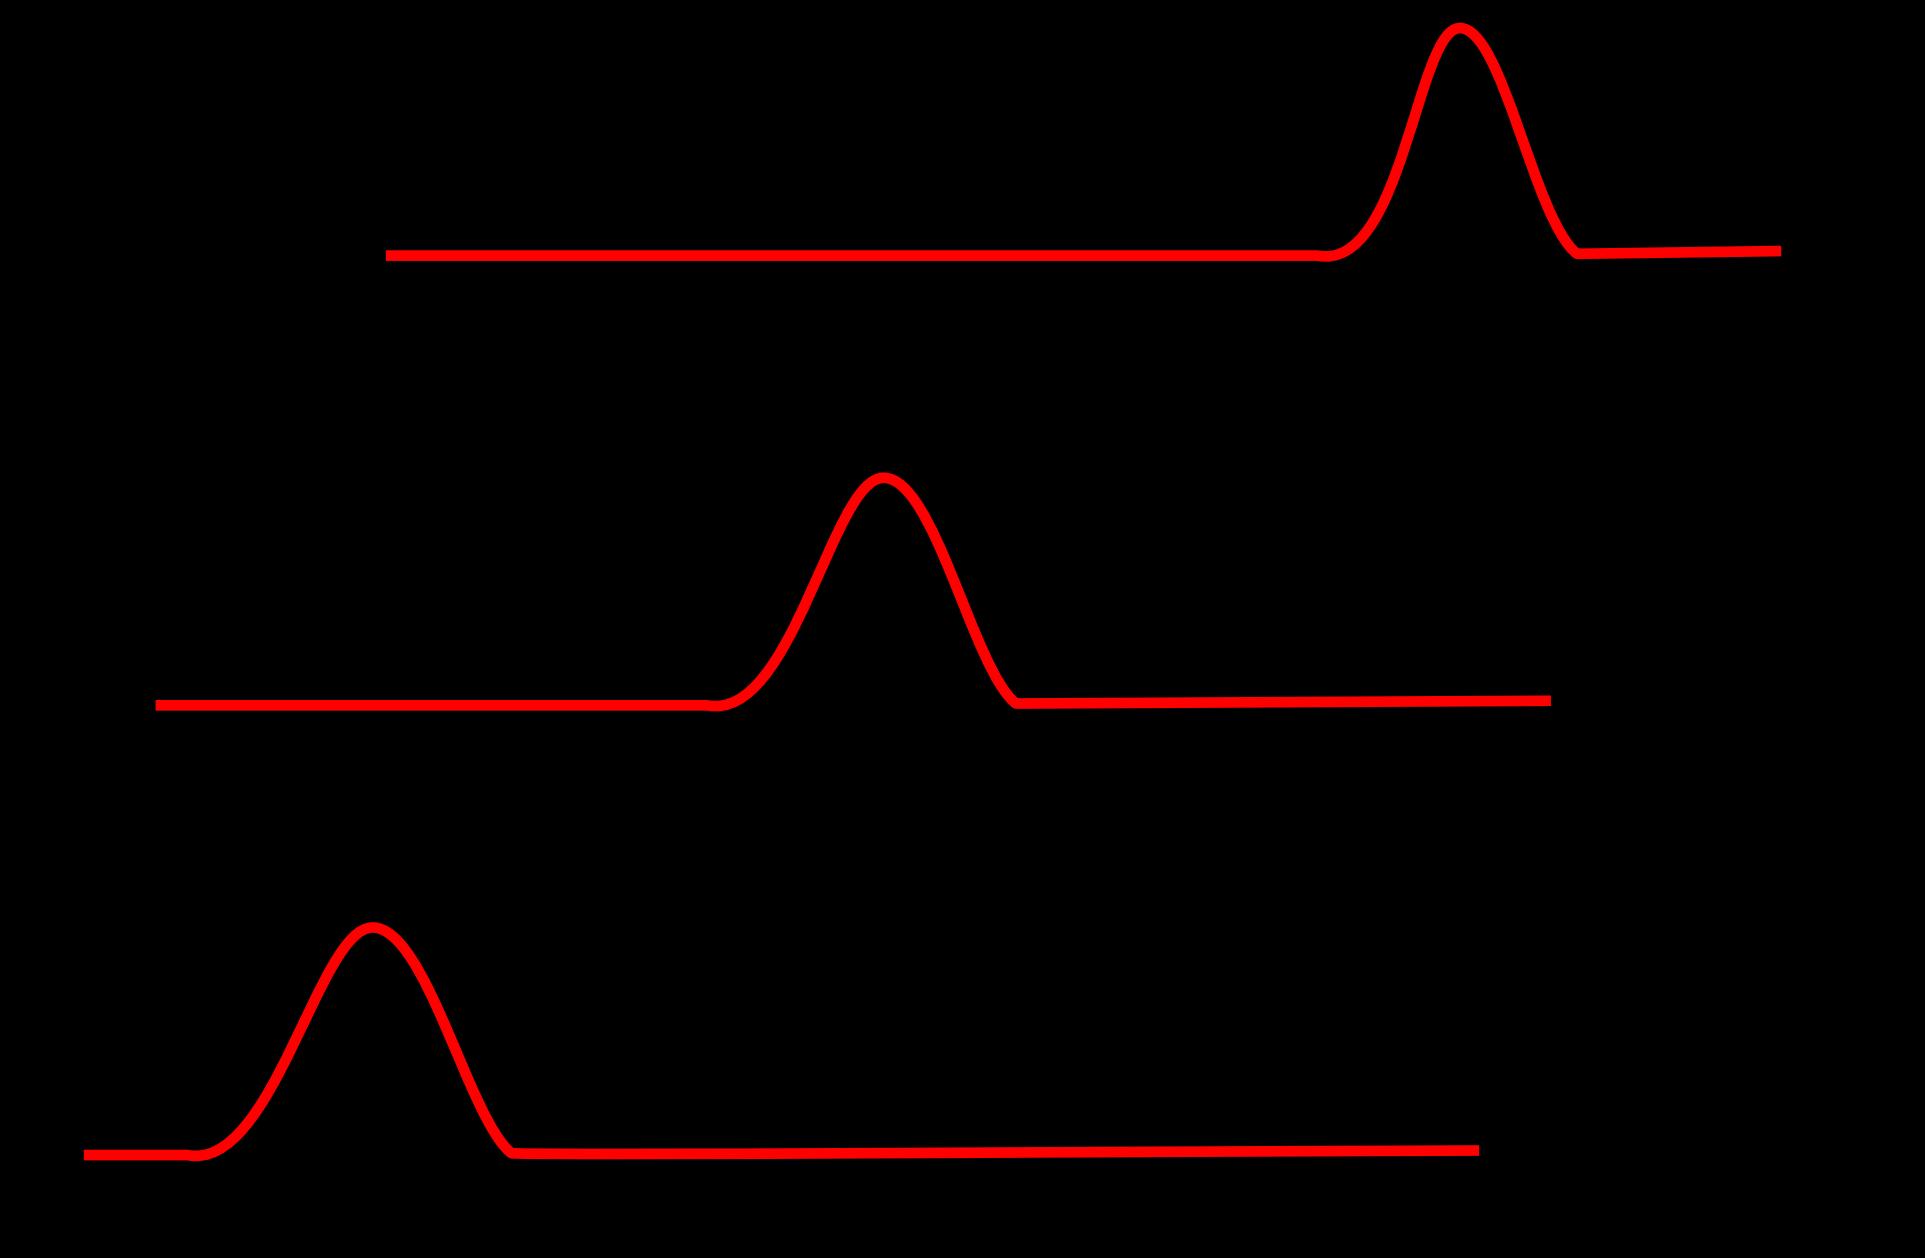
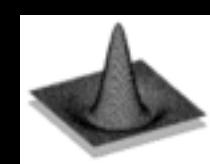
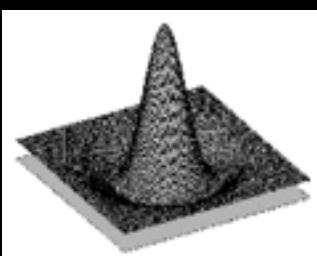
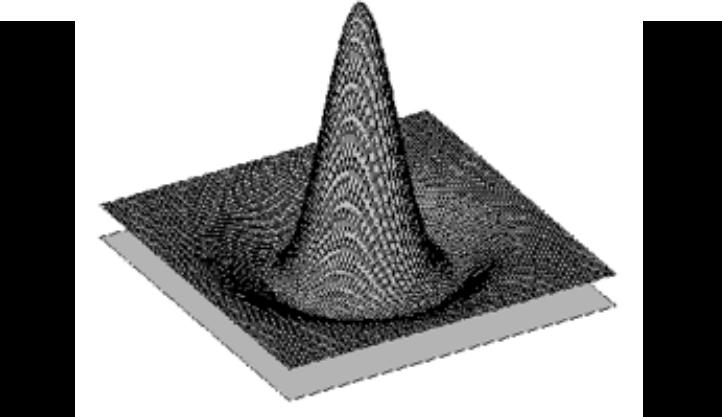
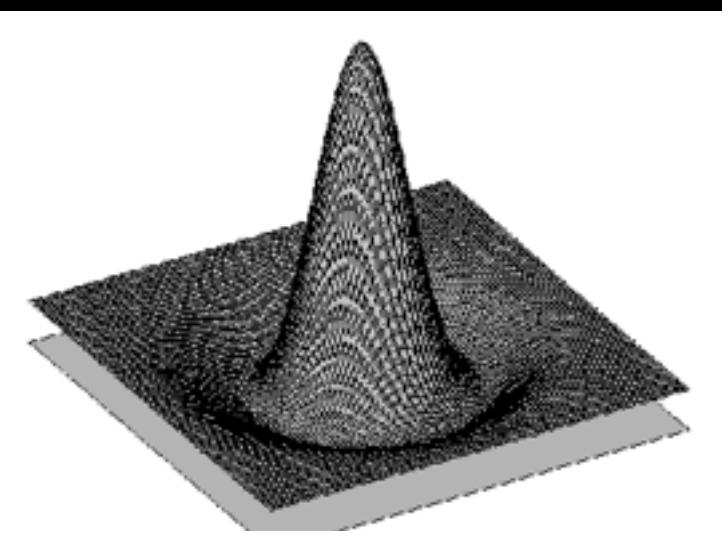
$$f(I_{i_1 \square i_m}(x, \sigma))$$



$$f(I_{i_1 \square i_m}(x', \sigma'))$$

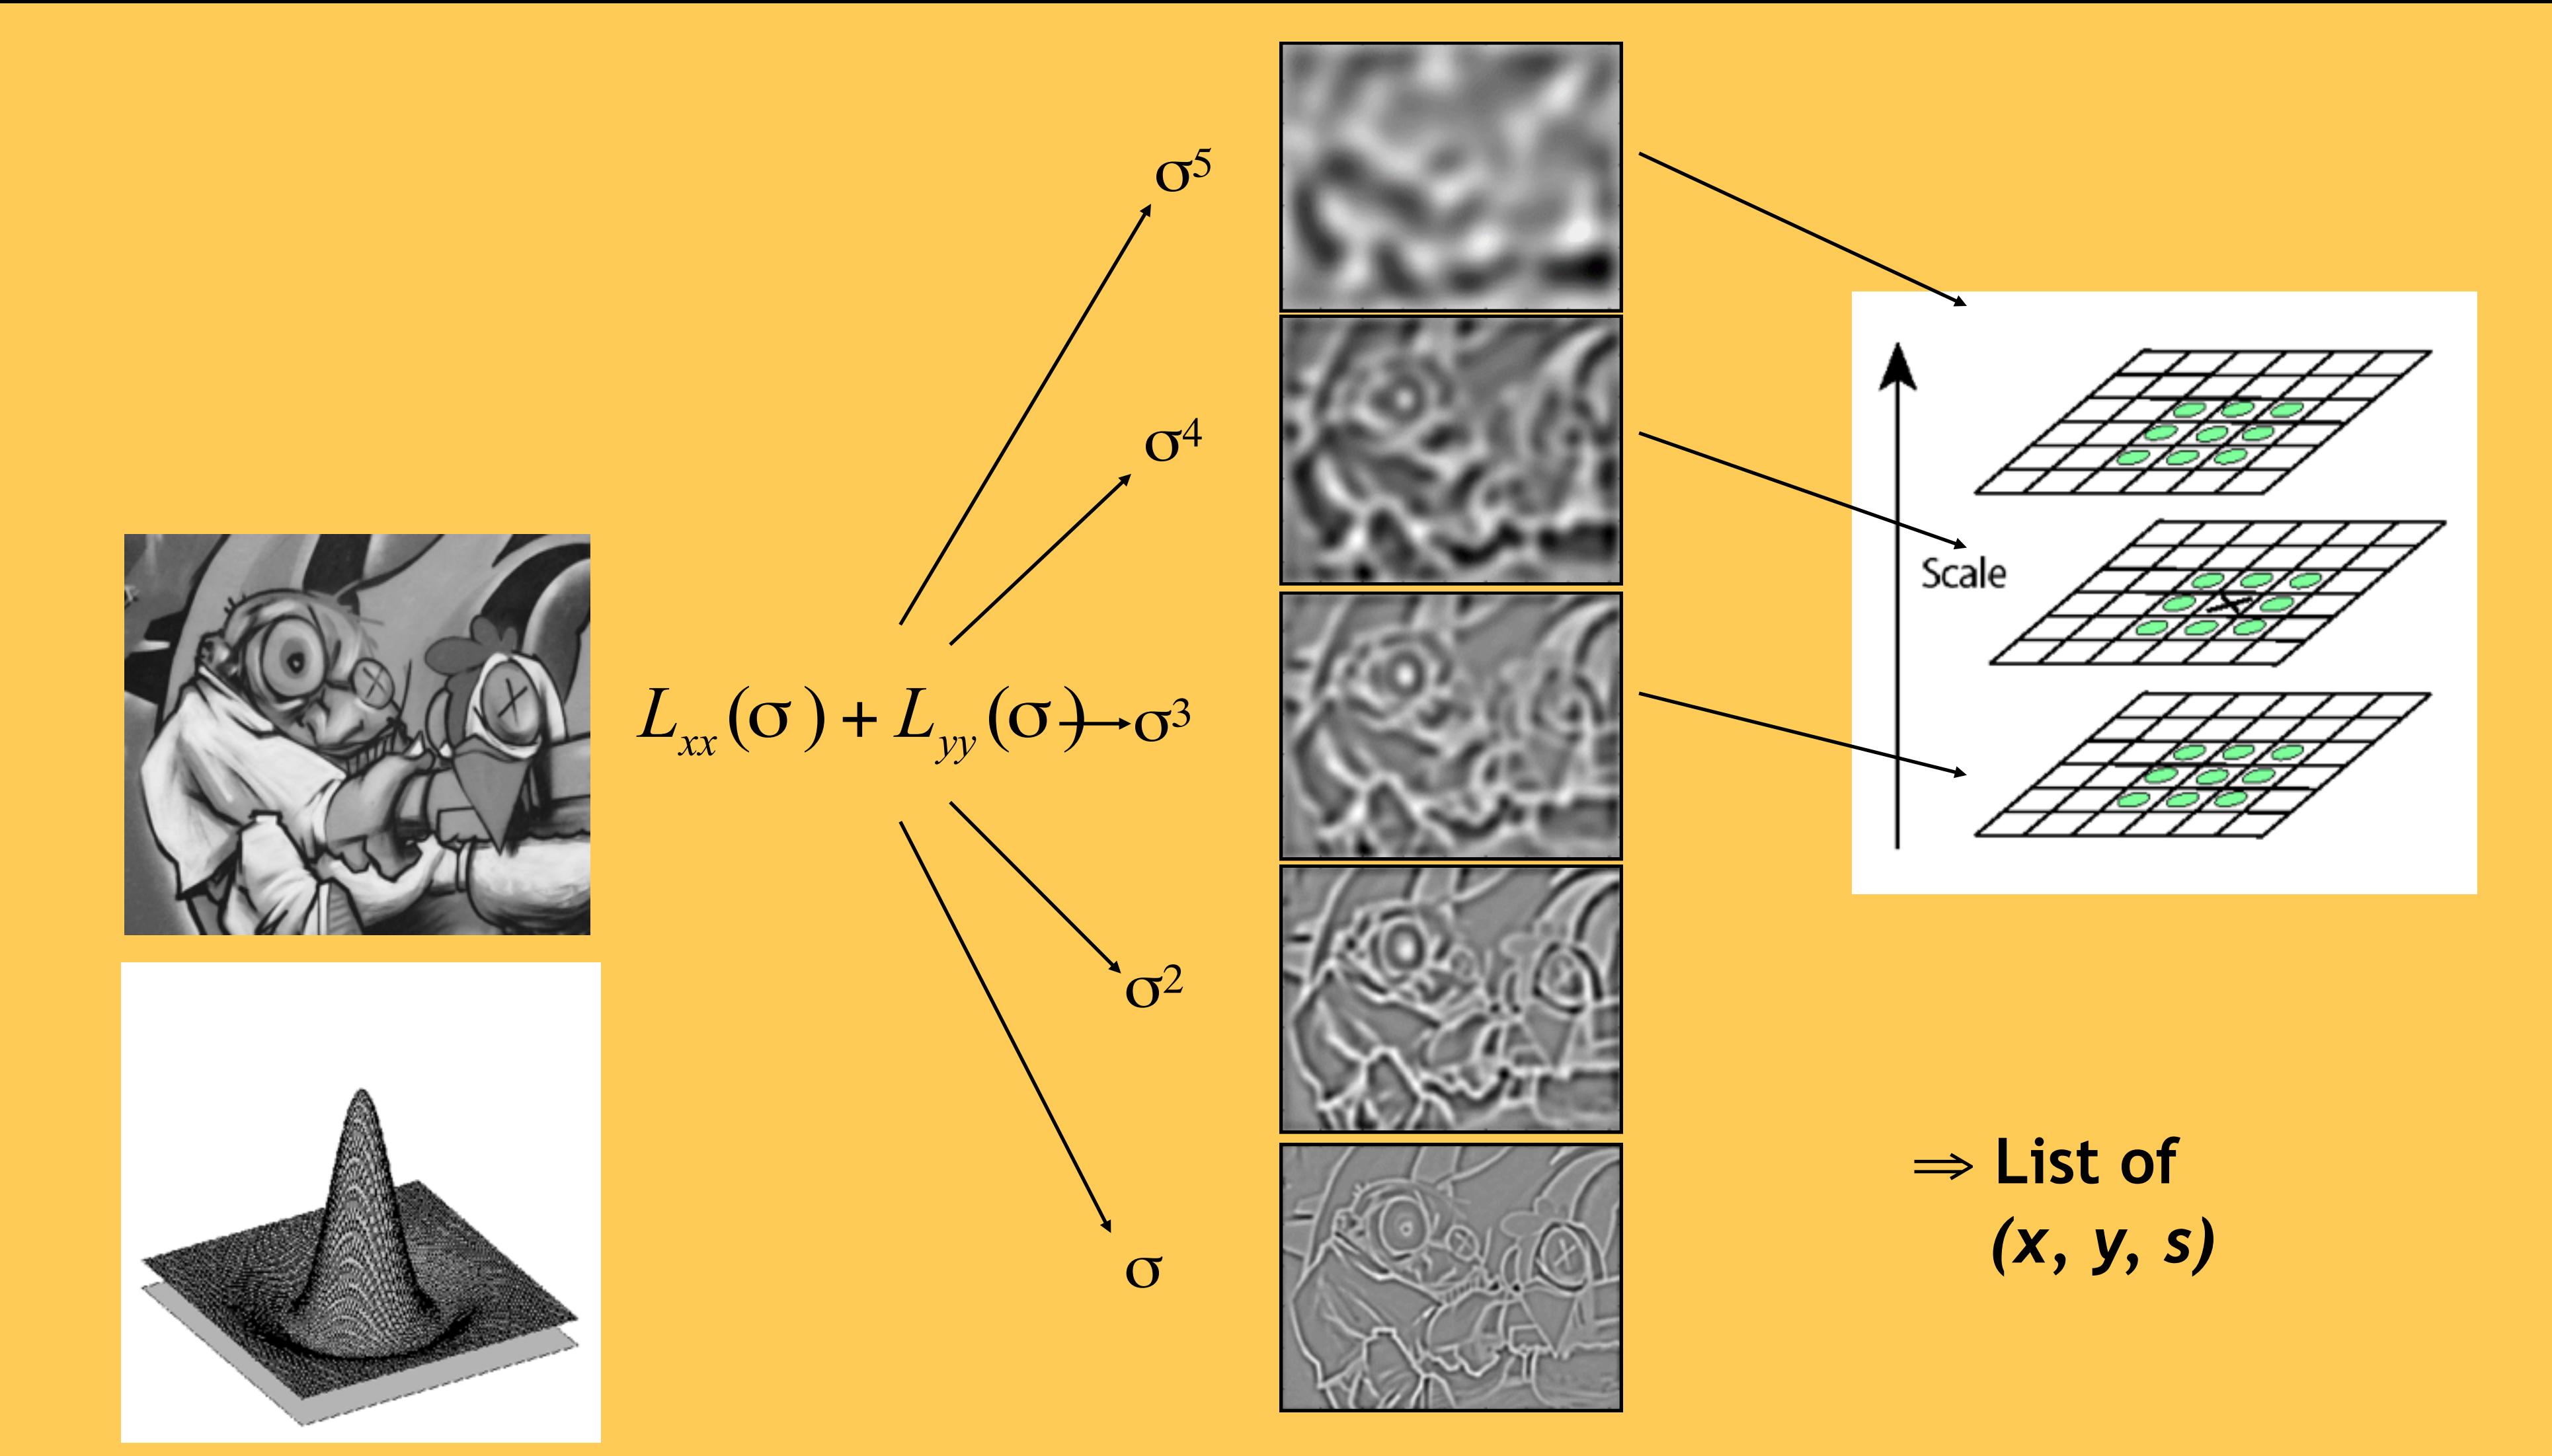
WHAT IS A USEFUL SIGNATURE FUNCTION?

- ▶ Laplacian-of-Gaussian = “blob” detector



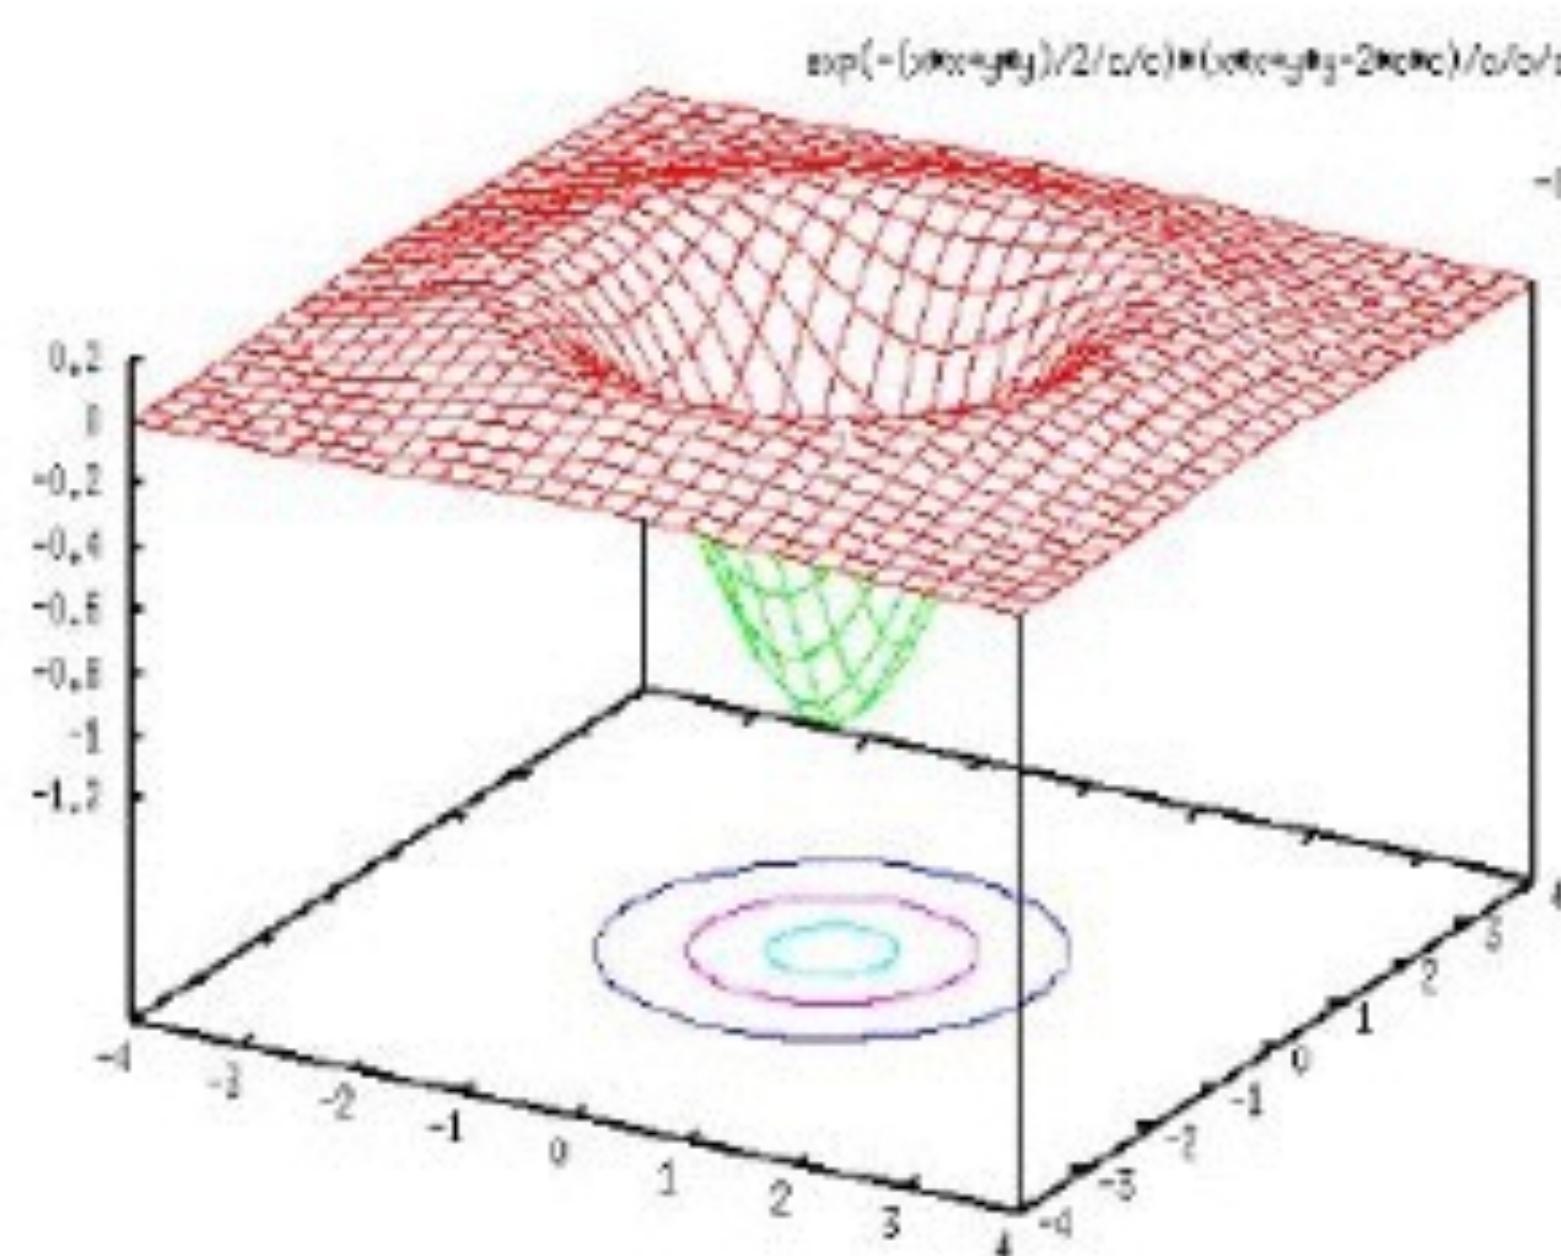
TEXT

FIND LOCAL MAXIMA IN POSITION-SCALE SPACE

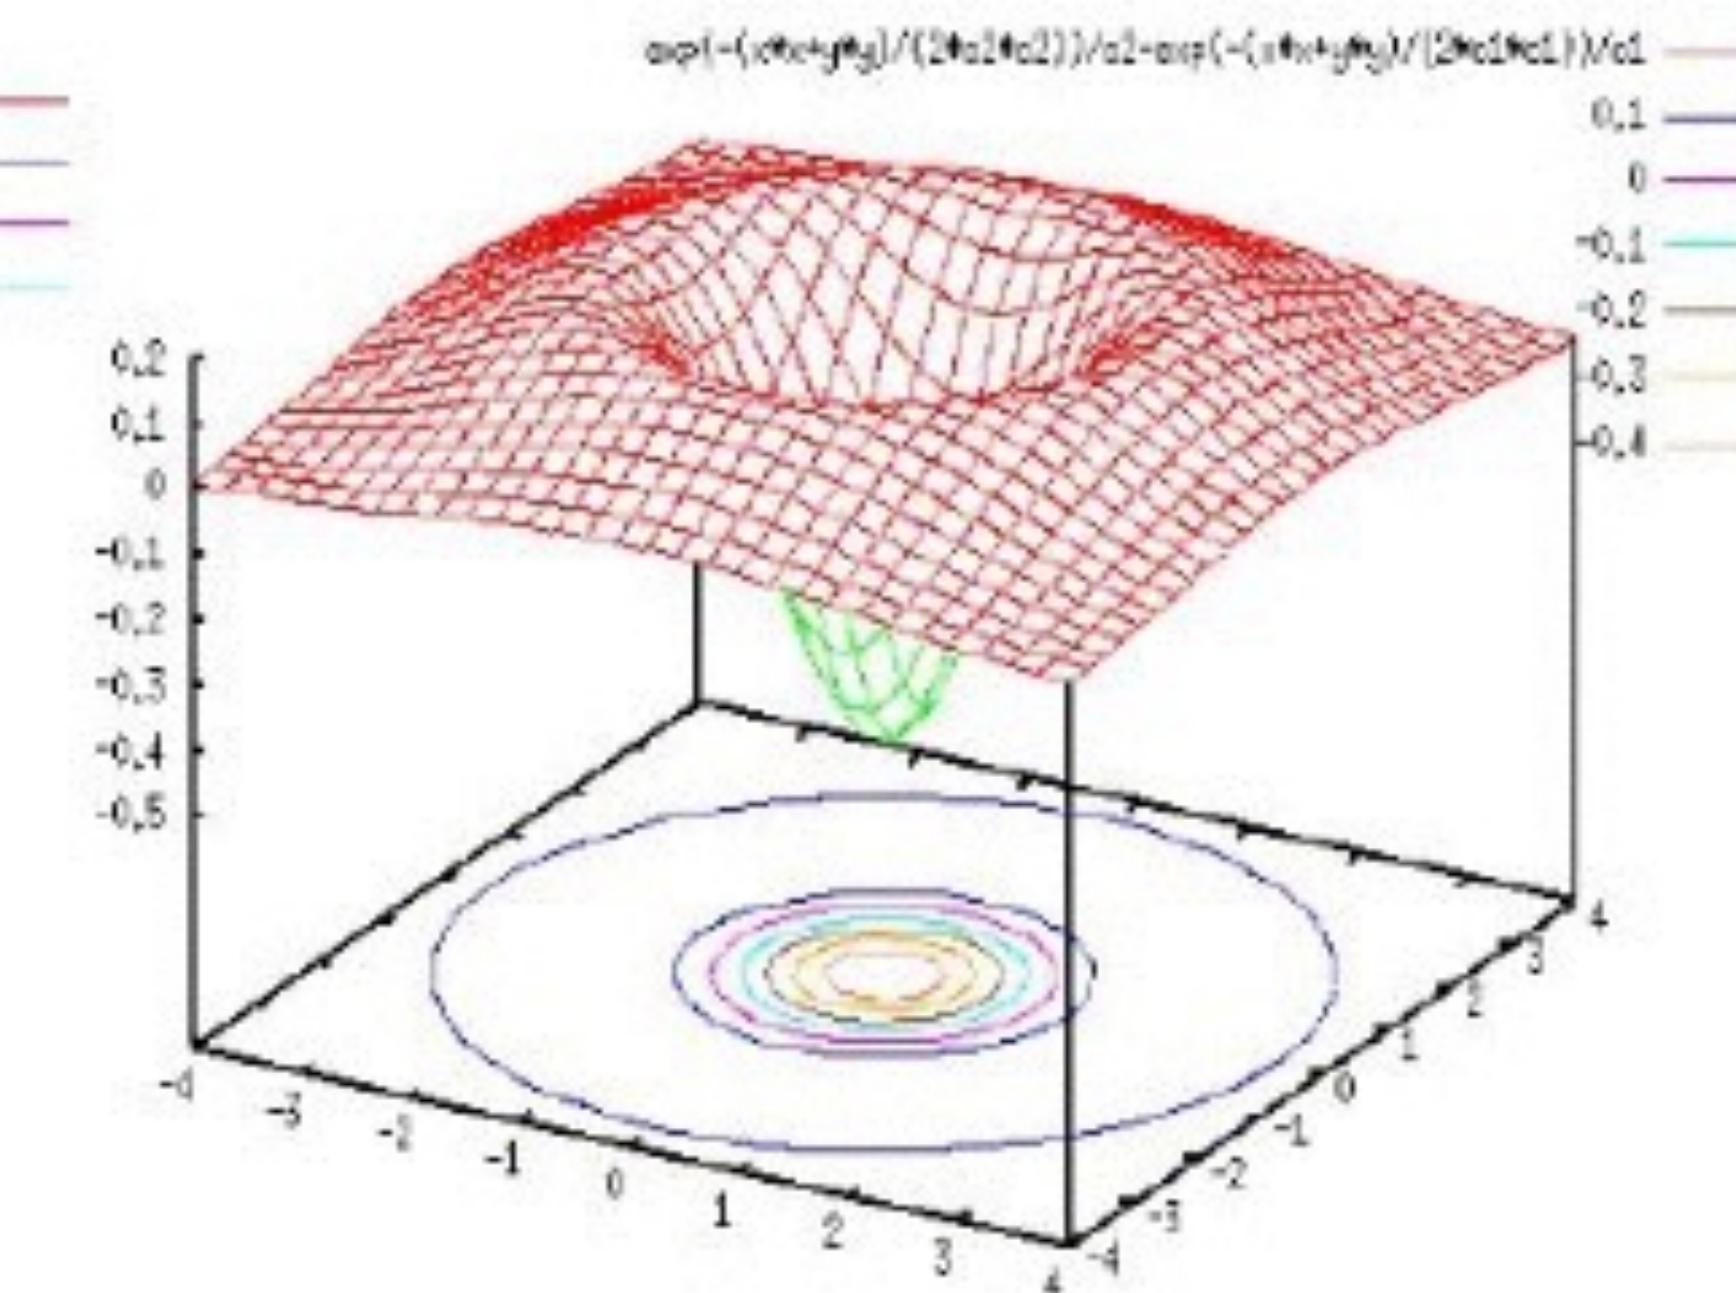


DIFFERENCE-OF-GAUSSIAN (DOG)

Laplacian of Gaussian

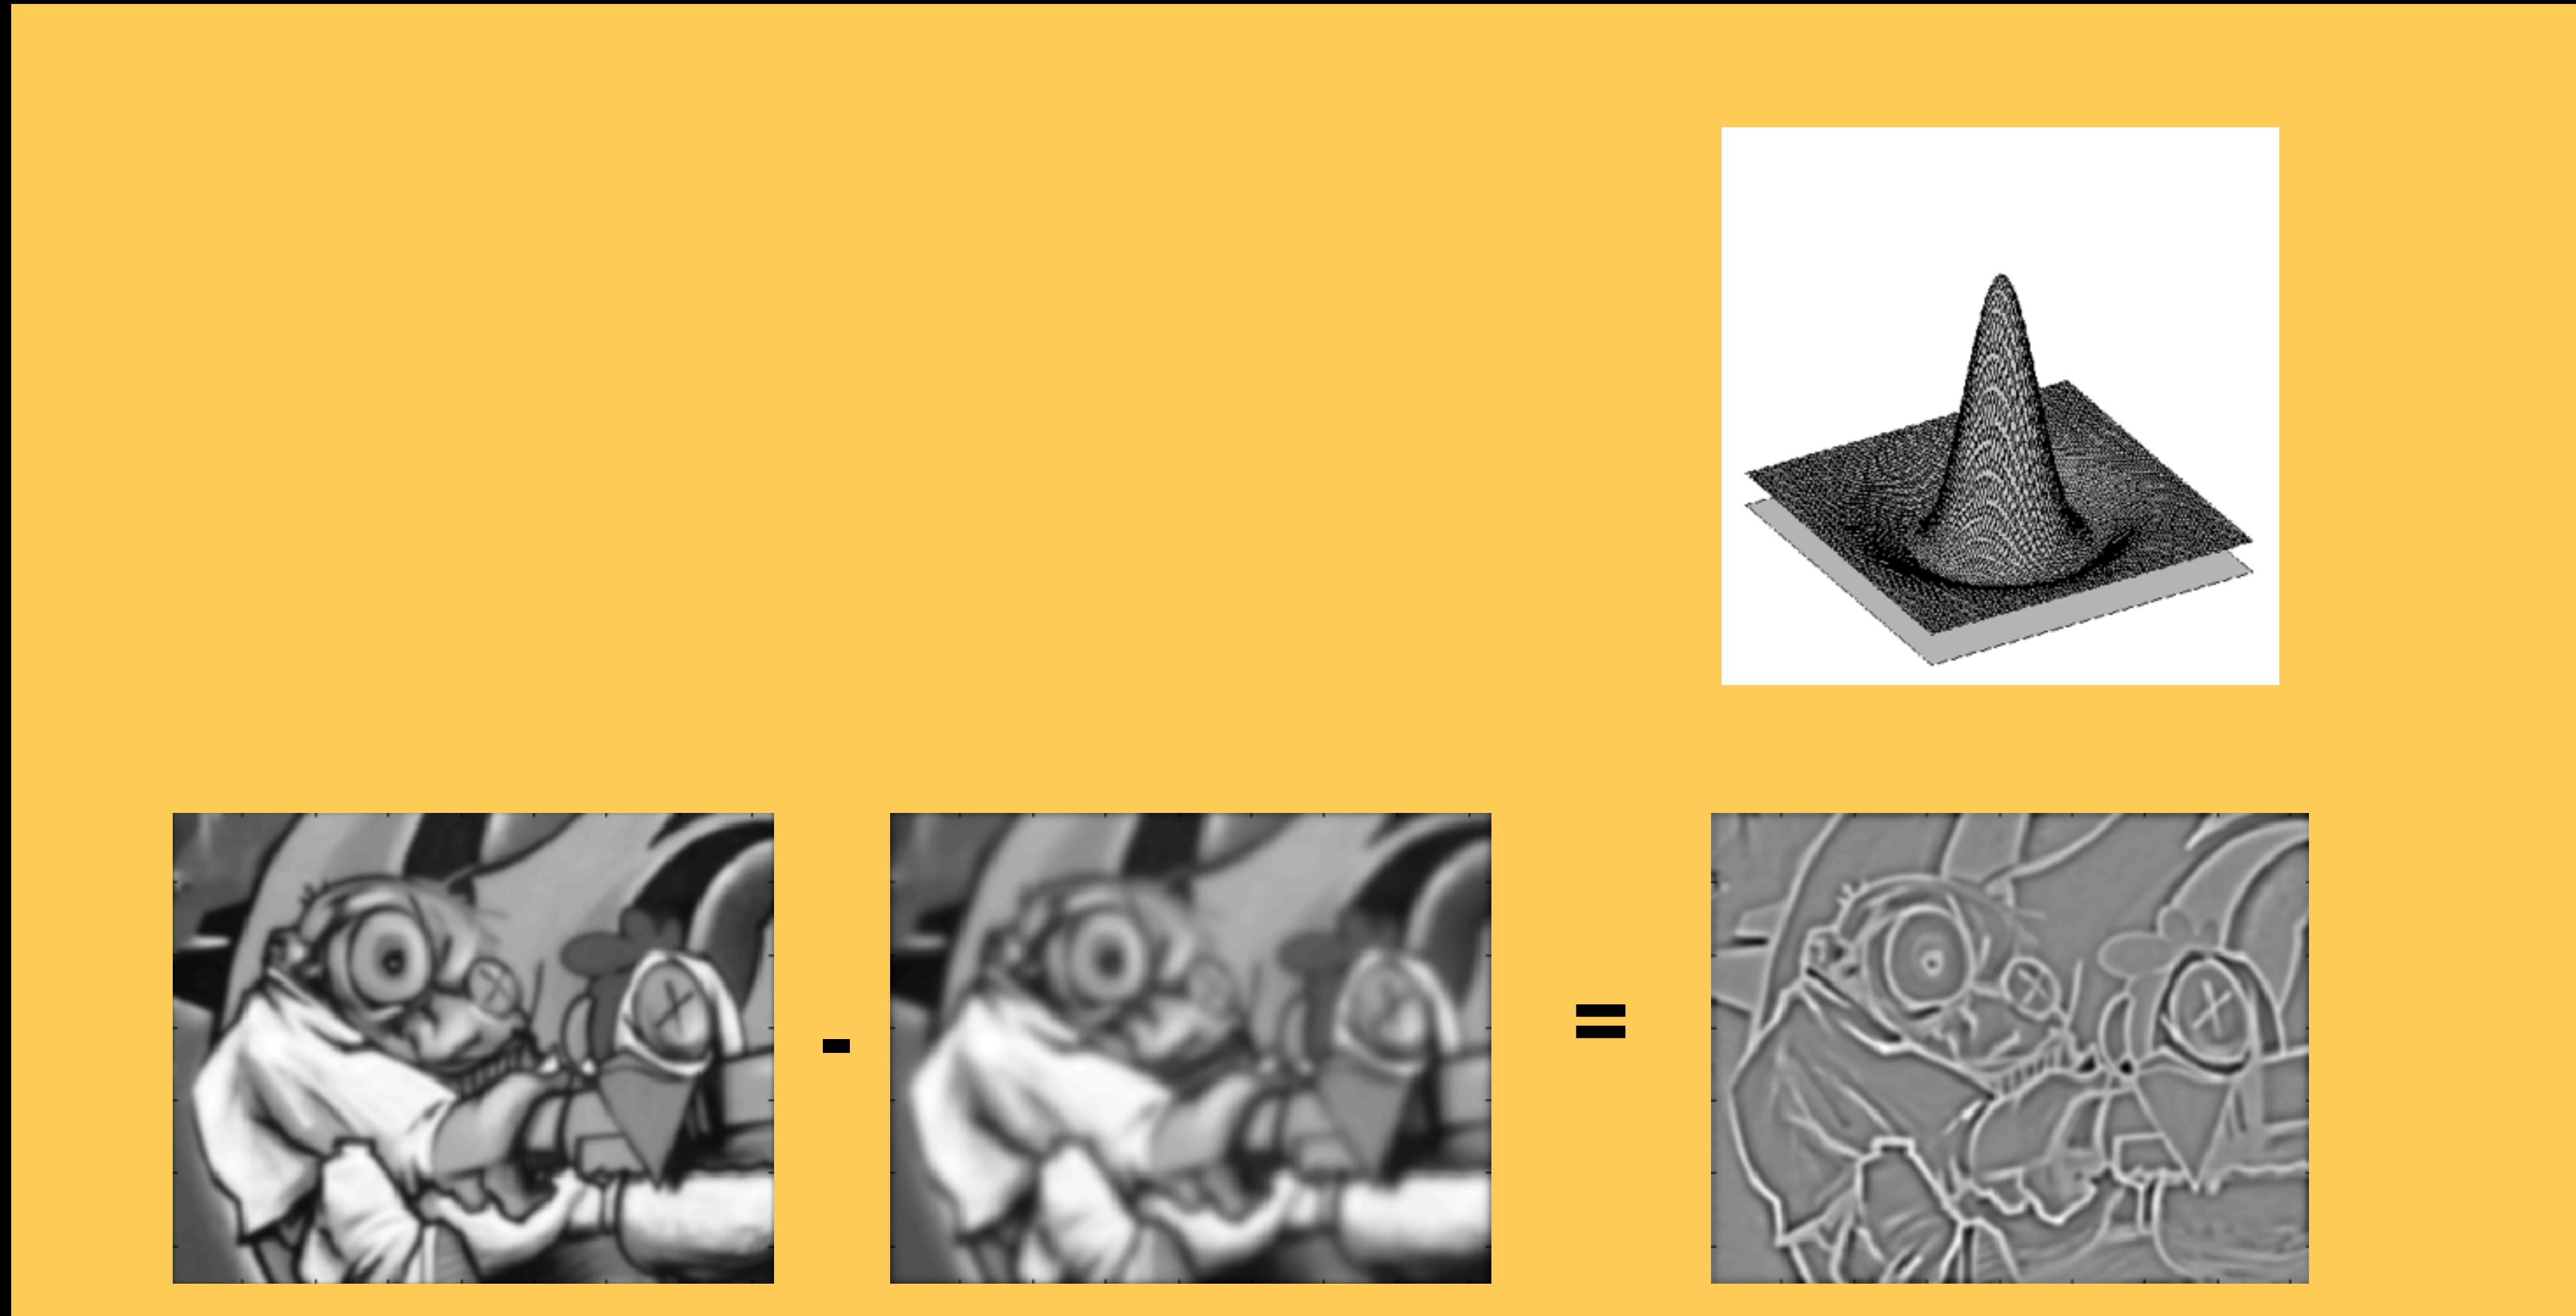


Difference of Gaussians



TEXT

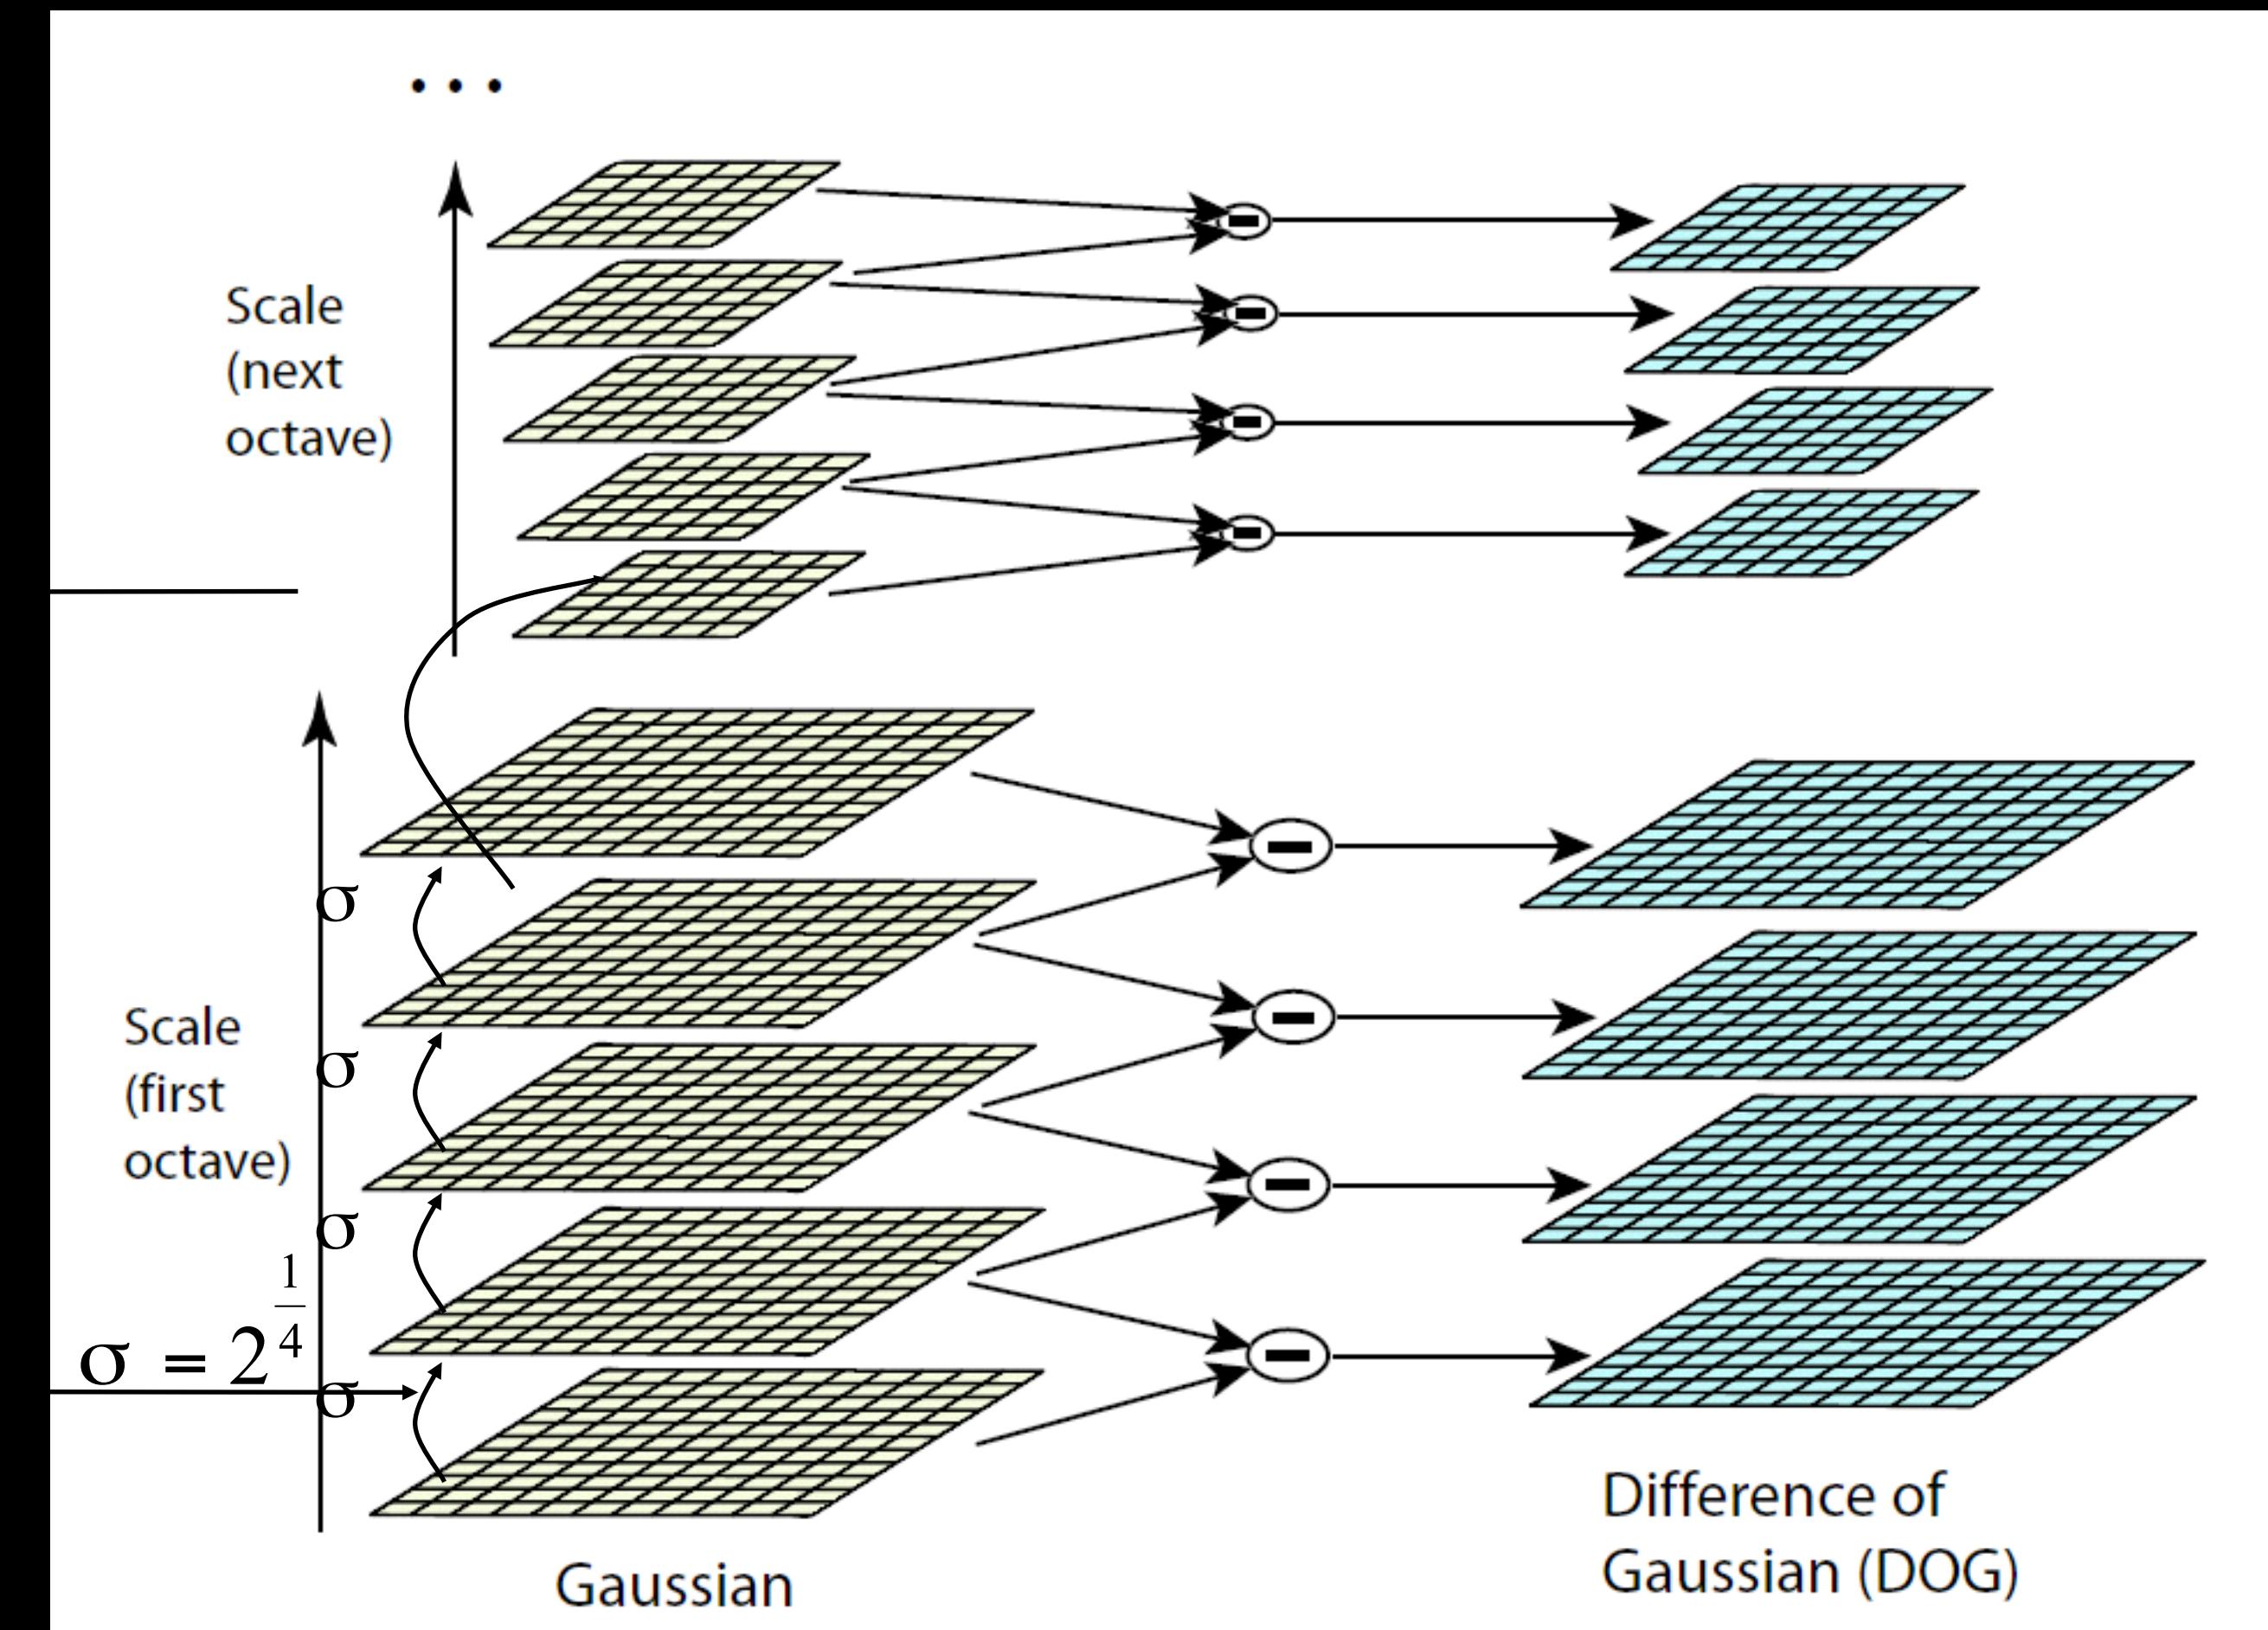
DIFFERENCE-OF-GAUSSIAN (DOG)



K. Grauman, B. Leibe

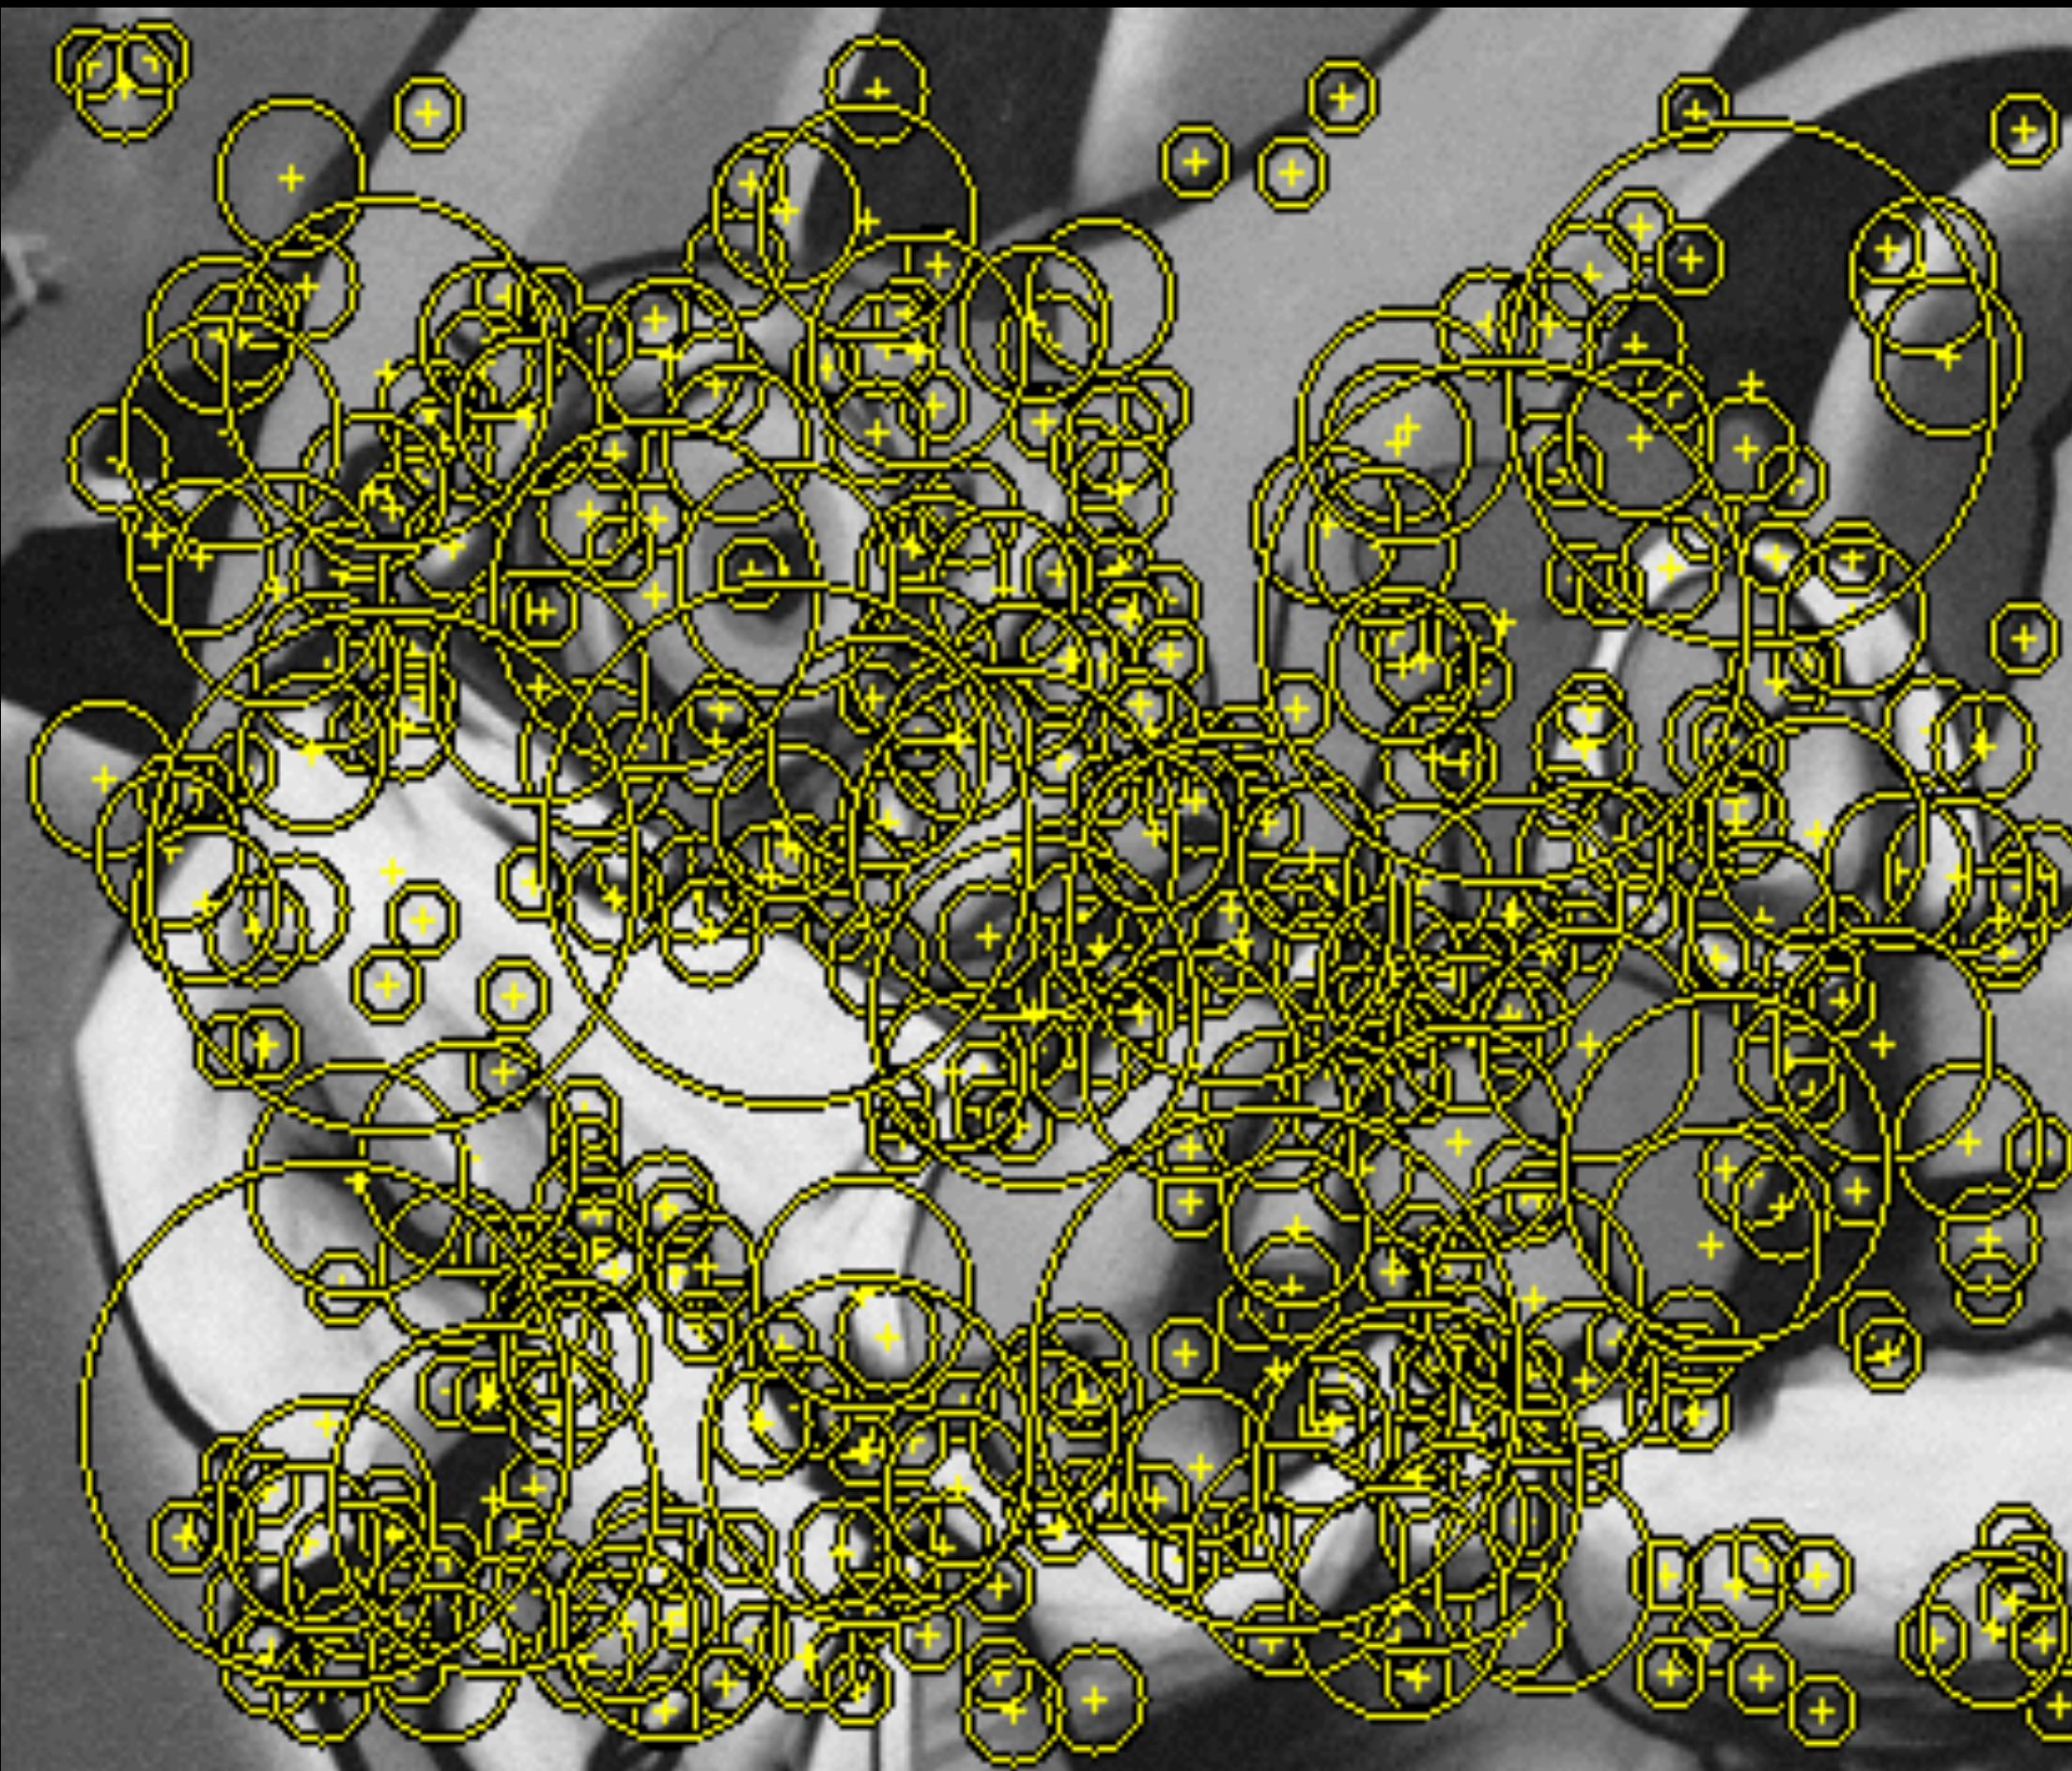
DOG – EFFICIENT COMPUTATION

▶ Computation in Gaussian scale pyramid



TEXT

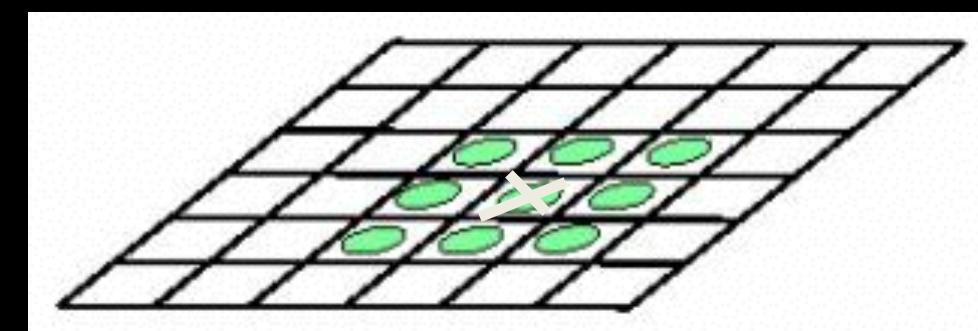
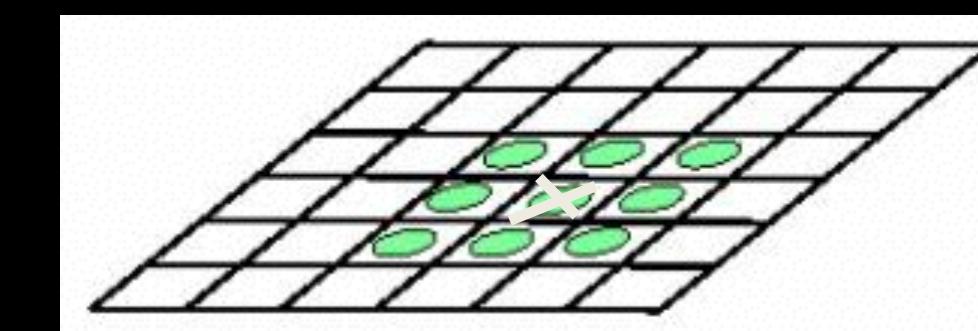
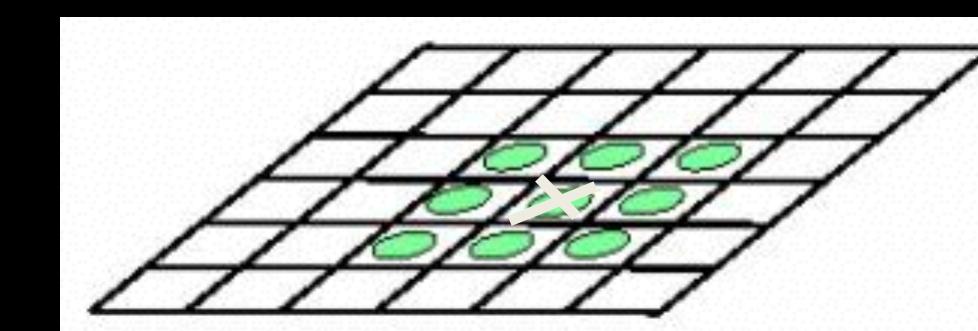
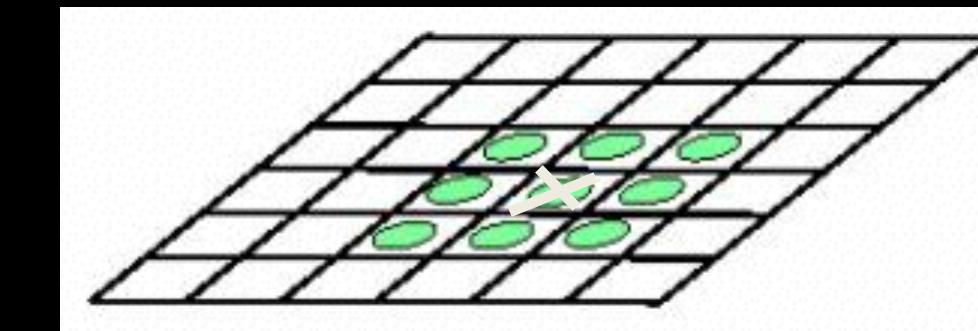
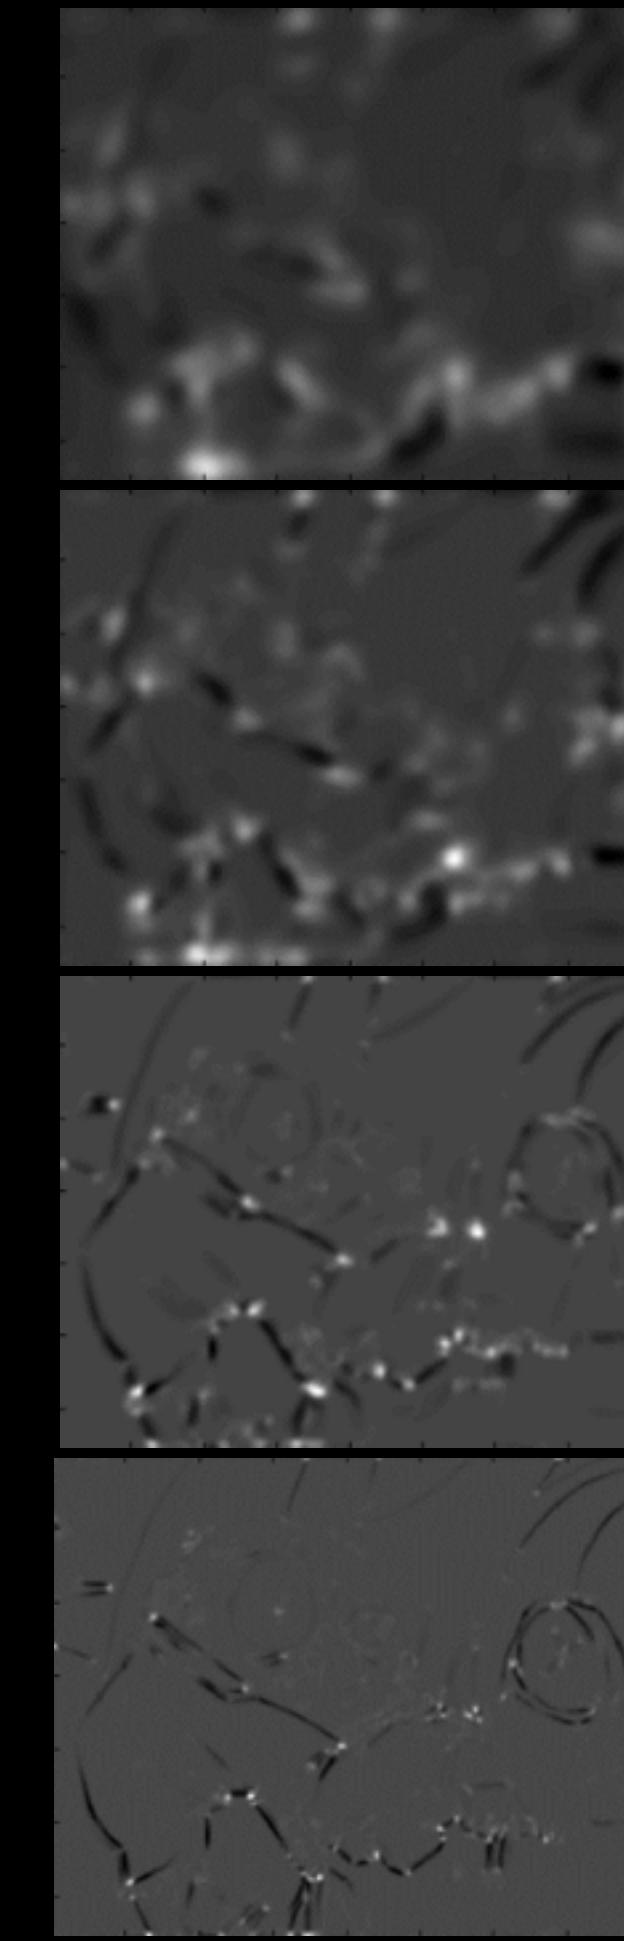
RESULTS: DIFFERENCE-OF-GAUSSIAN



TEXT

HARRIS-LAPLACE [MIKOLAJCZYK '01]

1. Initialization: Multiscale Harris corner detection

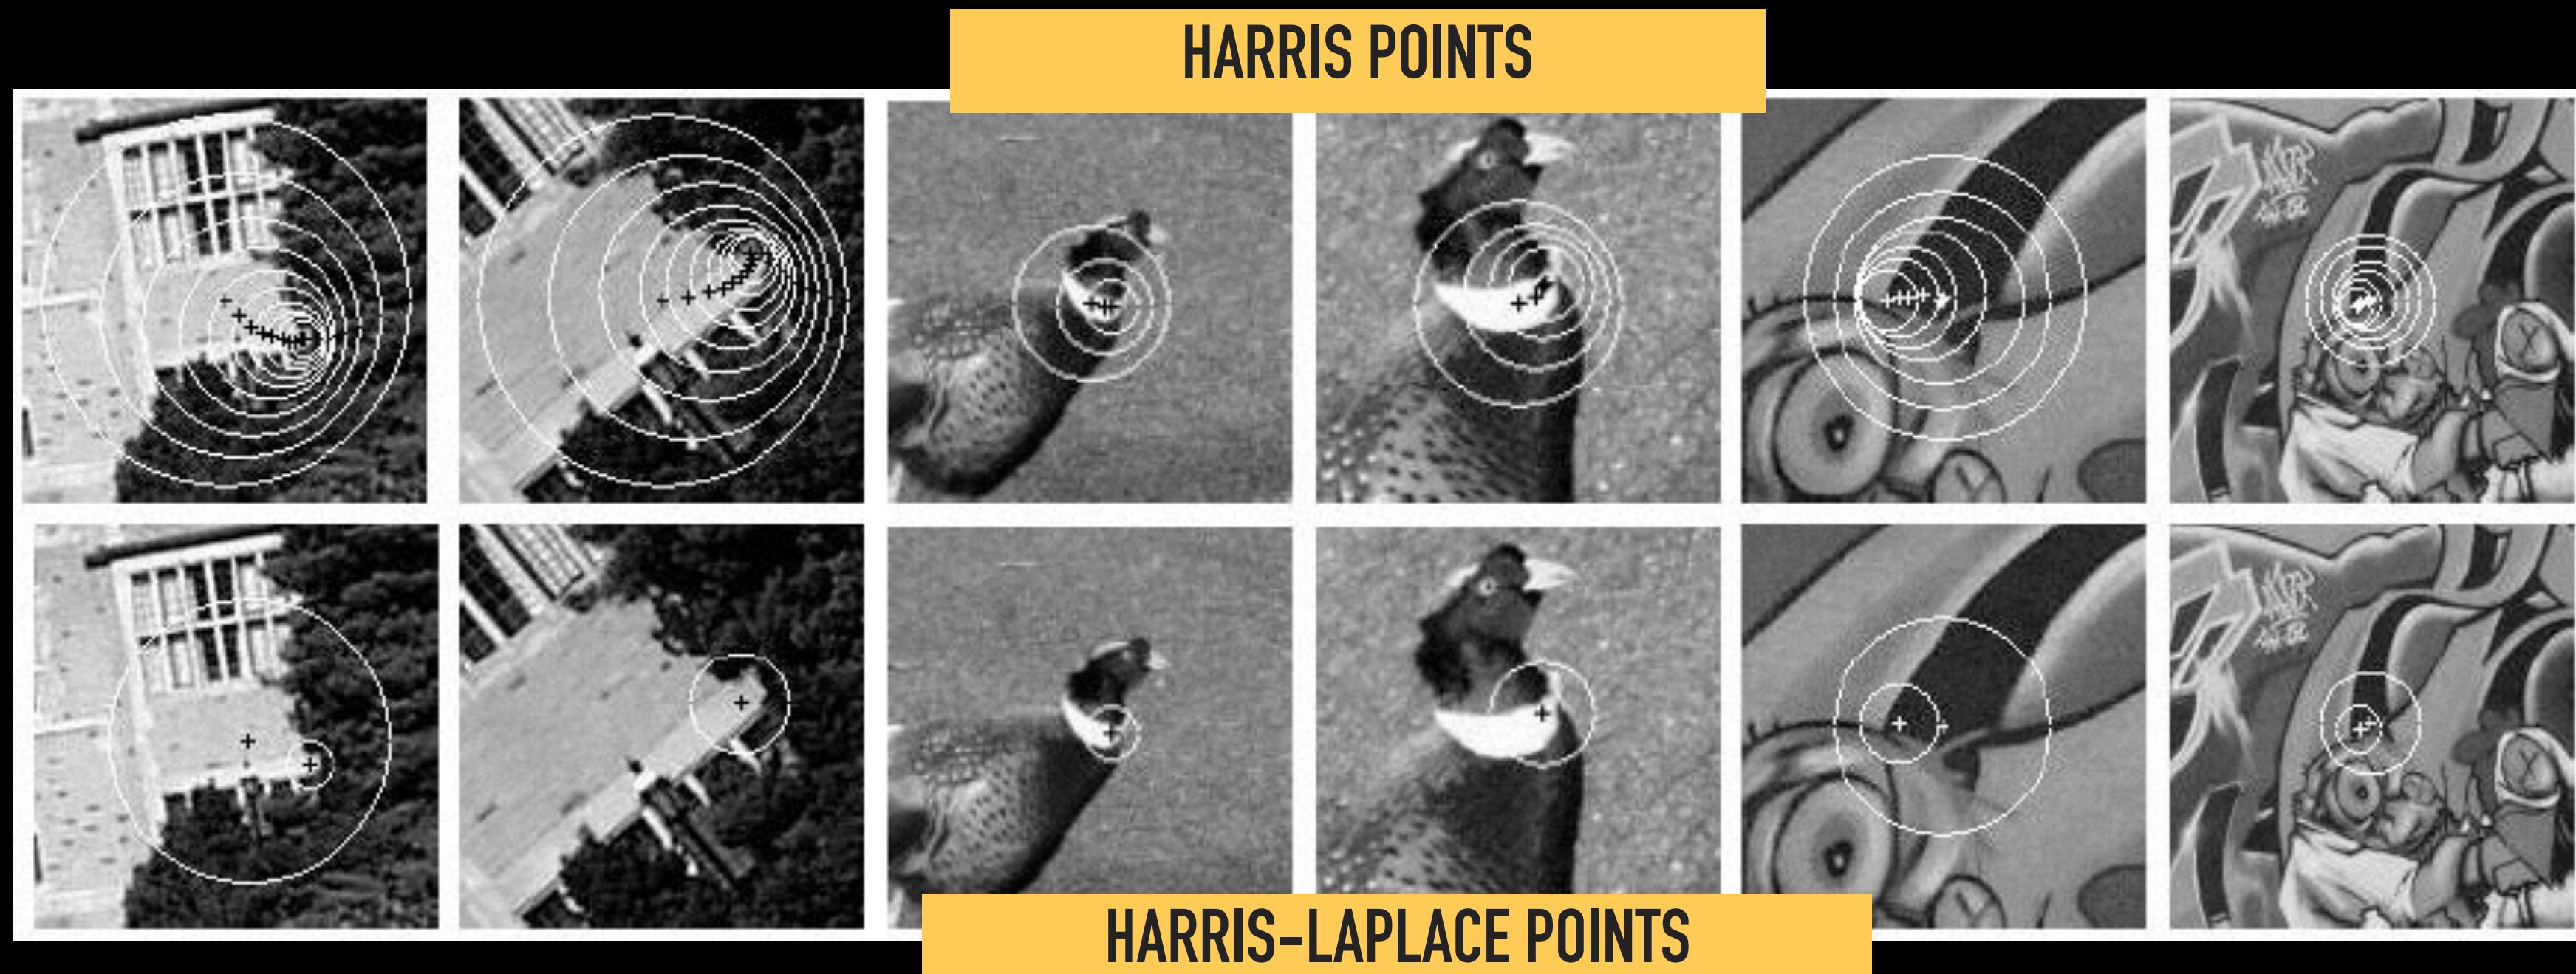


COMPUTING HARRIS FUNCTION

DETECTING LOCAL MAXIMA

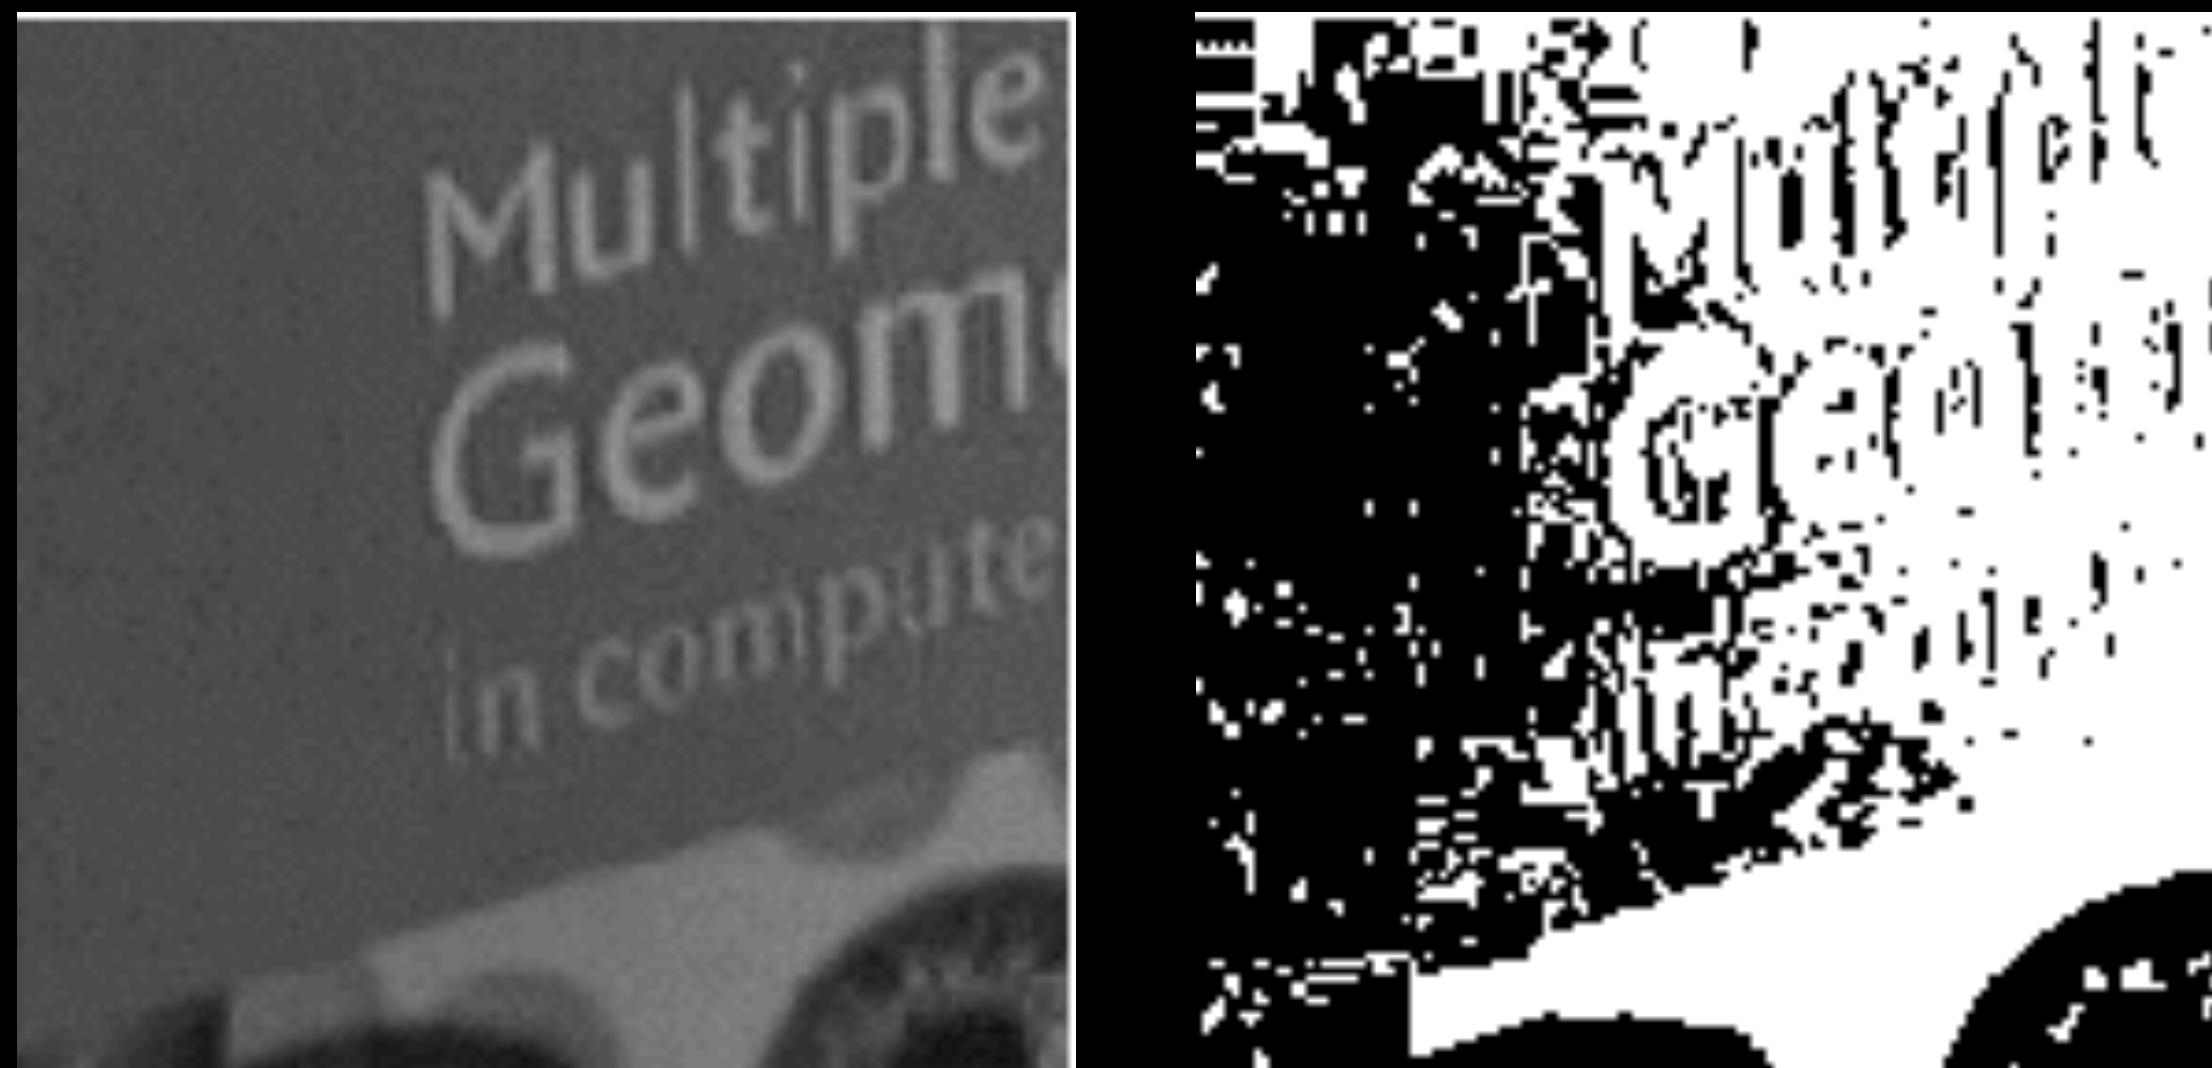
HARRIS-LAPLACE [MIKOŁAJCZYK '01]

1. Initialization: Multiscale Harris corner detection
2. Scale selection based on Laplacian
(same procedure with Hessian \Rightarrow Hessian-Laplace)

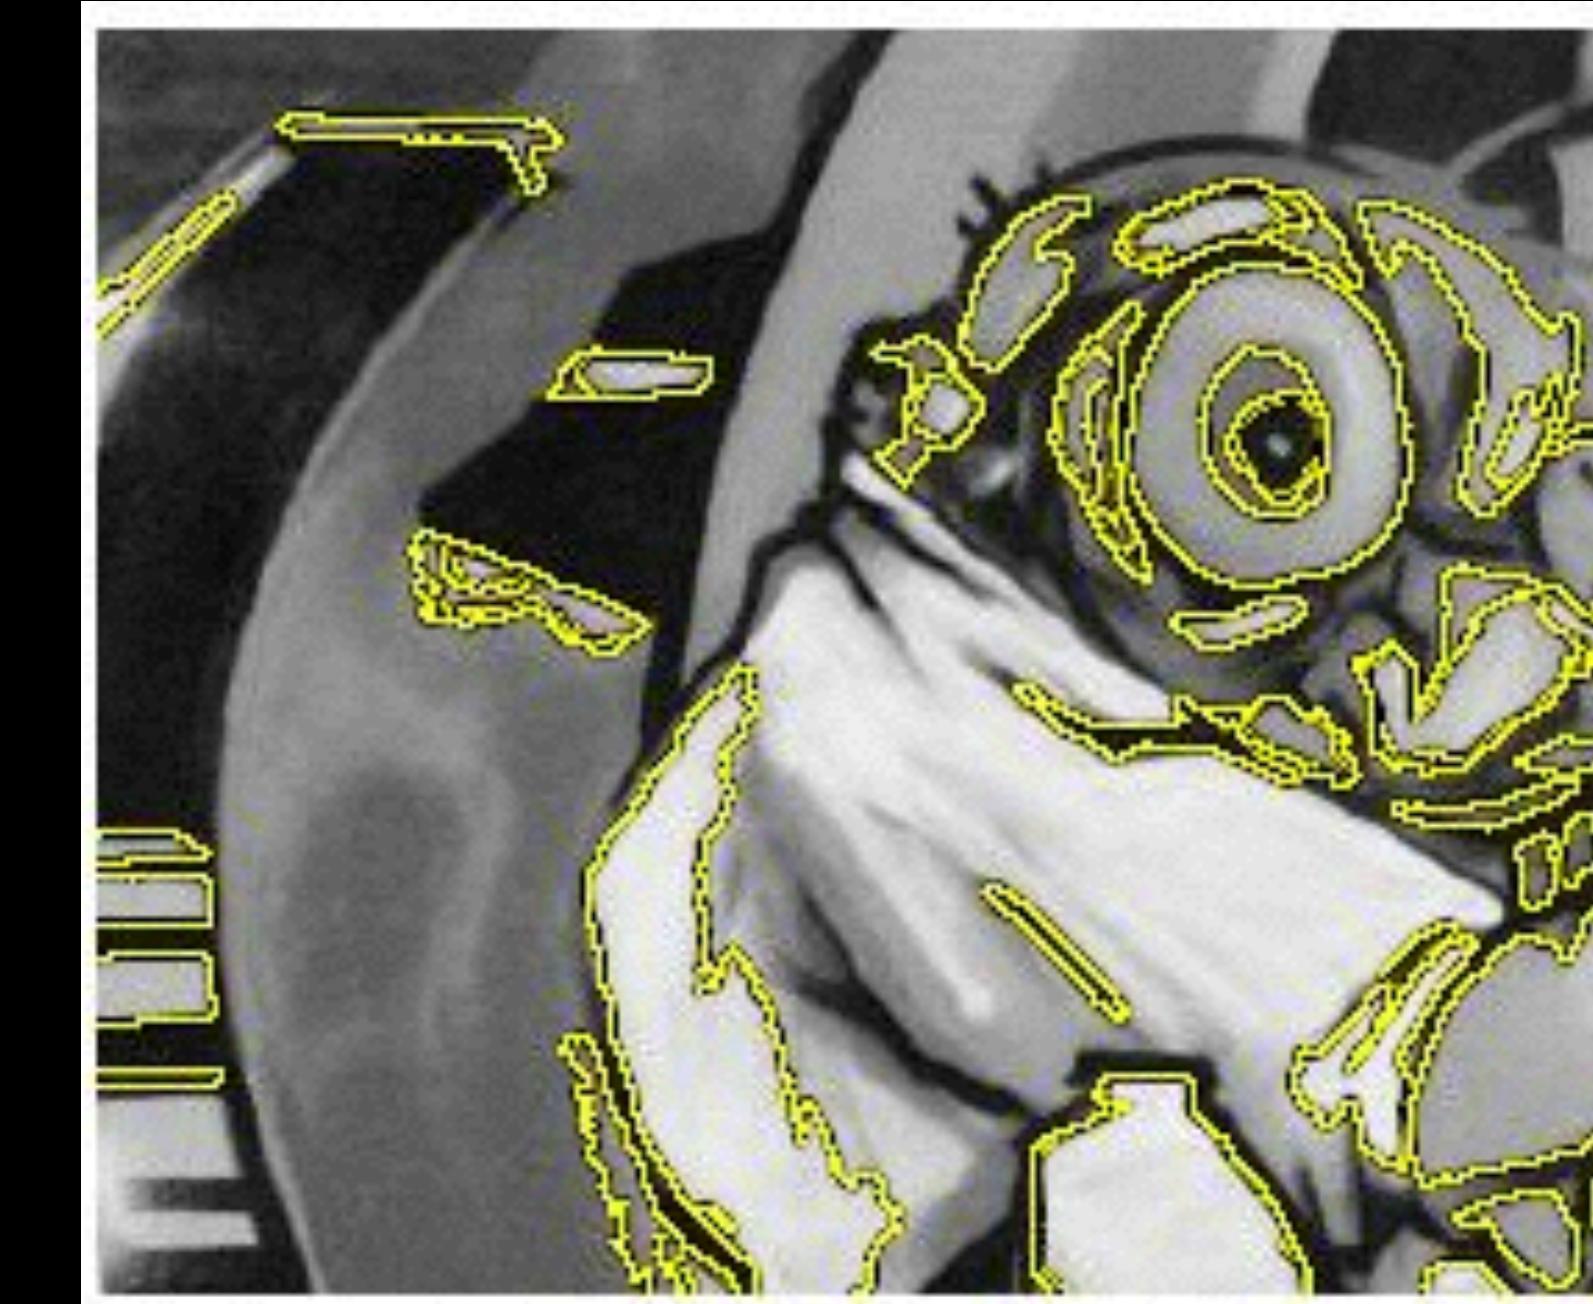
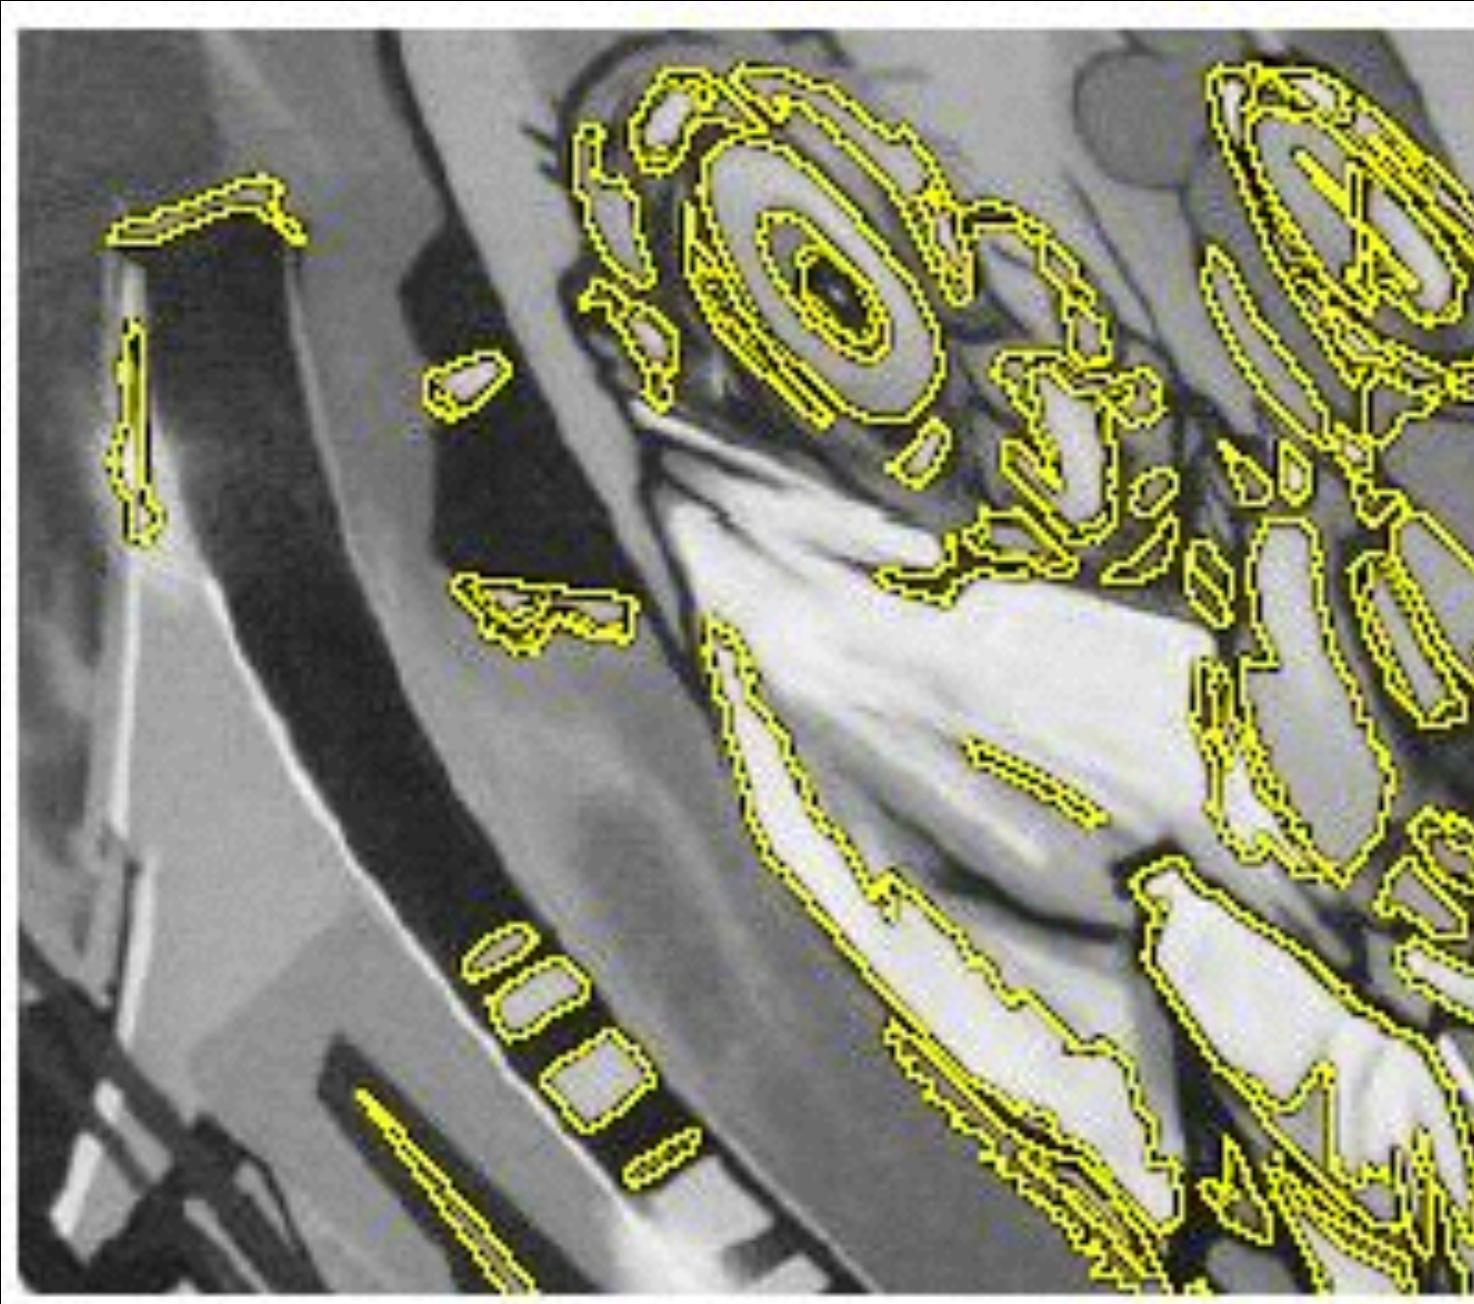


MAXIMALLY STABLE EXTREMAL REGIONS [MATAS '02]

- ▶ Based on Watershed segmentation algorithm
- ▶ Select regions that stay stable over a large parameter range



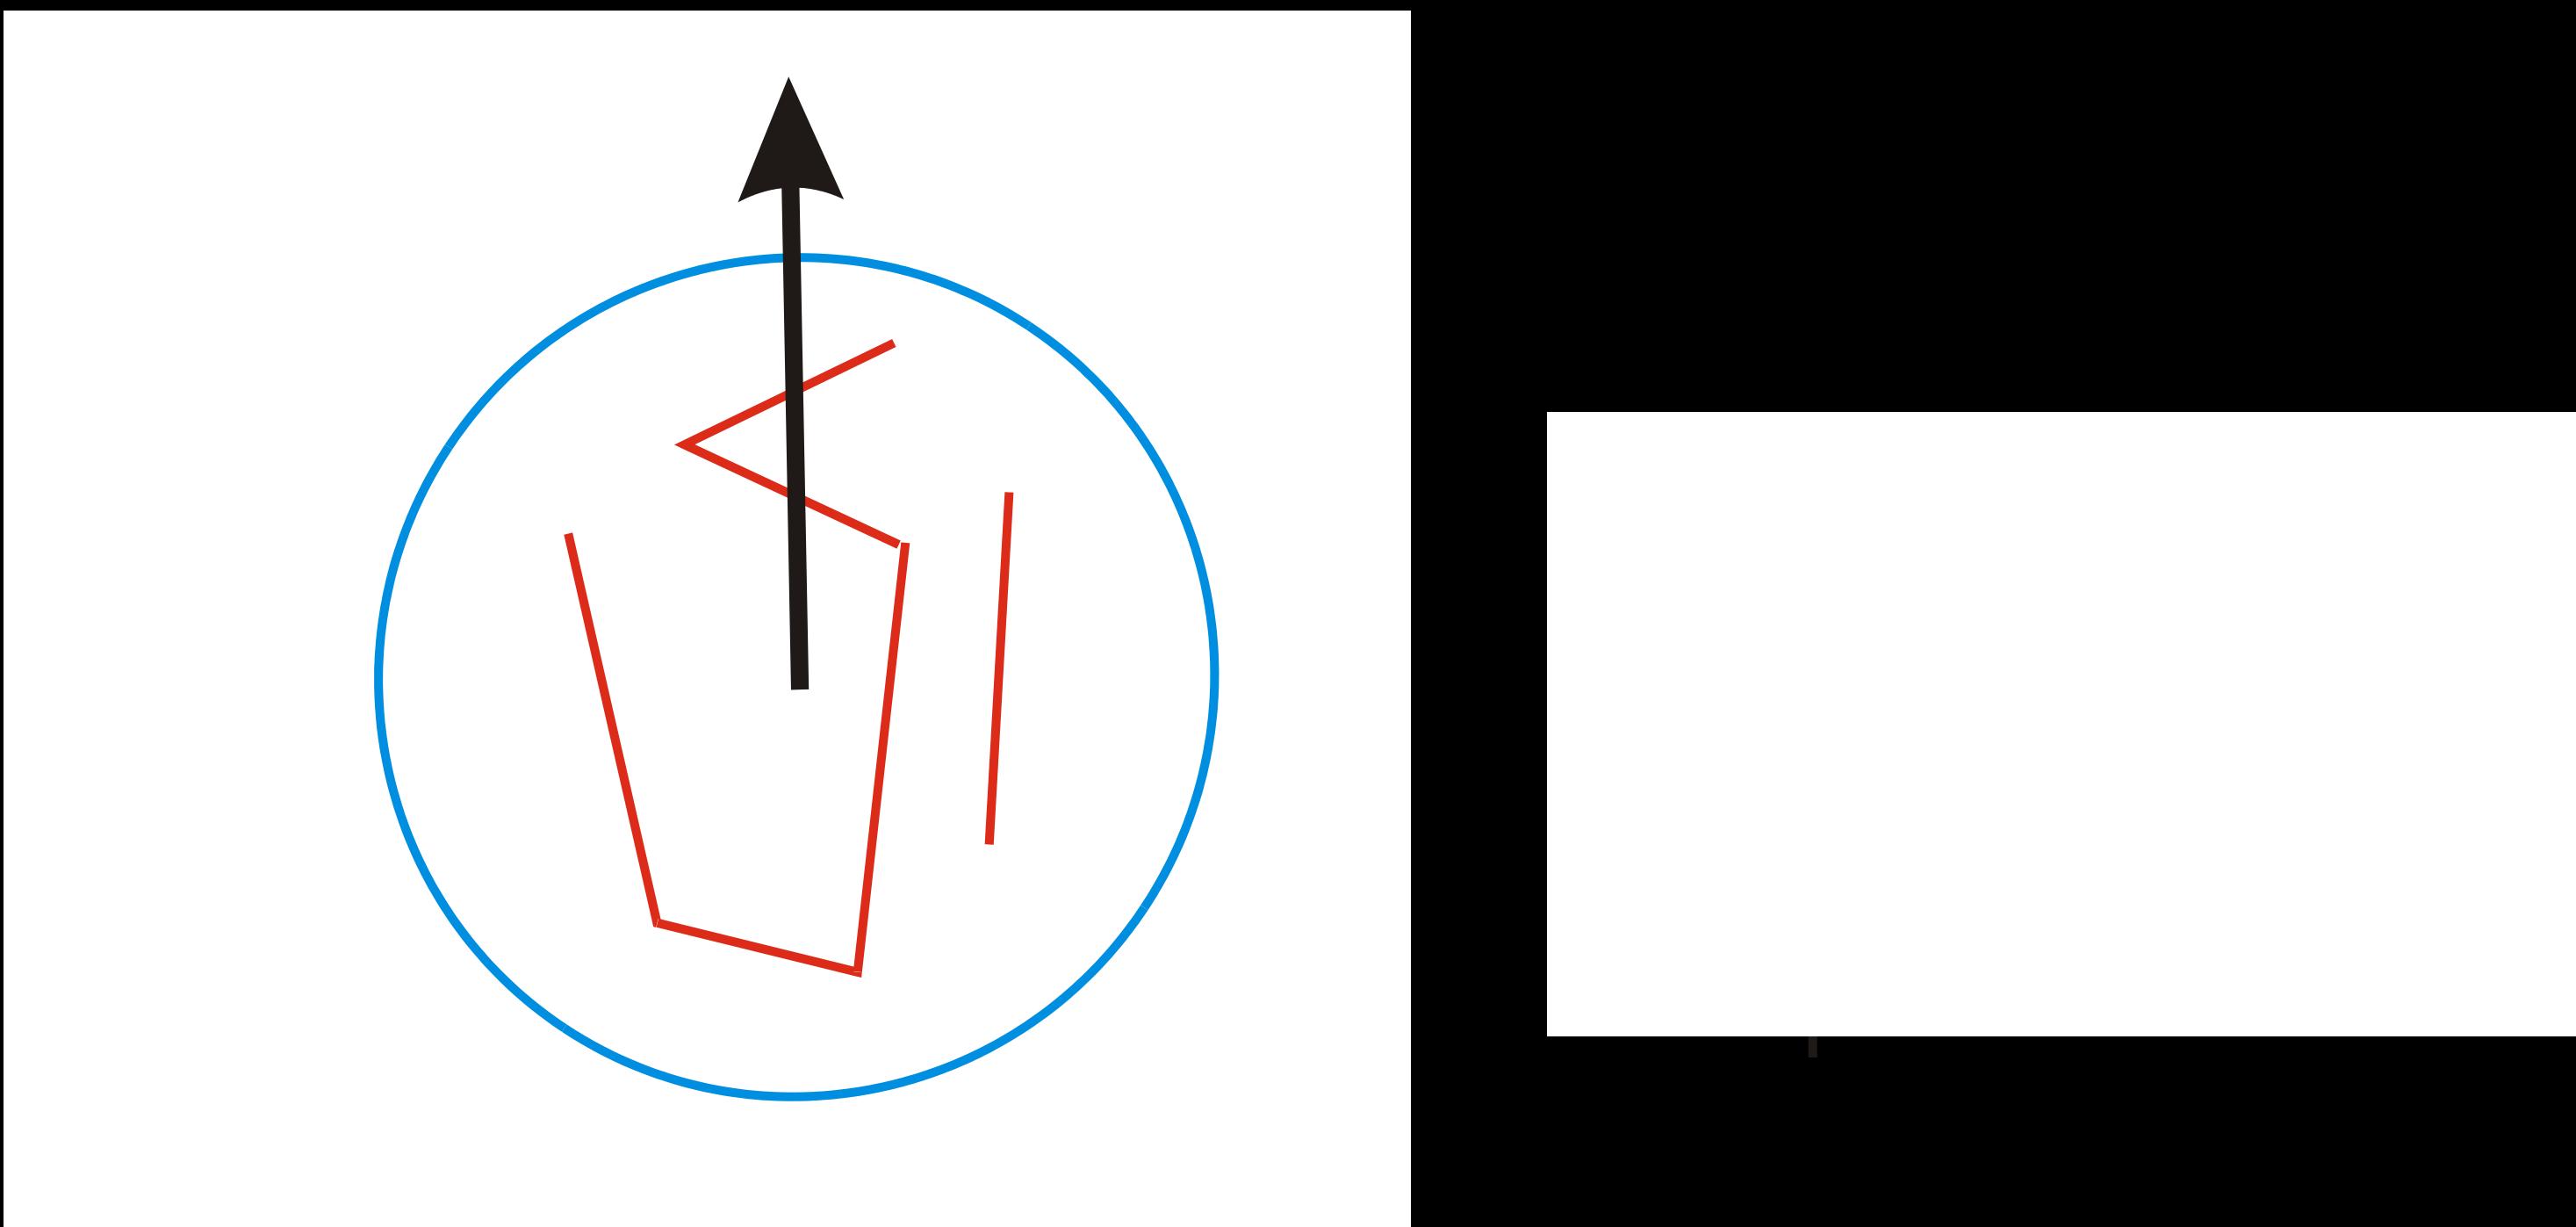
TEXT



ORIENTATION NORMALIZATION

- ▶ Compute orientation histogram
- ▶ Select dominant orientation
- ▶ Normalize: rotate to fixed orientation

[LOWE, SIFT, 1999]



AVAILABLE AT A WEB SITE NEAR YOU...

- ▶ For most local feature detectors, executables are available online:
- ▶ [http://www.robots.ox.ac.uk/~vgg/research/
affine](http://www.robots.ox.ac.uk/~vgg/research/affine)
- ▶ <http://www.cs.ubc.ca/~lowe/keypoints/>
- ▶ <http://www.vision.ee.ethz.ch/~surf>